



NOAA Processed Report NMFS-NWFSC-PR-2020-05

<https://doi.org/10.25923/5wpn-qt71>

***Stock Synthesis* User Manual**

Version 3.30.15

April 2020

U.S. DEPARTMENT OF COMMERCE

National Oceanic and Atmospheric Administration
National Marine Fisheries Service
Northwest Fisheries Science Center

NOAA Processed Report Series NMFS-NWFSC-PR

The Northwest Fisheries Science Center of NOAA's National Marine Fisheries Service uses the NOAA Processed Report NMFS-NWFSC-PR series to disseminate information only. Manuscripts have not been peer-reviewed and may be unedited. Documents within this series represent sound professional work, but do not constitute formal publications. They should only be footnoted as a source of information, and may not be cited as formal scientific literature. The data and any conclusions herein are provisional, and may be formally published elsewhere after appropriate review, augmentation, and editing.

NWFSC Processed Reports are available from the NOAA Institutional Repository, <https://repository.library.noaa.gov>.

Mention throughout this document of trade names or commercial companies is for identification purposes only and does not imply endorsement by the National Marine Fisheries Service, NOAA.

Recommended citation:

(Methot et al. 2020)¹

¹ Methot, R. D., Jr., C. R. Wetzel, I. G. Taylor, and K. Doering. 2020. *Stock Synthesis* User Manual Version 3.30.15. U.S. Department of Commerce, NOAA Processed Report NMFS-NWFSC-PR-2020-05.

<https://doi.org/10.25923/5wpm-qt71>



**NOAA
FISHERIES**

***Stock Synthesis* User Manual Version 3.30.15**

Richard D. Methot, Jr.,¹ Chantel R. Wetzel,² Ian G. Taylor,² and Kathryn Doering²

<https://doi.org/10.25923/5wpn-qt71>

April 2020

¹Office of the Science Director
Northwest Fisheries Science Center
2725 Montlake Boulevard East
Seattle, Washington 98112

²Fishery Resource Analysis and Monitoring Division
Northwest Fisheries Science Center
2725 Montlake Boulevard East
Seattle, Washington 98112

U.S. DEPARTMENT OF COMMERCE

National Oceanic and Atmospheric Administration
National Marine Fisheries Service
Northwest Fisheries Science Center

Contents

Introduction	1
File Organization	3
Input Files	3
Output Files	3
Auxiliary Files	4
Starting SS	6
Converting Files from SS v.3.24	7
Starter File	8
Starter File Options (starter.ss)	9
Forecast File	17
Forecast File Options (forecast.ss)	18
Benchmark Calculations	26
Forecast Recruitment Adjustment	28
Data File	30
Overview of Data File	30
Units of Measure	30
Time Units	31
Seasons	32
Subseasons and Timing of Events	32
Terminology	33
Model Dimensions	33
Fleet Definitions	35
Bycatch Fleets	36
Catch	39
Indices	41
Discard	44
Mean Body Weight or Length	46
Population Length Bins	48
Dirichlet Parameter Number and Effective Sample Sizes	51
Length Composition Data	52
Age Composition Option	54
Age Composition Bins	54
Ageing Error	54
Age Composition Specification	56
Conditional Age-at-Length	57
Mean Length or Body Weight-at-Age	58

Environmental Data	60
Generalized Size Composition Data	61
Tag-Recapture Data	63
Stock (Morph) Composition Data	65
Selectivity Empirical Data (future feature)	66
Excluding Data	66
Data Super-Periods	67
Control File	69
Overview of Control File	69
Parameter Line Elements	70
Terminology	72
Beginning of Control File Inputs	72
Weight-at-Age	74
Settlement Timing for Recruits and Distribution	74
Movement	77
Time Blocks	79
Auto-generation	79
Biology	80
Natural Mortality	80
Growth	81
Maturity-Fecundity	84
Hermaphroditism	86
Parameter offset method	87
Catch Multiplier	87
Ageing Error Parameters	87
Sex ratio	88
Read Biology Parameters	88
Time-Varying Parameter Specification and Setup	93
Time-Varying Growth Considerations	96
Seasonal Biology Parameters	97
Spawner-Recruitment	97
Spawner-Recruitment Functions	98
Spawner-Recruitment Parameter Setup	102
Recruitment Deviation Setup	104
Fishing Mortality Method	110
Initial Fishing Mortality	111
Catchability	112
Mirrored Q with offset	115
Float Q	115
Q Conversion Issues Between SS v.3.24 and v.3.30	116
Selectivity and Discard	116
Reading the Selectivity and Retention Parameters	119

Selectivity Patterns	120
Selectivity Pattern Details	123
Retention	133
Discard Mortality	134
Sex-Specific Selectivity	134
Dirichlet Multinomial Error for Data Weighting	136
Selectivity Time-Varying Options	136
Two-Dimensional Auto-Regressive Selectivity	136
Tag Recapture Parameters	138
Variance Adjustment Factors	140
Lambdas (Emphasis Factors)	142
Controls for Variance of Derived Quantities	144
Optional Inputs	146
Empirical Weight-at-Age (wtatage.ss)	146
runnumber.ss	148
profilevalues.ss	148
Likelihood components	149
Running SS	150
Command Line Interface	150
Example of DOS batch input file	150
Simple Batch	151
Complicated Batch	152
Running Without Estimation	152
Running Parameter Profiles	153
Re-Starting a Run	154
Debugging Tips	154
Keyboard Tips	155
Running MCMC	155
Output Files	157
Standard ADMB output files	157
SS Summary	157
SIS table	157
Derived Quantities	158
Virgin Spawning Biomass vs Unfished Spawning Biomass	158
Metric for Fishing Mortality	159
Equilibrium SPR	159
F std	160
F-at-Age	160
MSY and other Benchmark Items	161

Brief cumulative output	162
Bootstrap Data Files	162
Forecast and Reference Points (Forecast-report.sso)	164
Main Output File, Report.sso	165
Using R To View Model Output (r4ss)	171
Advanced SS Configuration Settings and Advice	176
Using Time-Varying Parameters	176
Parameterizing the Two-Dimensional Autoregressive Selectivity	180
Continuous seasonal recruitment	182
Detailed Information on SS Processes	184
Jitter	184
Parameter Priors	185
Forecast Module: Benchmark and Forecasting Calculations	189
Fishing Mortality in Stock Synthesis	193
Guidance on Population Dynamics Modeling	199
Data Weighting	199
Recruitment Variability and Bias Correction	204
Code Examples	212
Ageing Error Estimation	212
Survival Based SRR Code	212
Random Walk Selectivity: Pattern 17	213
Cubic Spline Selectivity	214
Deviation Link	215
Bibliography	217

Introduction

Fish population (aka "stock") assessment models determine the impact of past fishing on the historical and current abundance of the population, evaluate sustainable rates of removals (catch), and project future levels of catch reflecting one or more risk-averse catch rules. These catch rules are codified in regional Fishery Management Plans according to requirements of the Sustainable Fisheries Act. In the U.S., approximately 500 federally managed fish and shellfish populations are managed under approximately 50 Fishery Management Plans. About 200 of these populations are assessed each year, based on a prioritized schedule for their current status. Despite this, many minor species have never been quantitatively assessed. Although the pace is slower than that for weather forecasting, fish stock assessments are operational models for fisheries management.

Assessment models typically assimilate annual catches, data on fish abundance from diverse surveys and fishery sources, and biological information regarding fish body size and proportions at age. A suite of models is available depending on the degree of data availability and unique characteristics of the fish population or its fishery. Where feasible, environmental time series are used as indicators of changes in population or observation processes, especially to improve the accuracy of the projections of abundance and sustainable catch into the future. Such linkages are based principally on correlations given the challenge of conducting field observations on an appropriate scale. The frontier of model development is in the rapid estimation of parameters to include random temporal effects, in the simultaneous modeling of a suite of interacting species, and in the explicit treatment of the spatial distribution of the population.

Assessment models are loosely coupled to other models. For example, an ocean-temperature or circulation model or benthic-habitat map may be directly included in the pre-processing of the fish abundance survey. A time series of a derived ocean factor, like the North Atlantic Oscillation, can be included as an indicator of a change in a population process. Output of a multi-decadal time series of derived fish abundance can be an input to ecosystem and economic models to better understand cumulative impacts and benefits.

Stock Synthesis (SS) is an age- and size-structured assessment model in the class of models termed integrated analysis models. SS has a population sub-model that simulates a stock's growth, maturity, fecundity, recruitment, movement, and mortality processes, an observation sub-model estimates expected values for various types of data, a statistical sub-model characterizes the data's goodness of fit and obtains best-fitting parameters with associated variance, and a forecast sub-model projects needed management quantities. SS outputs the quantities, with confidence intervals, needed to implement risk-averse fishery control rules. The model is coded in C++ with parameter estimation enabled by automatic differentiation (admb). Windows, Linux, and iOS versions are available. Output processing and associated tools are in R, and a graphical interface is in QT. SS is available from NOAA's VLAB. The rich feature set in SS allows it to be configured for a wide range of situations. SS has become the basis for a large fraction of U.S. assessments and many other assessments around the world.

This manual provides a guide for using SS. The guide contains a description of the input and output files and usage instructions. An overview and technical description of the model itself is in Methot and Wetzel (2013). However, SS has continued to evolve and grow since the publication in 2013, with this manual reflecting the most up to date information regarding SS. The model and a graphical user interface are available from the NOAA VLAB. The VLAB site also provides a user forum for posting questions and for accessing various additional materials. An output processor package, `r4ss`, in R is available for download from CRAN or GitHub.

Additional guidance for new users can be found in the SS document library within the NOAA VLAB website. The "Begin Here - Introduction to Stock Synthesis" folder located in the Document Library contains step-by-step guidance for running Stock Synthesis.

File Organization

Input Files

1. starter.ss: required file containing filenames of the data file and the control file plus other run controls (required).
2. datafile: file containing model dimensions and the data (required)
3. control file: file containing set-up for the parameters (required)
4. forecast.ss: file containing specifications for reference points and forecasts (required)
5. ss.par: previously created parameter file that can be read to overwrite the initial parameter values in the control file (optional)
6. wtatage.ss: file containing empirical input of body weight by fleet and population and empirical fecundity-at-age (optional)
7. runnumber.ss: file containing a single number used as run number in output to CumReport.sso and in the processing of profilevalues.ss (optional)
8. profilevalues.ss: file contain special conditions for batch file processing (optional)

Output Files

1. data.ss_new: Contains a user-specified number of datafiles, generated through a parametric bootstrap procedure, and written sequentially to this file.
2. control.ss_new: Updated version of the control file with final parameter values replacing the initial parameter values.
3. starter.ss_new: New version of the starter file with annotations.
4. Forecast.ss_new: New version of the forecast file with annotations.
5. warning.sso: This file contains a list of warnings generated during program execution.
6. echoinput.sso: This file is produced while reading the input files and includes an annotated echo of the input. The sole purpose of this output file is debugging input errors.

7. Report.sso: This file is the primary report file.
8. ss_summary.sso: Output file that contains all the likelihood components, parameters, derived quantities, total biomass, summary biomass, and catch. This file offers an abridged version of the report file that is useful for quick model evaluation. This file is only available in SS v.3.30.08.03 and greater.
9. CompReport.sso: Observed and expected composition data in a list-based format.
10. Forecast-report.sso: Output of management quantities and for forecasts.
11. CumReport.sso: This file contains a brief version of the run output, output is appended to current content of file so results of several runs can be collected together. This is useful when a batch of runs is being processed.
12. Covar.sso: This file replaces the standard ADMB ss.cor with an output of the parameter and derived quantity correlations in database format.
13. ss.par: This file contains all estimated and fixed parameters from the model run.
14. ss.std, ss.rep, ss.cor etc.: Standard ADMB output files.
15. checkup.sso: Contains details of selectivity parameters and resulting vectors. This is written during the first call of the objective function.
16. Gradient.dat: New for SS v.3.30, this file shows parameter gradients at the end of the run.
17. rebuild.dat: Output formatted for direct input to Andre Punt's rebuilding analysis package. Cumulative output is output to REBUILD.SS (useful when doing MCMC or profiles).
18. SIS_table.sso: Output formatted for reading into the NMFS Species Information System.
19. Parmtrace.sso: Parameter values at each iteration.
20. posteriors.sso, derived_posteriors.sso, posterior_vectors.sso: Files associated with MCMC.

Auxiliary Files

These files are additional files (e.g. excel files) which allow for exploration or understanding of specific parameterization which can assist in selecting appropriate starting values. These files are available for download from the Vlab website.

1. SS3-OUTPUT.xls: Excel file with macros to read report.sso and display results.

2. SS_330_Control_Setup.xls:
 - (a) Shows how to setup an example control file for SS.
3. SS_330_Data_Input.xls:
 - (a) Shows how to setup an example data input file for SS.
4. SS_330_Starter&Forecast.xls:
 - (a) Shows how to setup an example starter and forecast input files for SS.

Starting SS

SS is typically run through the command line interface, although it can also be called from another program, R, the SS-GUI or a script file (such as a DOS batch file). SS is compiled for Windows, Mac, and Linux operating systems. The memory requirements depend on the complexity of the model you run, but in general, SS will run much slower on computers with inadequate memory. See the Running SS section on page 150 for additional notes on methods of running SS.

Communication with the program is through text files. When the program first starts, it reads the file `starter.ss`, which typically must be located in the same directory from which SS is being run. The file `starter.ss` contains required input information plus references to other required input files, as described in the File Organization section on page 3. The names of the control and data files must match the names specified in the `starter.ss` file. File names, including `starter.ss`, are case-sensitive on Linux and Mac systems but not on Windows. The `echoinput.sso` file outputs how the executable read each input file and can be used for troubleshooting when trying to setup a model correctly. Output from SS consists of text files containing specific keywords. Output processing programs, such as the SS GUI, Excel, or R can search for these keywords and parse the specific information located below that keyword in the text file.

Converting Files from SS v.3.24

Converting files from SS v.3.24 to SS v.3.30 can be easily performed by using the program `ss_trans.exe`. The following file structure and steps are recommended for converting model files:

1. Create "transition" folder. Place the 4 main model files (control, data, starter, and forecast) from version SS v.3.24 within the transition folder along with the SS transition executable (`ss_trans.exe`). One tip is to use the `control.ss_new` from the SS v.3.24 estimated model rather than the `control.ss` file which will set all parameter values at the previous estimated maximum likelihood estimated (MLE) parameters. Run the transition executable with phase = 0 within the starter file with the read par file turned off (option 0).
2. Create "converted" folder. Place the `ss_new` (`data.ss_new`, `control.ss_new`, `starter.ss_new`, `forecast.ss_new`) files created by the transition executable contained within the "transition" folder into this new folder. Rename the `ss_new` files to the appropriate suffixes and change the names in the `starter.ss` file accordingly.
3. Review the control (`control.ss_new`) file to determine that all model functions converted correctly. The structural changes and assumptions for a couple of the advanced model features are too complicated to convert automatically. See below for some known features that may not convert. When needed, it is recommended to modify the `control.ss_new` file, the converted control file, for only the features that failed to convert properly.
4. Change the max phase to a value greater than the last phase in which the a parameter is set to estimated within the control file. Run the new SS v.3.30 executable (`ss.exe`) within the "converted" folder using the renamed `ss_new` files created from the transition executable.
5. Compare likelihood and model estimates between the SS v.3.24 and SS v.3.30 model versions.

There are some options that have been substantially changed in SS v.3.30 which impedes the automatic converting of SS v.3.24 model files. Known examples of SS v.3.24 options that cannot be converted, but for which better alternatives are available in SS v.3.30 are:

1. The use of Q deviations,
2. Complex birth seasons,
3. Environmental effects on spawner-recruitment parameters,
4. Setup of time-varying quantities for models that used the no-longer-available features (e.g., logistic bound constraint).

Starter File

SS begins by reading the file `starter.ss`. The term `COND` appears in the "Typical Value" column of this documentation (it does not actually appear in the model files), it indicates that the following section is omitted except under certain conditions, or that the factors included in the following section depend upon certain conditions. In most cases, the description in the definition column is the same as the label output to the `ss_new` files.

Starter File Options (starter.ss)

Value	Options	Description
#C this is a starter comment	Must begin with #C then rest of the line is free form	All lines in this file beginning with #C will be retained and written to the top of several output files
data_file.dat		File name of the data file
control_file.ctl		File name of the control file
0	Initial Parameter Values: 0 = use values in control file; and 1 = use ss.par after reading setup in the control file.	Do not set equal to 1 if there have been any changes to the control file that would alter the number or order of parameters stored in the ss.par file. Values in ss.par can be edited, carefully. Do not run ss_trans.exe from a ss.par from SS v.3.24.
1	Run display detail: 0 = none other than ADMB outputs; 1 = one brief line of display for each iteration; and 2 = fuller display per iteration.	With option 2, the display shows value of each -logL component for each iteration and it displays where crash penalties are created
1	Detailed age-structure report: 0 = minimal (no Report file); 1 = include all output; and 2 = brief output.	Option 0 will forgo the writing of the Report file, but the ss_summary file will be written that has minimal derived and estimated quantities. This is a useful option for some data-limited assessment approaches (e.g., XSSS or SSS). Option 1 will write out the full Report file. Option 2 will write out select items in the Report file and will omit some more detailed sections (e.g., numbers-at-age).

Value	Options	Description
0	Write 1st iteration details: 0 = omit; and 1 = write detailed intermediate calculations to echoinput.sso during first call.	This output is largely unformatted and undocumented and is mostly used by the developer.
0	Parameter Trace: 0 = omit; 1 = write good iteration and active parameters; 2 = write good iterations and all parameters; 3 = write every iteration and all parameters; and 4 = write every iteration and active parameters.	This controls the output to parmtrace.sso. The contents of this output can be used to determine which values are changing when a model approaches a crash condition. It also can be used to investigate patterns of parameter changes as model convergence slowly moves along a ridge.
1	Cumulative Report: 0 = omit; 1 = brief; and 2 = full.	Controls reporting to the file Cumreport.sso. This cumulative report is most useful when accumulating summary information from likelihood profiles or when simply accumulating a record of all model runs within the current subdirectory
1	Full Priors: 0 = only calculate priors for active parameters; and 1 = calculate priors for all parameters that have a defined prior.	Turning this option on (1) adds the log likelihood contribution from all prior values for fixed and estimated parameters to the total negative log likelihood. With this option off (0), the total negative log likelihood will include the log likelihood for priors for only estimated parameters.

Value	Options	Description
1	Soft Bounds: 0 = omit; and 1 = use.	This option creates a weak symmetric beta penalty for the selectivity parameters. This becomes important when estimating selectivity functions in which the values of some parameters cause other parameters to have negligible gradients, or when bounds have been set too widely such that a parameter drifts into a region in which it has negligible gradient. The soft bound creates a weak penalty to move parameters away from the bounds.
1	Data File Output: 0 = none; 1 = output an annotated replicate of the input data file; 2 = add a second data file containing the model's expected values with no added error; and 3+ = add N-2 parametric bootstrap data files.	All output files are sequentially output to data.ss_new and will need to be parsed by the user into separate data files. The output of the input data file makes no changes, so retains the order of the original file. Output files 2-N contain only observations that have not been excluded through use of the negative year denotation, and the order of these output observations is as processed by the model. The N obs values are adjusted accordingly. At this time, the tag recapture data is not output to data.ss_new.
8	Turn off estimation: -1 = exit after reading input files; 0 = exit after one call to the calculation routines and production of sso and ss_new files; and <positive value> = exit after completing this phase.	The 0 option is useful for (1) quickly reading in a messy set of input files and producing the annotated control.ss_new and data.ss_new files, or (2) examining model output based solely on input parameter values. Similarly, the value option allows examination of model output after completing a specified phase. Also see usage note for restarting from a specified phase.
1000	MCMC burn interval	Number of iterations to discard at the start of an MCMC run.
200	MCMC thin interval	Number of iterations to remove between the main period of the MCMC run.

Value	Options	Description
0.0	Jitter: 0 = no jitter done to starting values; and >0 starting values will vary with larger jitter values resulting in larger changes from the parameter values in the control or par file.	The jitter function has been revised with SS v.3.30. Starting values are now jittered based on a normal distribution with the $\text{pr}(P_{\text{MIN}}) = 0.1\%$ and the $\text{pr}(P_{\text{MAX}}) = 99.9\%$. A positive value here will add a small random jitter to the initial parameter values. When using the jitter option, care should be given when defining the low and high bounds for parameter values and particularly -999 or 999 should not be used to define bounds for estimated parameters.
-1	SD Report Start: -1 = begin annual SD report in start year; and <year> = begin SD report this year.	
-1	SD Report End: -1 = end annual SD report in end year; -2 = end annual SD report in last forecast year; and <value> = end SD report in this year.	
2	Extra SD Report Years: 0 = none; and <value> = number of years to read.	In a long time series application, the model variance calculations will be smaller and faster if not all years are included in the SD reporting. For example, the annual SD reporting could start in 1960 and the extra option could select reporting in each decade before then.
COND: If Extra SD report years > 0		
1940 1950	Vector of years for additional SD reporting. The number of years need to equal the value specified in the above line (Extra SD Reporting).	

Value	Options	Description
0.0001	Final convergence	This is a reasonable default value for the change in log likelihood denoting convergence. For applications with much data and thus a large total log likelihood value, a larger convergence criterion may still provide acceptable convergence
0	Retrospective year: 0 = none; and -x = retrospective year relative to end year.	Adjusts the model end year and disregards data after this year. May not handle time varying parameters completely.
0	Summary biomass min age	Minimum integer age for inclusion in the summary biomass used for reporting and for calculation of total exploitation rate.
1	Depletion basis: 0 = skip; 1 = $X \cdot SB_0$; 2 = $X \cdot SB_{MSY}$; 3 = $X \cdot SB_{styr}$; and 4 = $X \cdot SB_{endyr}$.	Selects the basis for the denominator when calculating degree of depletion in SB. The calculated values are reported to the SD report. Relative to virgin spawning biomass. Relative to spawning biomass that achieves MSY. Relative to model start year spawning biomass. Relative to spawning biomass in the model end year.
1	Fraction (X) for depletion denominator	Value for use in the calculation of the ratio for $SSB_y / (X \cdot SSB_0)$.
1	SPR report basis: 0 = skip; 1 = use $1 - SPR_{target}$; 2 = use $1 - SPR$ at MSY;	SPR is the equilibrium SSB per recruit that would result from the current year's pattern and intensity of F's. The quantities identified by 1, 2, and 3 here are all calculated in the benchmarks section. Then the one specified here is used as the selected

Value	Options	Description
	3 = use 1-SPR at B_{target} ; and 4 = no denominator, so report actual 1-SPR values.	Denominator in a ratio with the annual value of $(1 - \text{SPR})$. This ratio (and its variance) is reported to the SD report output for the years selected above in the SD report year selection.
4	Annual F units: 0 = skip; 1 = exploitation rate in biomass; 2 = exploitation rate in numbers; 3 = sum(apical F's by fleet); 4 = population F for range of ages; and 5 = unweighted average F for range of ages.	In addition to SPR, an additional proxy for annual F can be specified here. As with SPR, the selected quantity will be calculated annually and in the benchmarks section. The ratio of the annual value to the selected (see F report basis below) benchmark value is reported to the SD report vector. Options 1 and 2 use total catch for the year and summary abundance at the beginning of the year, so combines seasons and areas. But if most catch occurs in one area and there is little movement between areas, this ratio is not informative about the F in the area where the catch is occurring. Option 3 is a simple sum of the full F's by fleet, so may provide non-intuitive results when there are multi areas or seasons or when the selectivities by fleet do not have good overlap in age. Option 4 is a real annual F calculated as a numbers weighted F for a specified range of ages (read below). The F is calculated as $Z - M$ where Z and M are each calculated as $\ln(N_{t+1}/N_t)$ with and without F active, respectively. The numbers are summed over all biology morphs and all areas for the beginning of the year, so subsumes any seasonal pattern.
COND: If F std reporting ≥ 4		Specify range of ages. Upper age must be less than max age because of incomplete handling of the accumulator age for this calculation.
3 7	Age range if F std reporting = 4.	
1	F report basis: 0 = not relative, report raw values; 1 = use F std value relative to $\text{SPR}_{\text{target}}$; 2 = use F std value relative to F_{MSY} ; and	Selects the denominator to use when reporting the F std report values. Note that order of these options differs from the biomass report basis options.

Value	Options	Description
	3 = use F std value relative to $F_{Btarget}$.	
0.01	MCMC output detail: 0 = default; 1 = output likelihood components and associated lambda values; 2 = expanded output; and 3 = make output subdirectory for each MCMC vector.	Specify format of MCMC output. This input requires the specification of two items; the output detail and a bump value to be added to the $\ln(R0)$ in the first call to MCMC. A bias adjustment of 1.0 is applied to recruitment deviations in the MCMC phase, which could result in reduced recruitment estimates relative to the MLE when a lower bias adjustment value is applied. A small value, called the "bump", is added to the $\ln(R0)$ for the first call to MCMC in order to prevent the stock from hitting the lower bounds when switching from MLE to MCMC. If you wanted to select the default output option and apply a bump value of 0.01 this is specified by 0.01 where the integer value represents the output detail and the decimal is the bump value.
0	Age-length-key (ALK) tolerance level, 0 >= values required	Value of 0 will not apply any compression. Values > 0 (e.g. 0.0001) will apply compression to the ALK which will increase the speed of calculations. The size of this value will impact the run time of your model, but one should be careful to ensure that the value used does not appreciably impact the estimated quantities relative to no compression of the ALK. The suggested value if applied is 0.0001.
	COND: Seed Value (i.e., 1234)	Specify a seed for data generation. This feature is not available in versions prior to 3.30.15. This is an optional input value which allows for the specification of a random number seed value. If you do not want to specify a seed, skip this input line and end the reading of the starter file with the check value (3.30).

Value	Options	Description
3.30	Model version check value.	A value of 3.30 indicates that the control and data files are currently in SS v3.30 format and a value of 999 indicates that the control and data file are in a previous SS v.3.24 version. The ss_trans.exe executable should be used and will convert the 3.24 version files to the new format in the control.ss_new and data.ss_new files. All ss_new files are in the SS v.3.30 format, so starter.ss_new has SS v.3.30 on the last line. The mortality-growth parameter section has a new sequence, so SS v.3.30 cannot read a ss.par file produced by SS v.3.24 and earlier, so please ensure that read par file option at the top of the starter file is set to 0. The Converting Files from SS v.3.24 section has additional information on model features that may impede file conversion.
End of Starter File		

Forecast File

The specification of options for forecasts is contained in the mandatory input file named `forecast.ss`. Additional details on the forecast module process are provided on page 189.

The term COND appears in the "Typical Value" column of this documentation (it does not actually appear in the model files), it indicates that the following section is omitted except under certain conditions, or that the factors included in the following section depend upon certain conditions. In most cases, the description in the definition column is the same as the label output to the `ss_new` files.

Forecast File Options (forecast.ss)

Value	Options	Description
1	Benchmarks/Reference Points: 0 = omit; 1 = calculate F_{SPR} , $F_{Btarget}$, and F_{MSY} ; and 2 = calculate F_{SPR} , F_{MSY} , $F_{0.10}$.	SS checks for consistency of the Forecast specification and the benchmark specification. It will turn benchmarks on if necessary and report a warning.
1	MSY Method: 1 = F_{SPR} as proxy; 2 = calculate F_{MSY} ; 3 = $F_{Btarget}$ as proxy or $F_{0.10}$; and 4 = $F_{end\ year}$ as proxy.	Specifies basis for calculating a single population level F_{MSY} value.
0.45	SPR_{target}	SS searches for F multiplier that will produce this level of spawning biomass per recruit (reproductive output) relative to unfished value.
0.40	Relative Biomass Target	SS searches for F multiplier that will produce this level of spawning biomass relative to unfished value. This is not “per recruit” and takes into account the spawner-recruitment relationship.
0 0 0 0 0 0 0 0 0 0	Benchmark Years: -999: start year; >0: absolute year; and <= 0: year relative to end year.	Requires 10 values, beginning and ending years for (1,2) biology (e.g., growth, natural mortality, maturity, fecundity), (3,4) selectivity, (5,6) relative Fs, (7,8) movement and recruitment distribution; (9,10) stock-recruitment parameters for averaging years in calculating benchmark quantities.

Value	Options	Description
1	Benchmark Relative F Basis: 1 = use year range; and 2 = set range for relF same as forecast below.	The specification does not affect year range for selectivity and biology.
2	Forecast: -1 = none, no forecast years; 0 = simple, single forecast year calculated; 1 = use F_{SPR} ; 2 = use F_{MSY} ; 3 = use $F_{Btarget}$ or $F_{0.10}$; 4 = set to average F scalar for the forecast relative F years below; and 5 = input annual F scalar.	This input is required but is ignored if benchmarks are turned off. This determines how forecast catches are calculated and removed from the population which is separate from the "MSY Method" above. If F_{MSY} is selected, it uses whatever proxy (e.g., F_{SPR} or F_{BTGT}) is selected in the "MSY Method" row.
10	N forecast years (must be ≥ 1)	At least one forecast year now required which differs from SS v.3.24 that allowed zero forecast years.
1	F scalar	Only used if Forecast option = 5 (input annual F scalar), but is a required line in the forecast file.
0 0 0 0 0 0	Forecast Years: >0 = absolute year; and <= 0 = year relative to end year.	Requires 6 values: beginning and ending years for selectivity, relative Fs, and recruitment distribution that will be used to create averages to use in forecasts. In future, hope to allow random effects to propagate into forecast. Please note, relative F for bycatch only fleets is scaled just like other fleets. More options for this in future.

Value	Options	Description
0	Forecast Selectivity Option: 0 = forecast selectivity is mean from year range; and 1 = forecast selectivity from annual time-varying parameters.	Determines the selectivity used in the forecast years. Selecting 1 will allow for application of time-varying selectivity parameters (e.g., random walk) to continue into the forecast period.
1	Control Rule: 1 = catch as function of SSB, buffer on F; 2 = F as function of SSB, buffer on F; 3 = catch as function of SSB, buffer on catch; and 4 = F is a function of SSB, buffer on catch.	Used to apply reductions ("buffer") to either the catch or F based on the control rule during the forecast period.
0.40	Control Rule Upper Limit	Relative biomass level to unfished biomass above which F is constant at control rule F_{target} .
0.10	Control Rule Lower Limit	Relative biomass level to unfished biomass below which F is set to 0 (management threshold).
0.75	Control Rule Buffer (multiplier or -1)	Control rule catch or F_{target} as a fraction of selected catch or F_{MSY} proxy. A value of -1 will allow the user to input a forecast year specific control rule fraction (added in v. 3.30.13).

Value	Options	Description
COND -1: Conditional input for annual control rule buffer	2019 0.8 2020 0.6 2021 0.5 -9999 0	Year and control rule buffer value. Can enter a value for each year, or starting sequence of years. The final control rule buffer value will apply to all sequent forecast years.
3	Number of forecast loops (1,2,3)	SS sequentially goes through the forecast up to three times. Maximum number of forecast loops: 1=OFL only, 2=ABC control rule and buffers, 3=set catches equal to control rule or input catch and redo forecast implementation error.
3	First forecast loop with stochastic recruitment	If this is set to 1 or 2, then OFL and ABC will be calculated as if there was perfect knowledge about recruitment deviations in the future.
0	Forecast recruitment: 0 = spawner recruit curve; 1 = value*(spawner recruit curve); 2 = value*(virgin recruitment); and 3 = recent mean from year range above.	Option 0, ignore input and do forecast recruitment as before SS v.3.30.10, if 1, then use next value as a multiplier applied after env/block/regime is applied, if 2, then use value as multiplier times adjusted virgin recruitment (after time-varying adjustments to R0), and if 3, then use value as the number of years from end of main recruitment deviations to average (mean is the recruitments, not the deviations). Need to set phase to -1 in control to get constant recruitment in MCMC.
1	Scalar or number of years of recent main recruitments to average.	This input depends upon option selected directly above. If option 1 or 2 selected this value should be a scalar value to be applied to recruitment. If option 3 is selected above this should be input as the number of years to average recruitment.

Value	Options	Description
0	Forecast loop control #5	Reserved for future model features.
2015	First year for caps and allocations	Should be after years with fixed inputs.
0	Implementation Error	The standard deviation of the log of the ratio between the realized catch and the target catch in the forecast. (set value >0.0 to cause implementation error deviations to be an estimated parameter that will add variance to forecast).
0	Rebuilder: 0 = omit U.S. West Coast rebuilder output; and 1 = do U.S. West Coast rebuilder output.	Creates a rebuild.dat file to be used for U.S. West Coast groundfish rebuilder program.
2004	Rebuilder catch (Year Declared): >0 = year first catch should be set to zero; and -1 = set to 1999.	Input line is required even if Rebuilder = 0, specified in the line above.
2004	Rebuilder start year (Year Initial): >0 = year for current age structure; and -1 = set to end year +1.	Input line is required even if Rebuilder = 0, specified two line above.
1	Fleet Relative F: 1 = use first-last allocation year; and 2 = read season(row) x fleet (column) set below.	

Value	Options	Description
2	Basis for maximum forecast catch: 2 = total catch biomass; 3 = retained catch biomass; 5 = total catch numbers; and 6 = retained total numbers.	The maximum basis for forecasted catch will be implemented for the for the "First year for caps and allocations" selected above. The maximum catch (biomass or numbers) by fleet is specified below on the "Maximum total forecast catch by fleet" line.
COND 2: Conditional input for fleet relative F (Enter: Season, Fleet, Relative F)		
1 1 0.6	Fleet allocation by relative F fraction.	The fraction of the forecast F value. For a multiple area model user must define a fraction for each fleet and each area. The total fractions must sum to one over all fleets and areas.
1 2 0.4		
-9999 0 0	Terminator line	
1 50	Maximum total forecast catch by fleet (in units specified above total catch/numbers, retained catch/numbers)	Enter fleet number and its maximum value. Last line of the entry must have fleet number = -9999.
-9999 -1		
-9999 -1	Maximum total catch by area -1 = no maximum	Enter area number and its max. Last line of the entry must have area number = -9999.
1 1	Fleet assignment to allocation group	Enter list of fleet number and its allocation group number if it is in a group. Last line of the entry must have fleet number = -9999.
-9999 -1		

Value	Options	Description
COND: if N allocation groups is >0		Enter a year and the allocation fraction to each group for that year. SS will fill those values to the end of the forecast, then read another year from this list. Terminate with -9999 in year field. Annual values are rescaled to sum to 1.0.
2002 1	Allocation to each group for each year of the forecast	
-9999 1		
-1	<p>Basis for forecast catch:</p> <p>-1 = Read basis with each observation, allows for a mixture of dead, retained, or F basis by different fleets for the fixed catches below;</p> <p>2 = Dead catch (retained + discarded);</p> <p>3 = Retained catch; and</p> <p>99 = Input apical F (the apical F value for the model years can be found in the EXPLOITATION section in the Report file).</p>	The dead or retained value in the forecast catch inputs will be interpreted in terms of numbers or biomass based on the units of the input catch for each fleet.
COND: == -1	Forecasted catches - enter one line per number of fixed forecast year catch	
2012 1 1 1200 2	Year & Season & Fleet & Catch or F value & Basis	
2013 1 1 1400 3	Year & Season & Fleet & Catch or F value & Basis	
-9999 0 0 0 0	Indicates end of inputted catches to read	
COND: > 0	Forecasted catches - enter one line per number of fixed forecast year catch	
2012 1 1 1200	Year & Season & Fleet & Catch or F value	
2013 1 1 1200	Year & Season & Fleet & Catch or F value	

Value	Options	Description
-9999 0 0 0	Indicates end of inputted catches to read	
999	End of Input	
End of Forecast File		

Benchmark Calculations

This feature of SS is designed to calculate an equilibrium fishing rate intended to serve as a proxy for the fishing rate that would provide maximum sustainable yield (MSY). Then in the forecast module these fishing rates can be used in the projections.

Four reference points can be calculated by SS. The first of which is the estimate of F_{MSY} within the model, while the others use proxies or an alternative estimated point.

- F_{MSY} : Search for the F that produces maximum equilibrium (e.g. dead catch).
- F_{SPR} : Search for the F that produces spawning biomass per recruit this is a specific fraction, termed SPR_{target} , of spawning biomass per recruit under unfished conditions. Note that this is in relative terms so it does not take into account the spawner-recruit relationship.
- $F_{Btarget}$: Search for the F that produces an absolute spawning biomass that is a specified fraction, termed relative biomass target, of the unfished spawning biomass. Note that this is in absolute terms so takes into account the spawner-recruit relationship.
- $F_{0.10}$: Search for the F that produces a slope in yield per recruit, dY/dF , that is 10% of the slope at the origin. Note that with SS, this option is mutually exclusive with $F_{Btarget}$. Only one will be calculated and the one that is calculated can serve as the proxy for F_{MSY} and forecasting.

Estimation

Each of the potential reference points is calculated by searching across a range of F multiplier levels, calculating equilibrium biomass and catch at that F , using Newton-Raphson method to calculate a better F multiplier value, and iterating a fixed number of times to achieve convergence on the desired level.

Calculations

The calculation of equilibrium biomass and catch uses the same code that is used to calculate the virgin conditions and the initial equilibrium conditions. This equilibrium calculation code takes into account all morph, timing, biology, selectivity, and movement conditions as they apply while doing the time series calculations. You can verify this by running SS to calculate F_{MSY} then hardwire initial F to equal this value, use the F_{method} approach 2 so each annual F is equal to F_{MSY} and then set forecast F to be the same F_{MSY} . Then run SS without estimation and no recruitment deviations. You should see that the population has an initial equilibrium abundance equal to B_{MSY} and stays at this level during the time series and forecast.

Catch Units

For each fleet, SS always calculates catch in terms of biomass (mt) and numbers (1000s) for encountered (selected) catch, dead catch, and retained catch. These three categories differ only when some fleets have discarding or are designated as a bycatch fleet. SS uses total dead catch biomass as the quantity that is principally reported and the quantity that is optimized when searching for F_{MSY} . The quantity "dead catch" may occasionally be referred to as "yield".

Biomass Units

The principle measure of fish abundance, for the purpose of reference point calculation, is female reproductive output. This is referred to as SSB (spawning stock biomass) and sometimes just "B" because the typical user settings have one unit of reproductive output (fecundity) per kg of mature female biomass. So when the output label says B_{MSY} , this is actually the female reproductive output at the proxy for F_{MSY} .

Fleet Allocation

An important concept for the reference point calculation is the allocation of fishing rate among fleets. Internally, this is $Bmark_relF(f, s)$ and it is the fraction of the F multiplier assigned to each fleet, f and season, s . The value, $F_multiplier * Bmark_relF(f, s)$, is the F level for a particular fleet in a particular season and for the age that has a selectivity of 1.0. Other ages will have different F values according to their selectivity.

- The $Bmark_relF$ values can be calculated by SS from a range of years specified in the input for Benchmark Years or it can be set to be the same as the $Forecast_RelF$, which in turn can be based on a range of years or can be input as a set of fixed values.
- Note that for Bycatch Fleets, the F's calculated by application of $Bmark_relF$ for a bycatch fleet can be overridden by a F value calculated from a range of years or a fixed F value that is input by the user. If such an override is selected for a bycatch fleet, that F value is not adjusted by changes to the F multiplier. This allows the user to treat a bycatch fleet as a constant background F while the optimal F for other fleets is sought. Also for bycatch fleets, there is user control for whether or not the dead catch from the bycatch fleet is included in the total dead catch that is optimized when searching for F_{MSY} .

Virgin vs. Unfished

The concept of unfished spawning biomass, SSB_{unf} , is important to the reference points calculations. Unfished spawning biomass can be potentially different than virgin spawning biomass, SSB_{virgin} .

- Virgin spawning biomass is calculated from the parameter values associated with the start year of the model configuration and it serves as the basis from which the population model starts and the basis for calculation of stock depletion.
- Unfished spawning biomass can be calculated for any year or range of years, so can change over time as R_0 , steepness, or biological parameters change.
- In the reference points calculation, the Benchmark Years input specifies the range of time over which various quantities are averaged to calculate the reference points. For biology, selectivity, F 's, and movement the values being averaged are the year-specific derived quantities. But for the stock-recruitment parameters (R_0 and steepness), the parameter values themselves are averaged over time.
- During the time series or forecast, the current year's unfished spawning output (SSB_{unf}) is used as the basis for the spawner-recruitment curve against which deviations from the spawner-recruitment curve are applied. So, if R_0 is made time-varying, then the spawner-recruit curve itself is changed. However, if the regime shift parameter is time-varying, then this is an offset from the spawner-recruitment curve and not a change in the curve itself. Changes in R_0 will change year-specific reference points and change the expected value for annual recruitments, but changes in regime shift parameter only change the expected value for annual recruitments.
- In reporting the time series of depletion level, the denominator can be based on virgin spawning output (SSB_{virgin}) or B_{MSY} . Note that B_{MSY} is based on unfished spawning output (SSB_{unf}) for the specified range of Benchmark years, not on SSB_{virgin} .

Forecast Recruitment Adjustment

Recruitment during the forecast years sometimes needs to be set at a level other than that determined by the spawner-recruitment curve. One way to do this is by an environmental or block effect on the regime shift parameter. A more straightforward approach is now provided by the special forecast recruitment feature described here. There are 4 options provided for this feature. These are:

- 0 = Do nothing: this is the default and will invoke no special treatment for the forecast recruitments.
- 1 = Multiplier on spawner-recruitment: the expected recruitment from the stock recruitment relationship is multiplied by this factor.
 - This is a multiplier, so null effect comes from a value of 1.0.

- The order of operations is to apply the SRR, then the regime effect, then this special forecast effect, then bias adjustment, then the deviations.
- In the spawner recruit output of the report.sso there are 4 recruitment values stored.
- 2 = Multiplier on virgin recruitment: The virgin recruitment is multiplied by this factor.
 - This is a multiplier, so null effect comes from a value of 1.0.
 - The order of operations is to apply any environmental or block effects to R0, then apply the special forecast effect, then bias adjustment, then the deviations.
 - Note that environmental or block effects on R0 are rare and are different than environment or block effects on the regime parameter.
- 3 = Mean recent recruitment: calculate the mean recruitment and use it during the forecast period.
 - Note that bias adjustment is not applied to this mean because the values going into the mean have already been bias adjusted.

This feature affects the expected recruitment in all years after the last year of the main recruitment deviations. This means that if the last year of main recruitment deviations is before end year, then the last few recruitments, termed "late", are also affected by this forecast option. For example, option 3 would allow you to set the last 2 years of the time series and all forecast years to have recruitment equal to the mean recruitment for the last 10 years of the main recruitment era.

Data File

Overview of Data File

1. Dimensions (years, ages, number of fleets, number of surveys, etc.)
2. Fleet and survey names, timing, etc.
3. Catch data (biomass or numbers)
4. Discard totals or rate data
5. Mean body weight or mean body length data
6. Length composition set-up
7. Length composition data
8. Age composition set-up
9. Age imprecision definitions
10. Age composition data
11. Mean length-at-age or mean bodyweight-at-age data
12. Generalized size composition (e.g., weight frequency) data
13. Environmental data
14. Tag-recapture data
15. Stock composition (e.g., morphs identified by otolith microchemistry) data
16. Selectivity observations (new placeholder, not yet implemented)

Units of Measure

The normal units of measure are as follows:

- Catch biomass - metric tons

- Body weight - kilograms
- Body length - usually in centimeters, weight-at-length parameters must correspond to the units of body length and body weight
- Survey abundance - any units if catchability (Q) is freely scaled; metric tons or thousands of fish if Q has a quantitative interpretation
- Output biomass - metric tons
- Numbers - thousands of fish, because catch is in metric tons and body weight is in kilograms
- Spawning biomass - metric tons of mature females if eggs/kg = 1 for all weights; otherwise has units that are based on the user-specified fecundity

Time Units

- Spawning:
 - Restricted to happening once per year at a specified date (in real months). To create multiple spawning events per year, the way to do this is to change the definition of a year (e.g. call a 6 month period a "year" when there are two spawning events about 6 months apart) which you can do by having the total duration of the seasons be less than 1 year.
- Recruitment:
 - Occurs at specified recruitment events that occur at user-specified dates (in real months).
 - There can be one to many recruitment events across a year; each producing a platoon as a portion of the total recruitment.
 - A settlement platoon enters the model at age 0 if settlement is between the time of spawning and the end of the year; it enters at age 1 if settlement is after the first of the year; these ages at settlement can be overridden in the settlement setup.
- Timing
 - All fish advance to the next older integer age on January 1, no matter when they were born during the year. Consult with your ageing lab to assure consistent interpretation.
- Parameters:
 - Time-varying parameters are allowed to change annually, not seasonally.

- Rates like growth and mortality are per year.

Seasons

Seasonal quantities in the model are calculated and treated in the following methods:

- Seasons are the time step during which constant rates apply.
- Catch and discard amounts are per season and F is calculated per season.
- The year can have just one annual season, or be subdivided into seasons of unequal length.
- Season duration is input in real months and is converted into fractions of an annum. Annual rate values are multiplied by the per annum season duration.
- If the sum of the input season durations do not equal 12.0, then the input durations are divided by the total duration input to rescale season duration to equal 1.0.
- Allows for a special situation in which the year could be only 0.25 in duration (e.g., seasons as years) so that spawning and time-varying parameters can occur more frequently.

Subseasons and Timing of Events

The treatment of subseasons in SS v.3.30 provide more precision in the timing of events compared to earlier model versions. SS v.3.24 and all earlier versions effectively had two subseasons per season because the age-length-key (ALK) for each observation used the mid-season mean length-at-age and spawning occurred at the beginning of a specified season.

Time steps in SS v.3.30 now can be broken into subseason and the ALK can be calculated multiple times over the course of a year:

ALK	ALK*	ALK*	ALK	ALK*	ALK
Subseason 1	Subseason 2	Subseason 3	Subseason 4	Subseason 5	Subseason 6

ALK* only re-calculated when there is a survey that subseason

- Even number (min = 2) of subseasons per season (regardless of season duration):
 - Two subseasons will mimic SS v.3.24

- Specifying more subseasons will give finer temporal resolution, but will slow the model down, the effect of which is mitigated by only calculating growth as needed.
- Survey timing is now cruise-specific and specified in units of months (e.g., April 15 = 4.5).
 - sstrans.exe will convert year, season in SS v.3.24 format to year, real month in SS v.3.30 format.
- Survey integer season and spawn integer season assigned at run time based on real month and season duration(s).
- Growth and the age-length-key (ALK) is calculated at the beginning and mid-season or when there is a defined subseason with data observation.
- Fishery body weight uses mid-subseason growth.
- Survey body weight and size composition is calculated using the nearest subseason.
- Reproductive output now has specified spawn timing (in months fraction) and interpolates growth to that timing.
- Survey numbers calculated at cruise survey timing using e^{-z} .
- Continuous Z for entire season. Same as applied in version SS v.3.24.

Terminology

The term COND appears in the "Typical Value" column of this documentation (it does not actually appear in the model files), it indicates that the following section is omitted except under certain conditions, or that the factors included in the following section depend upon certain conditions. In most cases, the description in the definition column is the same as the label output to the ss_new files.

Model Dimensions

Value	Description
#V3.30.XX.XX	Model version number. This is written by SS in the new files and a good idea to keep updated in the input files.

#C data using new survey	Data file comment. Must start with #C to be retained then written to top of various output files. These comments can occur anywhere in the data file, but must have #C in columns 1-2.
1971	Start year
2001	End year
1	Number of seasons per year
12	Vector with the number of months in each season. These do not need to be integers. Note: If the sum of this vector is close to 12.0, then it is rescaled to sum to 1.0 so that season duration is a fraction of a year. If the sum is not equal to 12.0, then the entered values are summed and rescaled to 1. So, with one season per year and 3 months per season, the calculated season duration will be 0.25, which allows a quarterly model to be run as if quarters are years. All rates in SS are calculated by season (growth, mortality, etc.) using annual rates and season duration.
2	The number of subseasons. Entry must be even and the minimum value is 2. This is for the purpose of finer temporal granularity in calculating growth and the associated age-length key.
1.5	Spawning month; spawning biomass is calculated at this time of year (1.5 means January 15) and used as basis for the total recruitment of all settlement events resulting from this spawning.
2	Number of sexes: 1 = current one gender, ignore fraction female input in the control file; 2 = current two gender, use fraction female in the control file; and -1 = one gender and multiply the spawning biomass by the fraction female in the control file.
20	Number of ages. The value here will be the plus-group age. SS starts at age 0.
1	Number of areas
2	Total number of fishing and survey fleets (which now can be in any order).

Fleet Definitions

The catch data input has been modified to improve the user flexibility to add/subtract fishing and survey fleets to a model set-up. The fleet setup input is transposed so each fleet is now a row. Previous versions (SS v.3.24 and earlier) required that fishing fleets be listed first followed by survey only fleets. In version SS v.3.30 all fleets now have the same status within the model structure and each has a specified fleet type (except for models that use tag recapture data, this will be corrected in future versions). Available types are; catch fleet, bycatch only fleet, or survey.

Inputs that define the fishing and survey fleets:

2 Number of fleets which includes survey in any order					
Fleet Type	Timing	Area	Catch Units	Catch Mult.	Fleet Name
1	-1	1	1	0	FISHERY1
3	1	1	2	0	SURVEY1

Fleet Type

Define the fleet type (e.g., fishery fleet, survey fleet):

- 1 = fleet with input catches;
- 2 = bycatch fleet (all catch discarded) and invoke extra input for treatment in equilibrium and forecast;
- 3 = survey: assumes no catch removals even if associated catches are specified below. If you would like to remove survey catch set fleet type to option = 1 with specific month timing for removals (defined below in Timing); and
- 4 = ignored (not yet implemented).

Timing

Timing for data observations has been revised in SS v.3.30:

- Fishery = -1 catch is treated as if it occurred over the whole season. A user can override this assumption for specific data observations (e.g., length or age) by specifying a month.
- Survey = 1 The fleet timing here for surveys is not used and only the month value associated with each observation is relevant (e.g., month specification in the indices of abundance or the month for composition data).

Area

An integer value indicating the area in which a fleet operates.

Catch Units

Ignored for survey fleets, their units are read later:

- 1 = biomass (in metric tons); and
- 2 = numbers (thousands of fish).

See Units of Measure for more information.

Catch Multiplier

Invokes use of a catch multiplier, which is then entered as a parameter in the mortality-growth parameter section. The estimated value or fixed value of the catch multiplier is used to adjust the observed catch:

- 0 = No catch multiplier used; and
- 1 = Apply a catch multiplier which is defined as an estimable parameter in the control file after the cohort growth deviation in the biology parameter section. The model's estimated retained catch will be multiplied by this factor before being compared to the observed retained catch.

A catch multiplier can be useful when trying to explore historical unrecorded catches or ongoing illegal and unregulated catches. The catch multiplier is a full parameter line in the control file and has the ability to be time-varying.

Bycatch Fleets

The option to include bycatch fleets was introduced in SS v.3.30.10. This is an optional input and if no bycatch is to be included in to the catches this section can be ignored.

A fishing fleet is designated as a bycatch fleet by indicating that its fleet type is 2. A bycatch fleet creates a fishing mortality, same as a fleet of type 1, but a bycatch fleet has all catch discarded so the input value for retained catch is ignored. However, an input value for retained catch is still needed to indicate that the bycatch fleet was active in that year and season. A catch multiplier cannot be used with bycatch fleets because catch multiplier works on retained catch. SS will expect that the retention function for this fleet will be set in the selectivity section to type 3, indicating that all selected catch is discarded dead. It is necessary to specify a selectivity pattern for the bycatch fleet and, due to generally lack of data, to externally derive values for the parameters of this selectivity.

All catch from a bycatch fleet is discarded, so one option to use a discard fleet is to enter annual values for the amount (not proportion) that is discarded in each time step. However, it is uncommon to have such data for all years. An alternative approach that has been used principally in the U.S. Gulf of Mexico is to input a time series of effort data for this fleet in the survey section (e.g., effort is a "survey" of F, for example, the shrimp trawl fleet in the Gulf of Mexico catches and discards small finfish and an effort time series is available for this fleet) and to input in the discard data section an observation for the average discard over time using the super year approach. Another use of bycatch fleet is to use it to estimate effect of an external source of mortality, such as a red tide event. In this usage there may be no data on the magnitude of the discards and SS will then rely solely on the contrast in other data to attempt to estimate the magnitude of the red tide kill that occurred. The benefit of doing this as a bycatch fleet, and not a block on natural mortality, is that the selectivity of the effect can be specified.

Bycatch fleets are not expected to be under the same type of fishery management controls as the retained catch fleets included in the model. This means that when SS enters into the reference point equilibrium calculations, it would be incorrect to have SS re-scale the magnitude of the F for the bycatch fleet as it searches for the F that produces, for example, F35%. Related issues apply to the forecast. Consequently, a separate set of controls is provided for bycatch fleets (defined below). Input is required for each fleet designated as fleet type = 2.

If a fleet above was set as a bycatch fleet (fleet type = 2), the following line is required:

Bycatch fleet input controls:

a: Fleet Index	b: Include in MSY	c: Fmult	d: F or First Year	e: Last Year	f: Not used
2	2	3	1982	2010	0

The above example set-up defines one fleet (fleet number 2) as a bycatch fleet with the dead catch from this fleet to not be included in the search for MSY (b: Include in MSY = 2). The level of F from the bycatch fleet in reference point and forecast is set to the mean (c: Fmult = 3) of the estimated F for the range of years from 1982-2010.

Fleet Index

Fleet number for which to include bycatch catch. Fleet number is assigned within the model based on the order of listed fleets in the Fleet Definition section. If there are multiple bycatch fleets, then a line for each fleet is required in the bycatch section.

Include in MSY

The options are:

- 1 = deadfish in MSY, ABC, and other benchmark and forecast output; and
- 2 = omit from MSY and ABC (but still include the mortality).

Fmult

The options are:

- 1 = F multiplier scales with other fleets;
- 2 = bycatch F constant at input value in column d; and
- 3 = bycatch F from range of years input in columns d and e.

F or First Year

The specified F or first year for the bycatch fleet.

F or Last Year

The specified F or last year for the bycatch fleet.

Not Used

This column is not yet used and is reserved for future features.

Bycatch Fleet Usage Instructions and Warnings

When implementing a bycatch fleet, changes to both the data and control file are needed.

The needed changes to the data file are:

1. Fleet type – set to value of 2.
2. Set bycatch fleet controls per information above.
3. Catch input – you must enter a positive value for catch in each year/season that you want a bycatch calculated. The entered value of catch will be ignored by SS, it is just a placeholder to invoke creating an F.
 - (a) Initial equilibrium – you may want to enter the bycatch amount as retained catch for the initial equilibrium year because there is no option to enter initial equilibrium discard in the discard section.

4. Discard input – It is useful, but not absolutely necessary, to enter the amount of discard to assist SS in estimating the F for the bycatch fleet.
5. Survey input - It is useful, but not absolutely necessary, to enter the effort time series by the bycatch fleet to assist SS in estimating the annual changes in F for the bycatch fleet.

The needed changes to the control file are:

1. The F method must be set to 2 in order for SS to estimate F with having information on retained catch.
2. Selectivity –
 - (a) A selectivity pattern must be specified and fixed (or estimated if composition data is provided).
 - (b) The discard column of selectivity input must be set to a value of 3 to cause all catch to be discarded.

In version 3.30.14 it was identified that there can be an interaction between the use of bycatch fleets and the search for the $F_{0.1}$ reference point which may results in the search failing. Changes to the search feature were implemented to make the search more robust, however, issue may still be encountered. In these instances it is recommended to not select the $F_{0.1}$ reference point calculation in the forecast file.

Catch

After reading the fleet-specific indicators, a list of catch values by fleet and season are read in by the model. The format for the catches is year and season that the catch will be attributed to, fleet, a catch value, and a year specific catch standard error. Only positive catches need to be entered, so there is no need for records corresponding to all years and fleets. To include an equilibrium catch value the year should be noted as -999 and this is now fleet and season specific.

There is no longer a need to specify the number of records to be read; instead the list is terminated by entering a record with the value of -9999 in the year field. The updated list based approach extends throughout the data file (e.g., catch, length- and age-composition data), the control file (e.g., lambdas), and the forecast file (e.g., total catch by fleet, total catch by area, allocation groups, forecasted catch).

In addition, it is possible to collapse the number of seasons. So, if a season value is greater than the number of seasons for a particular model, that catch is added to the catch for the final season. This is one way to easily collapse a seasonal model into an annual model. The alternative option is to the use

of season = 0. This will cause SS to distribute the input value of catch equally among the number of seasons. SS assumes that catch occurs continuously over seasons and hence is not specified as month in the catch data section. However, all other data types will need to be specified by month.

The new format for version SS v.3.30 for a 2 season model with 2 fisheries looks like the table below. The example is sorted by fleet, but the sort order does not matter. In data.ss_new, the sort order is fleet, year, season.

Catches by year, season for every fleet:

Year	Season	Fleet	Catch	Catch SE
-999	1	1	56	0.05
-999	2	1	62	0.05
1975	1	1	876	0.05
1975	2	1	343	0.05
...
-999	1	2	55	0.05
-999	2	2	22	0.05
1975	1	2	555	0.05
1975	2	2	873	0.05
...
-9999	0	0	0	0

- Catch can be in terms of biomass or numbers for each fleet, but cannot be mixed within a fleet.
- Catch is retained catch (aka landings). If there is discard also, then it is handled in the discard section below. This is the recommended setup which results in a model estimated retention curve based upon the discard data (specifically discard composition data). However, there may be instances where the data do not support estimation of retention curves. In these instances catches can be specified as all dead (retained + discard estimates).
- If there are challenges to estimating discards within the model, catches can be input as total dead without the use of discard data and retention curves.
- If there is reason to believe that the retained catch values underestimate the true catch, then it is possible in the retention parameter set up to create the ability for the model to estimate the degree of unrecorded catch. However, this is better handled with the new catch multiplier option.

Indices

Indices are data that are compared to aggregate quantities in the model. Typically the index is a measure of selected fish abundance, but this data section also allows for the index to be related to a fishing fleet's F , or to another quantity estimated by the model. The first section of the "Indices" setup contains the fleet number, units, error distribution, and whether additional output (SD Report) will be written to the Report file for each fleet that has index data.

Catch-per-unit-effort (CPUE) and Survey Abundance Observations:

Fleet/ Survey	Units	Error Distribution	SD Report
1	1	0	0
2	1	0	0
...

Units

The options for units for input data are:

- 0 = numbers;
- 1 = biomass;
- 2 = F ; and
 - Note the F option can only be used for a fishing fleet and not for a survey, even if the survey selectivity is mirrored to a fishing fleet. The values of these effort data are interpreted as proportional to the level of the fishery F values. No adjustment is made for differentiating between continuous F values versus exploitation rate values coming from Pope's approximation. A normal error structure is recommended so that the input effort data are compared directly to the model's calculated F , rather than to $\log(F)$. The resultant proportionality constant has units of $1/Q$ where Q is the catchability coefficient.
- ≥ 30 special survey types. These options bypass the calculation of survey selectivity so no selectivity parameters should be entered. The expected values for these types are:
 - 30 = spawning biomass/output (e.g., for an egg and larvae survey);
 - 31 = $\exp(\text{recruitment deviation})$, useful for environmental index affecting recruitment;
 - 32 = spawning biomass * $\exp(\text{recruitment deviation})$, for a pre-recruit survey occurring before density-dependence;

- 33 = recruitment, age-0 recruits;
- 34 = depletion (spawning biomass/virgin spawning biomass);
 - * Special survey option 34 automatically adjusts phases of parameters. To use the depletion survey approach, the user will need to make the following revisions to the SS data file: 1) add a new survey fleet, 2) define the survey type as option 34, 3) add two depletion survey data points, and initial unfished set equal to 1 for an unfished modeled year and one for a later year with the depletion estimates, 4) set the input CV value for each survey data point to a low value (e.g., 0.0001) to force the model to fit these data, and in the control file 5) add the survey to the control file in the Q set-up and selectivity sections with float set to 0 with parameter value set to 0.
 - * There are options for additional control over this in the control file Q setup section under the "link information" column where:
 - 0 = add 1 to phases of all parameters. Only R0 active in new phase 1. Mimics the default option of previous model versions;
 - 1 = only R0 active in phase 1. Then finish with no other parameters becoming active; useful for data-limited draws of other fixed parameters. Essentially, this option allows SS to mimic DB-SRA; and
 - 2 = no phase adjustments, can be used when profiling on fixed R0.
 - * Warning: the depletion survey approach has not been tested on multiple area models. This approach may present challenges depending upon the dynamics within each area.
- 35 = survey of a deviation vector ($e(survey(y)) = f(parm_dev(k, y))$), can be used for an environmental time-series with soft linkage to the index. The selected deviation vector is specified in Q section of the control file. The index of the deviation vector to which the index is related is specified in the 2nd column of the Q setup table (see Catchability).

Error Distribution

The options for error distribution form are:

- -1 = normal error;
- 0 = lognormal error; and
- >0 = Student's t-distribution in log space with degrees of freedom equal to this value. For DF>30, results will be nearly identical to that for lognormal distribution. A DF value of about 4

gives a fat-tail to the distribution. The standard error values entered in the data file must be the standard error in \log_e space.

Abundance indices typically have a lognormal error structure with units of standard error of $\log_e(\text{index})$. If the variance of the observations is available only as a coefficient of variation (CV), then the value of standard error can be approximated as $\sqrt{(\log_e(1 + (CV)^2))}$ where the CV is the standard error of the observation divided by the mean value of the observation.

For the normal error structure, the entered values for standard error are interpreted directly as a standard error in arithmetic space and not as a CV. Thus switching from a lognormal to a normal error structure forces the user to provide different values for the standard error input in the data file.

If the data exist as a set of normalized Z-scores, you can assert a lognormal error structure after entering the data as $\exp(\text{Z-score})$ because it will be logged by SS. Preferably, the Z-scores would be entered directly and the normal error structure would be used.

Enable SD Report

Indices with SD Report enabled will have the expected values for their historical values appear in the ss.std and ss.cor files. The default value for this option is 0.

- 0 = SD Report not enabled for this index; and
- 1 = SD Report enabled for this index.

Data Format

Year	Month	Fleet/Survey	Observation	SE
1991	7	3	80000	0.056
1995	7.2	3	65000	0.056
...
2000	7.1	3	42000	0.056
-9999	0	0	0	0

- For fishing fleets, CPUE is defined in terms of retained catch (biomass or numbers).
- For fishery independent surveys, retention/discard is not defined so CPUE is implicitly in terms of total CPUE.

- If a survey has its selectivity mirrored to that of a fishery, only the selectivity is mirrored so the expected CPUE for this mirrored survey does not use the retention curve (if any) for the fishing fleet.
- If the fishery or survey has time-varying selectivity, then this changing selectivity will be taken into account when calculating expected values for the CPUE or survey index.
- Year values that are before start year or after end year are excluded from model, so the easiest way to include provisional data in a data file is to put a negative sign on its year value.
- Duplicate survey observations for the same year are not allowed.
- Observations that are to be included in the model but not included in the negative log likelihood need to have a negative sign on their fleet ID. Previously the code for not using observations was to enter the observation itself as a negative value. However, that old approach prevented use of a Z-score environmental index as a “survey”. This approach is best for single or select years from an index rather than an approach to remove a whole index. Removing a whole index from the model should be done through the use of lambdas at the bottom of the control file which will eliminate the index from model fitting.
- Observations can be entered in any order, except if the super-year feature is used.
- Super-periods are turned on and then turned back off again by putting a negative sign on the season. Previously, super-periods were started and stopped by entering -9999 and the -9998 in the SE field. See the Data Super-Period section of this manual for more information.
- If the statistical analysis used to create the CPUE index of a fishery has been conducted in such a way that its inherent size/age selectivity differs from the size/age selectivity estimated from the fishery’s size and age composition, then you may want to enter the CPUE as if it was a separate survey and with a selectivity that differs from the fishery’s estimated selectivity. The need for this split arises because the fishery size and age composition should be derived through a catch-weighted approach (to appropriately represent the removals by the fishery) and the CPUE should be derived through an area-weighted approach to better serve as a survey of stock abundance.

Discard

If discard is not a feature of the model specification, then just a single input is needed:

0	Number of fleets with discard observations
---	--

If discard is being used, the input syntax is:

1	Number of fleets with discard observations			
Fleet	Units	Error Distribution		
1	2	-1		
Year	Month	Fleet	Observation	Standard Error
1980	7	1	0.05	0.25
1991	7	1	0.10	0.25
-9999	0	0	0	0

Note that although the user must specify a month for the observed discard data, the unit for discard data is in terms of a season rather than a specific month. So, if using a seasonal model, the input month values must corresponding to some time during the correct season. The actual value will not matter because the discard amount is calculated for the entirety of the season. However, discard length or age observations will be treated by entered observation month.

Discard Units

The options are:

- 1 = values are amount of discard in either biomass or numbers according to the selection made for retained catch;
- 2 = values are fraction (in biomass or numbers) of total catch discarded, biomass/number selection matches that of retained catch; and
- 3 = values are in numbers (thousands) of fish discarded, even if retained catch has units of biomass.

Discard Error Distribution

The four options for discard error are:

- >0 = degrees of freedom for Student's t-distribution used to scale mean body weight deviations. Value of error in data file is interpreted as CV of the observation;
- 0 = normal distribution, value of error in data file is interpreted as CV of the observation;
- -1 = normal distribution, value of error in data file is interpreted as standard error of the observation;

- -2 = lognormal distribution, value of error in data file is interpreted as standard error of the observation in log space; and
- -3 = truncated normal distribution (new with SS v.3.30, needs further testing), value of error in data file is interpreted as standard error of the observation. This is a good option for low observed discard rates.

Discard Notes

- Year values that are before start year or after end year are excluded from model, so the easiest way to include provisional data in a data file is to put a negative sign on its year value.
- Negative value for fleet causes it to be included in the calculation of expected values, but excluded from the log likelihood.
- Zero (0.0) is a legitimate discard observation, unless lognormal error structure is used.
- Duplicate discard observations from a fleet for the same year are not allowed.
- Observations can be entered in any order, except if the super-period feature is used.
- Note that in the control file you will enter information for retention such that 1-retention is the amount discarded. All discard is assumed dead, unless you enter information for discard mortality. Retention and discard mortality can be either size-based or age-based (new with SS v.3.30).

Cautionary Note

The use of CV as the measure of variance can cause a small discard value to appear to be overly precise, even with the minimum standard error of the discard observation set to 0.001. In the control file, there is an option to add an extra amount of variance. This amount is added to the standard error, not to the CV, to help correct this problem of underestimated variance.

Mean Body Weight or Length

This is the overall mean body weight or length across all selected sizes and ages. This may be useful in situations where individual fish are not measured but mean weight is obtained by counting the number of fish in a specified sample, e.g., a 25 kg basket.

Mean Body Weight Data Section:

1	Use mean body size data (0/1)					
COND > 0:						
30	Degrees of freedom for Student's t-distribution used to evaluate mean body weight deviation.					
Year	Month	Fleet	Partition	Type	Observation	CV
1990	7	1	0	1	4.0	0.95
1990	7	1	0	1	1.0	0.95
-9999	0	0	0	0	0	0

Partition

Mean weight data and composition data require specification of what group the sample originated from (e.g., discard, retained, discard + retained).

- 0 = whole catch in units of weight (discard + retained);
- 1 = discarded catch in units of weight; and
- 2 = retained catch in units of weight.

Type

Specify the type of data:

- 1 = mean length; and
- 2 = mean body weight.

Observation - Units

Units must correspond to the units of body weight, normally in kilograms, (or mean length in cm). The expected value of mean body weight (or mean length) is calculated in a way that incorporates effect of selectivity and retention.

Error

Error is entered as the CV of the observed mean body weight (or mean length)

Population Length Bins

The first part of the length composition section sets up the bin structure for the population. These bins define the granularity of the age-length key and the coarseness of the length selectivity. Fine bins create smoother distributions, but a larger and slower running model. First read a single value to select one of three population length bin methods, then any conditional input for options 2 and 3:

1	Use data bins to be read later. No additional input here.	
2	generate from bin width min max, read next:	
	2	Bin width
	10	Lower size of first bin
	82	Lower size of largest bin
The number of bins is then calculated from: $(\text{max Lread} - \text{min Lread})/(\text{bin width}) + 1$		
3	Read 1 value for number of bins, and then read vector of bin boundaries	
	37	Number of population length bins to be read
	10 12 14 ... 82	Vector containing lower edge of each population size bin

Notes

There are some items for users to consider when setting up population length bins:

- For option 2, bin width should be a factor of min size and max size. For options 2 and 3, the data length bins must not be wider than the population length bins and the boundaries of the bins do not have to align. The transition matrix between population and data length bins is output to echoinput.sso.
- The mean size at settlement (virtual recruitment age) is set equal to the min size of the first population length bin.
- When using more, finer population length bins, SS will create smoother length selectivity curves and smoother length distributions in the age-length key, but run more slowly (more calculations to do).
- The mean weight-at-length, maturity-at-length and size-selectivity are based on the mid-length of the population bins. So these quantities will be rougher approximations if broad bins are defined.

- Provide a wide enough range of population size bins so that the mean body weight-at-age will be calculated correctly for the youngest and oldest fish. If the growth curve extends beyond the largest size bin, then these fish will be assigned a length equal to the mid-bin size for the purpose of calculating their body weight.
- While exploring the performance of models with finer bin structure, a potentially pathological situation has been identified. When the bin structure is coarse (note that some applications have used 10 cm bin widths for the largest fish), it is possible for a selectivity slope parameter or a retention parameter to become so steep that all of the action occurs within the range of a single size bin. In this case, the model will see zero gradient of the log likelihood with respect to that parameter and convergence will be hampered.
- A value read near the end of the starter.ss file defines the degree of tail compression used for the age-length key, called ALK tolerance. If this is set to 0.0, then no compression is used and all cells of the age-length key are processed, even though they may contain trivial (e.g., 1×10^{-13}) fraction of the fish at a given age. With tail compression of, say 0.0001, SS will at the beginning of each phase calculate the min and max length bin to process for each age of each morphs ALK and compress accordingly. Depending on how many extra bins are outside this range, you may see speed increases near 10-20%. Large values of ALK tolerance, say 0.1, will create a sharp end to each distribution and likely will impede convergence. It is recommended to start with a value of 0 and if model speed is an issue, explore values greater than 0 and evaluate the trade-off between model estimates and run time. The user is encouraged to explore this feature.

Length Composition Data Structure

Enter a code to indicate whether or not length composition data will be used:

1	Use length composition data (0/1)
---	-----------------------------------

If the value 0 is entered, then skip all length related inputs below and skip to the age data setup section. Otherwise continue:

Specify bin compression and error structure for length composition data for each fleet:

Min. Tail Compress.	Constant added to prop.	Combine males & females	Compress. Bins	Comp. Error Dist.	Dirichlet Param. Select	Min. Sample Size
0	0.0001	0	0	0	0	1
0	0.0001	0	0	0	0	1

Minimum Tail Compression

Compress tails of composition until observed proportion is greater than this value; negative value causes no compression; Advise using no compression if data are very sparse, and especially if the set-up is using age composition within length bins because of the sparseness of these data.

Added Constant to Proportions

Constant added to observed and expected proportions at length and age to make logL calculations more robust. Tail compression occurs before adding this constant. Proportions are renormalized to sum to 1.0 after constant is added.

Combine Males & Females

Combine males into females at or below this bin number. This is useful if the sex determination of very small fish is doubtful so allows the small fish to be treated as combined sex. If Combine Males & Females > 0, then add males into females for bins 1 through this number, zero out the males, set male data to start at the first bin above this bin. Note that Combine Males & Females > 0 is entered as a bin index, not as the size associated with that bin. Comparable option is available for age composition data.

Compress Bins

This option allows for the compression of length or age bins beyond a specific length or age by each data source. As an example, a value of 5 in the compress bins column would condense the final five length bins for the specified data source.

Composition Error Distribution

The options are:

- 0 = Multinomial Error; and
- 1 = Dirichlet Multinomial Error.
 - The Dirichlet Multinomial Error distribution requires the addition of a parameter lines for the natural log of the effective sample size multiplier (θ) at the end of the selectivity parameter section in the control file. See the Dirichlet parameter in the control file for information regarding setup.
 - The Parameter Select option needs be used to specify which data sources should be weighted together or separate.

Dirichlet Parameter Select

Value that indicates the groups of composition data for estimation of the Dirichlet parameter for weighting composition data.

- 0 = Default; and
- 1-N = Only used for the Dirichlet option. Set to a sequence of numbers from 1 to N where N is the total number of combinations of fleet and age/length. That is, if you have 3 fleets with length data, but only 2 also have age data, you would have values 1 to 3 in the length comp setup and 4 to 5 in the age comp setup. You can also have a data weight that is shared across fleets by repeating values in Parameter Select.

Minimum Sample Size

The minimum value (floor) for all sample sizes. This value must be at least 0.001. Conditional age-at-length data may have observations with sample sizes less than 1. SS v.3.24 had an implicit minimum sample size value of 1.

Dirichlet Parameter Number and Effective Sample Sizes

If the Dirichlet multinomial error distribution is selected, indicate here which of a list of Dirichlet multinomial parameters will be used for this fleet. So each fleet could use a unique Dirichlet multinomial parameter, or all could share the same, or any combination of unique and shared. The requested number of Dirichlet multinomial parameters are specified as parameter lines in the control file immediately after the selectivity parameter section. Please note that age-compositions Dirichlet multinomial parameters are continued after length-compositions, so a model with one fleet and both data types would presumably require two new Dirichlet multinomial parameters.

The Dirichlet estimates the effective sample size as $N_{eff} = \frac{1}{1+\theta} + \frac{N\theta}{1+\theta}$ where θ is the estimated parameter and N is the input sample size. Stock Synthesis estimates the log of the Dirichlet multinomial parameter such that $\hat{\theta}_{fishery} = e^{-0.6072} = 0.54$ where assuming $N = 100$ for the fishery would result in an effective sample size equal to 35.7.

This formula for effective sample size implies that, as the Stock Synthesis parameter $\ln(\text{DM_theta})$ goes to large values (i.e., 20), then the adjusted sample size will converge to the input sample size. In this case, small changes in the value of the $\ln(\text{DM_theta})$ parameter has no action, and the derivative of the negative log-likelihood is zero with respect to the parameter, which means the Hessian will be singular and cannot be inverted. To avoid this non-invertible Hessian when the $\ln(\text{DM_theta})$ parameter becomes large, turn it off while fixing it at the high value. This is equivalent to turning off down-weighting of fleets where evidence suggests that the input sample sizes are reasonable.

For additional information about the Dirichlet multinomial please see Thorson et al. (2017) and the detailed Data Weighting section.

Length Composition Data

Composition data within SS can be entered as proportions, numbers, or values of observations by length bin based on data expansions.

The data bins do not need to cover all observed lengths. The selection of data bin structure should be based on the observed distribution of lengths and the assumed growth curve. If growth asymptotes at larger lengths, having additional length bins across these sizes may not contribute information to the model and may slow model run time. Additionally, the lower length bin selection should be selected such that, depending on the size selection, to allow for information on smaller fish and possible patterns in recruitment. While set separately users should ensure that the length and age bins align. It is recommended to explore multiple configurations of length and age bins to determine the impact of this choice on model estimation.

Specify the length composition data as:

28	Number of length bins for data					
<hr/>						
26 28 30 ... 80	Vector of length bins associated with the length data.					
<hr/>						
The above data bins will combine 0-26 cm fish, >26 - 28 cm, ..., and 80+ cm fish.						
<hr/>						
Example of a single length composition observation:						
<hr/>						
Year	Month	Fleet	Sex	Partition	Nsamp	data vector
<hr/>						
1986	1	1	3	0	20	<female then male data>
...
-9999	0	0	0	0	0	0

Sex

If model has only one sex defined in the set-up, all observations must have sex set equal to 0 or 1 and the data vector by year will equal the number of the user defined data bins. This also applies to the age data.

In a 2 sex model, the data vector always has female data followed by male data, even if only one of the two sexes has data that will be used. The below description applies to a 2 sex model:

- Sex = 0 means combined male and female (must already be combined and information placed in the female portion of the data vector) (male entries must exist for correct data reading, then will be ignored).
- Sex = 1 means female only (male entries must exist for correct data reading, then will be ignored).
- Sex = 2 means male only (female entries must exist and will be ignored after being read).
- Sex = 3 means data from both sexes will be used and they are scaled so that they together sum to 1.0; i.e. sex ratio is preserved.

Partition

Partition indicates samples from either discards, retained, or combined.

- 0 = combined;
- 1 = discard; and
- 2 = retained.

Excluding Data

- If the value of year is negative, then that observation is not transferred into the working array. This feature is the easiest way to include observations in a data file but not to use them in a particular model scenario.
- If the value of fleet in the length or age composition observed data line is negative, then the observation is processed and its expected value and log likelihood are calculated, but this log likelihood is not included in the total log likelihood. This feature allows the user to see the fit to a provisional observation without having that observation affect the model.

Note

When processing data to be input into SS, all observed fish of sizes smaller than the first bin should be added to the first bin and all observed fish larger than the last bin should be condensed into the last bin.

SS no longer requires that the number of length composition data lines be specified. Entering -9999 at the end of the data matrix will indicate to the model the end of length composition lines to be read.

Each observation can be stored as one row for ease of data management in a spreadsheet and for sorting of the observations. However, the 6 header values, the female vector and the male vector could each be on a separate line because ADMB reads values consecutively from the input file and will move to the next line as necessary to read additional values.

The composition observations can be in any order and replicate observations by a year for a fleet are allowed (unlike survey and discard data). However, if the super-period approach is used, then each super-periods' observations must be contiguous in the data file.

Age Composition Option

The age composition section begins by reading the number of age bins. If the value 0 is entered for the number of age bins, then SS skips reading the bin structure and all reading of other age composition data inputs.

17	Number of age bins; can be equal to 0 if age data are not used; do not include a vector of agebins if the number of age bins is set equal to 0.
----	---

Age Composition Bins

If a positive number of age bins is read, then SS reads the bin definition next.

1 2 3 ... 20 25	Vector of ages
-----------------	----------------

The bins are in terms of observed age (here age) and entered as the lower edge of each bin. Each ageing imprecision definition is used to create a matrix that translates true age structure into age structure. The first and last age' bins work as accumulators. So in the example any age 0 fish that are caught would be assigned to the age = 1 bin.

Ageing Error

Here, the capability to create a distribution of age (e.g., age with possible bias and imprecision) from true age is created. One or many age error definitions can be created. For each, there is input of a vector of mean age and standard deviation of age. For one definition, the input vectors can be replaced by vectors created from estimable parameters. In the future, capability to read a full age – age matrix could be created. The dimension of the ageing error matrix requires the column length

match the population maximum age specified at the top of the data file. However, the maximum age for binning of age data may be lower than the population maximum age.

2	Number of ageing error matrices to generate				
Example with no bias and very little uncertainty at age:					
Age-0	Age-1	Age-2	...	Max Age	
-1	-1	-1	...	-1	#Mean Age
0.001	0.001	0.001	...	0.001	#SD
Example with no bias and some uncertainty at age:					
0.5	1.5	2.5	...	Max Age + 0.5	#Mean Age
0.5	0.65	0.67	...	4.3	#SD Age
Example with bias and uncertainty at age:					
0.5	1.4	2.3	...	Max Age + Age Bias	#Mean Age
0.5	0.65	0.67	...	4.3	#SD Age

In principle, one could have year or laboratory specific matrices for ageing error. For each matrix, enter a vector with mean age for each true age; if there is no ageing bias, then set age equal to true age + 0.5. Alternatively, -1 value for mean age means to set it equal to true age plus 0.5. The addition of +0.5 is needed so that fish will get assigned to the intended integer age. The length of the input vector is equal to the population maximum age plus one (0-max age), with the first entry being for age 0 fish and the last for fish of population maximum age. The following line is a vector with the standard deviation of age for each true age with a normal distribution assumption.

SS is able to create one ageing error matrix from parameters, rather than from an input vector. The range of conditions in which this new feature will perform well has not been evaluated, so it should be considered as a preliminary implementation and subject to modification. To invoke this option, for the selected ageing error vector, set the standard deviation of ageing error to a negative value for age 0. This will cause creation of an ageing error matrix from parameters and any age or size-at-age data that specify use of this age error pattern will use this matrix. Then in the control file, add a full parameter line below the cohort growth deviation parameter (or the movement parameter lines if used) in the mortality growth parameter section. These parameters are described in the control file section of this manual.

Age Composition Specification

If age data are included in the model, the following set-up is required, similar to the length data section.

Specify bin compression and error structure for age composition data for each fleet:

Min. Tail Compress.	Constant added to prop.	Combine males & females	Compress. Bins	Comp. Error Dist.	Dirichlet Param. Select	Min. Sample Size
0	0.0001	1	0	0	0	1
0	0.0001	1	0	0	0	1

Specify method by which length bin range for age obs will be interpreted:

- 1 Bin method for age data
 - 1 = value refers to population bin index
 - 2 = value refers to data bin index
 - 3 = value is actual length (which must correspond to population length bin boundary)

An example age composition observation:

Year	Month	Fleet	Sex	Partition	Age Err	Lbin lo	Lbin hi	Nsamp	Data Vector
1987	1	1	3	0	2	-1	-1	79	<enter data values>
-9999	0	0	0	0	0	0	0	0	0

Syntax for Sex, Partition, and data vector are same as for length. The data vector has female values then male values, just as for the length composition data.

Age Error

Age error (Age Err) identifies which ageing error matrix to use to generate expected value for this observation.

Lbin Low and Lbin High

Lbin lo and Lbin hi are the range of length bins that this age composition observation refers to. Normally these are entered with a value of -1 and -1 to select the full size range. Whether these are entered as population bin number, length data bin number, or actual length is controlled by the value of the length bin range method above.

- Entering value of 0 or -1 for Lbin lo converts Lbin lo to 1;
- Entering value of 0 or -1 for Lbin hi converts Lbin hi to Maxbin;
- It is strongly advised to use the -1 codes to select the full size range. If you use explicit values, then the model could unintentionally exclude information from some size range if the population bin structure is changed.
- In reporting to the comp_report.sso, the reported Lbin_lo and Lbin_hi values are always converted to actual length.

Excluding Data

As with the length composition data, a negative year value causes the observation to not be read into the working matrix, a negative value for fleet causes the observation to be included in expected values calculation, but not in contribution to total log likelihood, a negative value for month causes start-stop of super-period.

Conditional Age-at-Length

Use of conditional age-at-length will greatly increase the total number of age composition observations and associated model run time but there can be several advantages to inputting ages in this fashion. First, it avoids double use of fish for both age and size information because the age information is considered conditional on the length information. Second, it contains more detailed information about the relationship between size and age so provides stronger ability to estimate growth parameters, especially the variance of size-at-age. Lastly, where age data are collected in a length-stratified program, the conditional age-at-length approach can directly match the protocols of the sampling program.

However, simulation research has shown that the use of conditional age-at-length data can result in biased growth estimates in the presence of unaccounted for age-based movement when length-based

selectivity is assumed (Lee et al., 2017), when other age-based processes (e.g., mortality) are not accounted for (Lee et al., 2019), or based on the age sampling protocol (Piner et al., 2016). Understanding how data are collected (e.g., random, length-conditioned samples) and the biology of the stock is important when using conditional age-at-length data for a fleet.

In a two sex model, it is best to enter these conditional age-at-length data as single sex observations (sex = 1 for females and = 2 for males), rather than as joint sex observations (sex = 3). Inputting joint sex observations comes with a more rigid assumption about sex ratios within each length bin. Using separate vectors for each sex allows 100% of the expected composition to be fit to 100% observations within each sex, whereas with the sex = 3 option, you would have a bad fit if the sex ratio were out of balance with the model expectation, even if the observed proportion at age within each sex exactly matched the model expectation for that age. Additionally, inputting the conditional age-at-length data as single sex observations isolates the age composition data from any sex selectivity as well.

Conditional age-at-length data are entered within the age composition data section and can be mixed with marginal age observations for other fleets of other years within a fleet. To treat age data as conditional on length, Lbin_lo and Lbin_hi are used to select a subset of the total size range. This is different than setting Lbin_lo and Lbin_hi both to -1 to select the entire size range, which treats the data entered on this line within the age composition data section as marginal age composition data.

An example conditional age-at-length composition observations:

Year	Month	Fleet	Sex	Partition	Age Err	Lbin lo	Lbin hi	Nsamp	Data Vector
1987	1	1	1	0	2	10	10	18	<data values>
1987	1	1	1	0	2	12	12	24	<data values>
1987	1	1	1	0	2	14	14	16	<data values>
1987	1	1	1	0	2	16	16	30	<data values>
-9999	0	0	0	0	0	0	0	0	0

In this example observation, the age data is treated as on being conditional on the 2 cm length bins of 10-11.99, 12-13.99, 14-15.99, and 16-17.99cm. If there are no observations of ages for a specific sex within a length bin for a specific year, that entry may be omitted.

Mean Length or Body Weight-at-Age

SS also accepts input of mean length-at-age or mean body weight-at-age. This is done in terms of observed age, not true age, to take into account the effects of ageing imprecision on expected mean

size-at-age. If the value of the Age Error column is positive, then the observation is interpreted as mean length-at-age. If the value of the Age Error column is negative, then the observation is interpreted as mean body weight-at-age and the abs(Age Error) is used as Age Error.

1 Use mean size-at-age observation (0 = none, 1 = read data matrix)

An example observation:

Yr	Month	Fleet	Sex	Part.	Age Err.	Ignore	Data Vector (Female - Male)	Sample Size (Female - Male)
1989	7	1	3	0	1	999	<Mean Size values>	<Sample Sizes>
...								
-9999	0	0	0	0	0	0	0 0 0 0 0 0	0 0 0 0 0 0

Note

- Negatively valued mean size entries will be ignored in fitting. This feature allows the user to see the fit to a provisional observation without having that observation affect the model.
- A number of fish value of 0 will cause mean size value to be ignored in fitting. This feature allows the user to see the fit to a provisional observation without having that observation affect the model.
- Negative value for year causes observation to not be included in the working matrix. This feature is the easiest way to include observations in a data file but not to use them in a particular model scenario.
- Each sexes' data vector and N fish vector has length equal to the number of age bins.
- The "Ignore" column is not used (set aside for future options) but still needs to have default values in that column (any value).
- Where age data are being entered as conditional age-at-length and growth parameters are being estimated, it may be useful to include a mean length-at-age vector with nil emphasis to provide another view on the model's estimates.
- An experiment that may be of interest might be to take the body weight-at-age data and enter it to SS as empirical body weight-at-true age in the wtatage.ss file, and to contrast results to entering the same body weight-at-age data here and to attempt to estimate growth parameters, potentially time-varying, that match these body weight data.

- If using mean size-at-age data, please see the lambda usage notes regarding issues for model fitting depending upon other data within the model.

Environmental Data

SS accepts input of time series of environmental data. Parameters can be made to be time-varying by making them a function of one of these environmental time series.

Parameter values can be a function of an environmental data series:

1	Number of environmental variables		
COND > 0 Example of 2 environmental observations:			
	Year	Variable	Value
	1990	1	0.10
	1991	1	0.15
	-9999	0	0

Note

- Any years for which environmental data are not read are assigned a value of 0.0. None of the current link functions in SS contain a link parameter that acts as an offset. Therefore, you should subtract the mean from your data before entering it into SS. This lessens the problem with missing observations, but does not eliminate it. A better approach for dealing with missing observations is to use a different approach for the environmental effect on the parameter. Set up the parameter to have random deviations for all years, then enter the zero-centered environmental information as a special survey of type 35 and set up the catchability of that survey to be a link to the deviation vector. This is definitely a more complex approach and not much used, but it is superior in treatment of missing values and superior in allowing for error in the environmental relationship.
- It is permissible to include a year that is one year before the start year in order to assign environmental conditions for the initial equilibrium year. But this works only for recruitment parameters, not biology or selectivity parameters.
- Environmental data can be read for up to 100 years after the end year of the model. Then, if the recruitment-environment link has been activated, the future recruitments will be influenced by any future environmental data. This could be used to create a future "regime shift" by setting historical values of the relevant environmental variable equal to zero and future values equal to 1, in which case the magnitude of the regime shift would be dictated by the value of the

environmental linkage parameter. Note that only future recruitment and growth can be modified by the environmental inputs; there are no options to allow environmentally-linked selectivity in the forecast years.

- Some model derived quantities like summary biomass and recruitment deviation are assigned to some negative valued environmental variables. This is a stepping stone towards creating ability for parameters to be density-dependent.

Generalized Size Composition Data

A flexible feature with SS is a generalized approach to size composition information. It was designed initially to provide a means to include weight frequency data, but was implemented to provide a generalized capability. The user can define as many size frequency methods as necessary.

- Each method has a specified number of bins.
- Each method has "units" so the frequencies can be in units of biomass or numbers.
- Each method has "scale" so the bins can be in terms of weight or length (including ability to convert bin definitions in pounds or inches to kg or cm).
- The composition data is input as females then males, just like all other composition data in SS. So, in a two-sex model, the new composition data can be combined sex, single sex, or both sex.
- If a retention function has been defined, then the new composition data can be from the combined discard + retained, discard only or retained only.
- There are two options for treating fish that in population size bins that are smaller than the smallest size frequency bin.
 - Option 1: By default, these fish are excluded (unlike length composition data where the small fish are automatically accumulated up into the first bin).
 - Option 2: If the first size bin is given as a negative value, then, accumulation is turned on and the negative of the entered value is used as the lower edge of the first size bin.

Example entry:

2	N of size frequency methods, if this value is 0, then omit all entries below.
---	---

COND > 0

25	15	Nbins per method
2	3	Units per each method (1 = biomass, 2 = numbers)
3	3	Scale per each method (1 = kg, 2 = lbs, 3 = cm, 4 = inches)
1e-9	1e-9	Min compression to add to each observation (entry for each method)
2	2	N observations per weight frequency method

Then enter the lower edge of the bins for each method. The two row vectors shown below contain the bin definitions for methods 1 and 2 respectively:

-26	28	30	32	34	36	38	40	42	...	60	62	64	68	72	76	80	90
-26	28	30	32	34	36	38	40	42	44	46	48	50	52	54			

An example observation is below. Note that the format is identical to the length composition data, including sex and partition options, except for the addition of the first column to indicate the size frequency method.

Method	Year	Month	Fleet	Sex	Part	Sample Size	<composition females then males>
1	1975	1	1	3	0	43	<data>
1	1977	1	1	3	0	43	<data>
1	1979	1	1	3	0	43	<data>
1	1980	1	1	3	0	43	<data>

Note

- There is no tail compression for generalized size frequency data.
- Super-period capability is enabled in same way as for length and age composition data.
- By choosing units=2 and scale=3, the size composition method can be nearly identical to the length composition method if the bins are set identically.
- Bin boundaries can be real numbers so obviously do not have to align with population length bin boundaries. SS interpolates as necessary.

- Size bins cannot be defined to be narrower than the population binwidth; an untrapped error will occur.
- Because the transition matrix can depend upon weight-at-length, it is calculated internally for each sex and for each season because weight-at-length can differ between sexes and can vary seasonally.

Tag-Recapture Data

The ability to analyze tag-recapture data is available with SS. Each released tag group is characterized by an area, time, sex and age at release. Each recapture event is characterized by a time and fleet. SS fleets each operate in only one area, hence, it is not necessary to record the area of recapture. Inside the model, the tagged cohort is apportioned across all growth patterns in that area at that time (with options to apportion to only one sex or to both). The tag cohort combined with the growth pattern then behaves according to the movement and mortality of that growth pattern. The number of tagged fish is modeled as a negligible fraction of the total population. This means that a tagging event does not move fish from an untagged group to a tagged group. Instead, it acts as if the tags are seeded into the population with no impact at all on the total population abundance or mortality. The choice to require assignment of a predominant age at release for each tag group is a pragmatic coding and model efficiency choice. By assigning a tag group to a single age, rather than distributing it across all possible ages according to the size composition of the release group, it can be tracked as a single diagonal cohort through the age x time matrix with minimal overhead to the rest of the model. Tags are considered to be released at the beginning of a season and recaptures follow the timing of the fleet that made the recapture.

Example set-up for tagging data:

1	Do tags - if this value is 0, then omit all entries below.							
COND = 1 All subsequent tag-recapture entries must be omitted if "Do Tags" = 0								
3	Number of tag groups							
7	Number of recapture events							
2	Mixing latency period: N periods to delay before comparing observed to expected recoveries (0 = release period)							
10	Max periods (months) to track recoveries, after which tags enter accumulator							
Release Data								
TG	Area	Year	Month	<tfill>	Sex	Age	N Release	
1	1	1980	1	999	0	24	2000	
2	1	1995	1	999	1	24	1000	
3	1	1985	1	999	2	24	10	
Recapture Data								
TG	Year		Month		Fleet		Number	
1	1982		1		1		7	
1	1982		1		2		5	
1	1985		1		2		0	
2	1997		1		1		6	
2	1997		2		1		4	
3	1986		1		1		7	
3	1986		2		1		5	

Note

- The release data must be entered in tag group order.
- <tfill> values are place holders and are replaced by program generated values for model time.
- Analysis of the tag-recapture data has one negative log likelihood component for the distribution of recaptures across areas and another negative log likelihood component for the

decay of tag recaptures from a group over time, hence informative about mortality. More on this is in the control file.

- Warning for earlier versions of SS: A shortcoming in the recapture calculations when also using Pope's F approach was identified and corrected in version 3.30.14.

Stock (Morph) Composition Data

It is sometimes possible to observe the fraction of a sample that is composed of fish from different stocks. These data could come from genetics, otolith microchemistry, tags, or other means. The growth pattern feature in SS allows definition of cohorts of fish that have different biological characteristics and which are independently tracked as they move among areas. SS now incorporates the capability to calculate the expected proportion of a sample of fish that come from different growth patterns, "morphs". In the inaugural application of this feature, there was a 3 area model with one stock spawning and recruiting in area 1, the other stock in area 3, then seasonally the stocks would move into area 2 where stock composition observations were collected, then they moved back to their natal area later in the year.

Stock composition by growth pattern (morph) data can be entered in SS as follows:

1	Do morph composition, if zero, then do not enter any further input below.						
COND = 1							
3	Number of observations						
2	Number of morphs						
0.0001	Minimum Compression						
Year	Month	Fleet	Null	Nsamp	Data by N Morphs		
1980	1	1	0	36	0.4	0.6	
1981	1	1	0	40	0.44	0.54	
1982	1	1	0	50	0.37	0.63	

Note

- The number of stocks entered with these data must match the number of growth patterns (morphs) in the control file.

- Each data line for unique observations should enter data for morph 1 first followed sequentially for each morph included in the model.
- The expected value is combined across sexes. The entered data values will be normalized to sum to one within SS.
- The "null" flag is included here in the data input section and is a reserved spot for future features.
- Note that there is a specific value of minimum compression to add to all values of observed and expected.
- Warning for earlier versions of SS: A flaw was identified in the calculation of accumulation by morph. This has been corrected in version 3.30.14. Older versions were incorrectly calculating the catch by morph using the expectation around age-at-length which already was accounting for the accumulation by morph.

Selectivity Empirical Data (future feature)

It is sometimes possible to conduct field experiments or other studies to provide direct information about the selectivity of a particular length or age relative to the length or age that has peak selectivity, or to have a prior for selectivity that is more easily stated than a prior on a highly transformed selectivity parameter. This section provides a way to input data that would be compared to the specified derived value for selectivity. This is a placeholder at this time, required to include in the data file and will be fully implemented soon.

Selectivity data feature is under development for a future option and is not yet implemented. The input line still must be specified in SS as follows:

0	Do data read for selectivity (future option)
---	--

End of Data File

999	#End of data file marker
-----	--------------------------

Excluding Data

Data that are before the model start year or greater than the retrospective year are not moved into the internal working arrays at all. So if you have any alternative observations that are used in some model

runs and not in others, you can simply give them a negative year value rather than having to comment them out. The first output to `data.ss_new` has the unaltered and complete input data. Subsequent reports to `data.ss_new` produce expected values or bootstraps only for the data that are being used. Additional information on bootstrapping is available in Bootstrap Data Files Section.

Data that are to be included in the calculations of expected values, but excluded from the calculation of negative log likelihood, are flagged by use of a negative value for fleet number.

Data Super-Periods

The super-period capability allows the user to introduce data that represent a blend across a set of time steps and to cause the model to create an expected value for this observation that uses the same set of time steps. The option is available for all types of data and a similar syntax is used.

All super-period observations must be contiguous in the data file. All but one of the observations in the sequence will have a negative value for fleet ID so the data associated with these dummy observations will be ignored. The observed values must be combined outside of the model and then inserted into the data file for the one observation with a positive fleet number.

Super-periods are started with a negative value for month, and then stopped with a negative value for month, observations within the super-period are designated with a negative fleet field. The standard error or input sample size field is now used for weighting of the expected values. An error message is generated if the super-period does not contain one observation with a positive fleet field.

An expected value for the observation will be computed for each selected time period within the super-period. The expected values are weighted according to the values entered in the standard error (or input sample size) field for all observations except the single observation holding the combined data. The expected value for that year gets a relative weight of 1.0. So in the example below, the relative weights are: 1982, 1.0 (fixed); 1983, 0.85; 1985, 0.4; 1986, 0.4. These weights are summed and rescaled to sum to 1.0, and are output in the `echoinput.sso` file.

Not all time steps within the extent of a super-period need be included. For example, in a three season model, a super-period could be set up to combine information from season 2 across 3 years, e.g., skip over the season 1 and season 3 for the purposes of calculating the expected value for the super-period. The key is to create a dummy observation (negative fleet value) for all time steps, except 1, that will be included in the super-period and to include one real observation (positive fleet value; which contains the real combined data from all the specified time steps).

Super-period example:

Year	Month	Fleet	Obs	SE	Comment
1982	-2	3	34.2	0.3	Start super-period. This observation has positive fleet value, so is expected to contain combined data from all identified periods of the super-period. The standard error (SE) entered here is use as the SE of the combined observation. The expected value for the survey in 1982 will have a relative weight of 1.0 (default) in calculating the combined expected value.
1983	2	-3	55	0.3	In super-period; entered observation is ignored. The expected value for the survey in 1983 will have a relative weight equal to the value in the standard error field (0.85) in calculating the combined expected value.
1985	2	-3	88	0.40	Note that 1984 is not included in the super-period Relative weight for 1985 is 0.4
1986	-2	-3	88	0.40	End super-period

A time step that is within the time extent of the super-period can still have its own separate observation. In the above example, the survey observation in 1984 could be entered as a separate observation, but it must not be entered inside of the contiguous block of super-period observations. For composition data (which allow for replicate observations), a particular time steps' observations could be entered as a member of a super-period and as a separate observation.

The super-period concept can also be used to combine seasons within a year with multiple seasons. This usage could be preferred if fish are growing rapidly within the year so their effective age selectivity is changing within year as they grow; fish are growing within the year so fishery data collected year round have a broader size-at-age modes than a mid-year model approximation can produce; and it could be useful in situations with very high fishing mortality.

Control File

Overview of Control File

These listed model features are denoted in the control file in the following order:

1. Number of growth patterns and platoons
2. Design matrix for assignment of recruitment to area/settlement event/growth pattern
3. Design matrix for movement between areas
4. Definition of time blocks that can be used for time-varying parameters
5. Controls for all time-varying parameters
6. Specification for growth and fecundity
7. Natural mortality growth parameters, weight-at-length, maturity, and fecundity, for each sex
8. Hermaphroditism parameter line (if used)
9. Recruitment distribution parameters for each area, settlement event, and growth pattern
10. Cohort growth deviation
11. Movement between areas (if used)
12. Age error parameter line (if used)
13. Catch multiplier (if used)
14. Fraction female
15. Setup for any mortality-growth parameters are time-varying
16. Seasonal effects on biology parameters
17. Spawner-recruitment parameters
18. Setup for any stock recruitment parameters are time-varying
19. Recruitment deviations

20. F ballpark value in specified year
21. Method for calculating fishing mortality (F)
22. Initial equilibrium F for each fleet
23. Catchability (Q) setup for each fleet and survey
24. Catchability parameters
25. Setup for any catchability parameters are time-varying
26. Length selectivity, retention, discard mortality setup for each fleet and survey
27. Age selectivity setup for each fleet and survey
28. Parameters for length selectivity, retention, discard mortality for each fleet and survey
29. Parameters for age selectivity, retention, discard mortality for each fleet and survey
30. Setup for any selectivity parameters that are time-varying
31. Tag-recapture parameters
32. Variance adjustments
33. Lambdas for likelihood components

The order in which they appear in the control file has grown over time rather opportunistically, so it may not appear particularly logical at this time, especially various aspects of recruitment distribution and growth. When the same information is entered via the SS GUI, it is organized more logically and then written in this form to the text control file.

Parameter Line Elements

The primary role of the SS control file is to define the parameters to be used by the model. The general syntax of the 14 elements of a long parameter line is described here. If used, time-varying parameter lines use only the first seven elements of a parameter line and will be referred to as a short parameter line. Three types of time-varying properties can be applied to a base parameter: blocks or trend, environmental linkage, and random deviation. Each parameter line contains:

Column	Element	Description
1	LO	Minimum value for the parameter
2	HI	Maximum value for the parameter
3	INIT	Initial value for the parameter. If the phase (described below) for the parameter is negative the parameter is fixed at this value. If the ss.par file is read, it overwrites these INIT values.
4	PRIOR	Expected value for the parameter. This value is ignored if the prior type is 0 (no prior) or 1 (symmetric beta). If the selected prior type (described below) is lognormal, this value is entered in log space.
5	PRIOR SD	Standard deviation for the prior, used to calculate likelihood of the current parameter value. This value is ignored if prior type is 0.
6	PRIOR TYPE	0 = none, 1 = symmetric beta; 2 = full beta; 3 = lognormal without bias adjustment; 4 = lognormal with bias adjustment; 5 = gamma; and 6 = normal.
7	PHASE	Phase in which parameter begins to be estimated. A negative value causes the parameter to retain its INIT value (or value read from the ss.par file).
8	Env var & Link	Create a linkage to an input environmental time-series
9	Dev link	Invokes use of the deviation vector in the linkage function
10	Dev min yr	Beginning year for the deviation vector
11	Dev max yr	Ending year for the deviation vector
12	Dev phase	Phase for estimation for elements in the deviation vector
13	Block	Time block or trend to be applied
14	Block function	Functional form for the block offset.

Note that relative to SS v.3.24, the order of PRIOR SD and PRIOR TYPE have been switched and the PRIOR TYPE options have been renumbered.

The full parameter line (14 in length) syntax for the mortality-growth, spawn-recruitment, catchability, and selectivity sections provides additional controls to give the parameter time-varying properties. If a parameter (a full parameter line of length 14) is set up to be time-varying (i.e., parameter time blocks, annual deviations), short parameter lines, the first 7 elements, are required to be specified immediately after the main parameter block (i.e., mortality-growth parameter section). Additional information regard time-varying parameters and how to implement them in SS is available below, with in depth information provided in the Using Time-Varying Parameters section.

Terminology

The term COND appears in the "Typical Value" column of this documentation (it does not actually appear in the model files), it indicates that the following section is omitted except under certain conditions, or that the factors included in the following section depend upon certain conditions. In most cases, the description in the definition column is the same as the label output to the ss_new files.

Beginning of Control File Inputs

Typical Value	Description and Options
#C comment	Comments beginning with #C at the top of the file will be retained and included in output.
0	<p>0 = Do not read the weight-at-age (wtatage.ss) file;</p> <p>1 = Read the weight-at-age (wtatage.ss) file, also read and use the growth parameters; and</p> <p>2 = Future option to read the weight-at-age (wtatage.ss) file, then omit reading and using growth parameters and all length-based data.</p> <p>Additional information on the weight-at-age file and the expected formatting can be found in the Empirical Weight-at-Age section.</p>

Typical Value	Description and Options
1	<p>Number (N) of growth patterns (GP), also referred to as morphs:</p> <p>These are collections of fish with unique biological characteristics (growth, mortality, weight-length, reproduction). The GP x Sex x Settlement Events constitute unique growth patterns that are tracked in SS. They are assigned these characteristics at birth and retain them throughout their lifetime. At recruitment, growth pattern members are distributed across areas (if any) and they retain their biological characteristics even if they move to another area in which a different cohort with different biological characteristics might predominate. For example, one could assign a fast-growing growth pattern to recruit predominately in a southern areas and a slow-growing growth pattern to a northern area. The natural mortality and growth parameters are specified for each growth pattern in the mortality-growth parameter section in the order of females growth pattern 1 to growth pattern N followed by males growth pattern 1 to growth pattern N in a two sex model.</p>
1	<p>Number of platoons within a growth pattern/morph:</p> <p>This allows exploration of size-dependent survivorship. A value of 1 will not create additional platoons. Odd-numbered values (i.e., 3, 5) will break the overall morph into that number of platoons creating a smaller, larger, and mean growth platoon. The higher the number of platoons the slower SS will run, so values above 5 not advised. The fraction of each morph assigned to each platoon is custom-input or designated to be a normal approximation. When multiple platoons are designated, an additional input is the ratio of between platoon to within platoon variability in size-at-age. This is used to partition the total growth variability. For the platoons, their size-at-age is calculated as a factor (determined from the between-within variability calculation) times the size-at-age of the central morph which is determined from the growth parameters for that Growth Pattern x Sex.</p>
COND > 1	Following 2 lines are conditional on N platoons > 1.
0.7	Platoon between/within standard deviation ratio. Ratio of the amount of variability in length-at-age between platoons to within platoons.
0.2 0.6 0.2	Distribution among platoons. Enter either a custom vector or enter a vector of length N with the first value of -1 to get a normal approximation: (0.15, 0.70, 0.15) for 3 platoons, or 5 platoons (0.031, 0.237, 0.464, 0.237, 0.031).

Weight-at-Age

SS added the capability to read empirical body weight-at-age for the population and each fleet, starting in version 3.04, in lieu of generating these weights internally from the growth parameters, weight-at-length, and size-selectivity. The values are read from a separate file named, wtatage.ss. This file is only required to exist if this option is selected. See the Empirical Weight-at-Age section for additional information on file formatting for empirical weight-at-age.

Settlement Timing for Recruits and Distribution

In older versions of SS one value of spawning biomass was calculated annually at the beginning of one specified spawning season and this spawning biomass produced one annual total recruitment value. The annual recruitment value was then distributed among seasons, areas, and growth types according to other model parameters.

Additional control of the seasonal timing was added in SS v.3.30 and now there now is an explicit elapsed time between spawning and recruitment. Spawning still occurs, just once per year, which defines a single spawning biomass for the stock-recruitment curve but its timing can be at any specified time, not just the beginning of a season. Recruitment of the progeny from an annual spawning can now enter the population in one or more settlement events, at some point after spawning as defined by the user.

Typical Value	Description and Options
1	Recruitment distribution method. This section controls which combinations of growth pattern x area x settlement will get a portion of the total recruitment coming from each spawning. Options: 1 = no longer available (used the SS v.3.24 or earlier setup); 2 = main effects for growth pattern, settle timing, and area; 3 = each settle entity; and 4 = none, no parameters (only if growth pattern x settlement x area = 1).
1	Spawner-Recruitment (not implement yet, but required), options: 1 = global; and 2 = by area (by area is not yet implemented; there is a conceptual challenge to doing the equilibrium calculation when there is fishing).

Typical Value		Description and Options	
1		Number of recruitment settlement assignments. Must be at least 1 even if only 1 settlement and 1 area because the timing of that settlement must be specified.	
0		Future feature, not implement yet but required.	
Growth Pattern	Month	Area	Age at settlement
1	5.5	1	0

The above example specifies settlement to mid-May (month 5.5). Note that normally the calendar age at settlement is 0 if settlement happens between the time of spawning and the end of that year, and at age 1 if settlement is in the year after spawning.

Below is an example set-up where there are multiple settlement events, with one occurring the following year after spawning:

3	Number of recruitment settlement events		
0	Unused option		
Growth Pattern	Month	Area	Age (for each settlement assignment)
1	11.0	1	0
1	12.0	1	0
1	1.0	1	1

Details regarding settlement of recruits and timing:

- Recruitment happens in specified settlement events (growth pattern, month, area).
- Number of unique settlement timings is calculated at runtime.
- Now there is explicit elapsed time between spawning and recruitment.
- Growth and natural mortality of the platoon begins at time of settlement, which is its real age 0.0 for growth; but pre-settlement fish exist from the beginning of the season of settlement, so can be caught if selected.
- Age at recruitment now user-controlled (should be 0 if in year of spawning).

- All fish become integer age 1 (for age determination) on their first January 1st.
- Recruitment can occur >12 months after spawning which is achieved by setting the settlement age to a value greater than 1.0.

The distribution of recruitment among these settlement events is controlled by recruitment apportionment parameters. There must be a parameter line for each growth pattern, then for each area, then for each settlement. All of these are required, but only those growth pattern x area x settlements designated to receive recruits in the recruitment design matrix will have the parameter used in the recruitment distribution calculation. For the recruitment apportionment, the parameter values are the natural log of apportionment weight. The sum of all apportionment weights is calculated for each growth pattern x area x settlements that have been designated to receive recruits in the recruitment design matrix. Then the apportionment weights are scaled to sum to 1.0 so that the total recruitment from the spawning event is distributed among the cells designated to receive recruitment. These distribution parameters can be time-varying, so the fraction of the recruits that occur in a particular growth pattern, area, or settlement can change from year to year.

Recruitment Distribution and Parameters

SS apports recruits according to:

$$\text{apportionment}_i = \frac{e^{p_i}}{\sum_{i=1}^N e^{p_i}} \quad (1)$$

where p_i is the proportion of recruits to area i and N is the number of settlement events. These parameters are defined in the mortality-growth parameter section.

Tips for fixing or estimating the recruitment apportionment:

- Set the value for one of these parameters, p_i , to 0.0 and not estimate it so that other parameters will be estimated (if not fixed) relative to its fixed value.
- Give the estimated parameters a min-max so they have a good range relative to the base parameter (i.e., of min = -5 and max = 5).
- In order to get a different distribution of recruitments in different years, you will need to make at least one of the recruitment distribution parameters time-varying.

In a seasonal model, all cohorts graduate to the age of 1 when they first reach January 1, even if the seasonal structure of the model has them being spawned in the late fall. In general, this means that SS operates under the assumption that all age data have been adjusted so that fish are age 0 at the time of spawning and all fish graduate to the next age on January 1. This can be problematic if the ageing structures deposit a ring at another time of year. Consequently, you may need to add or subtract a

year to some of your age data to make it conform to the SS structure, or more ideally you may need to define the SS calendar year to start at the beginning of the season at which ring deposition occurs. Talk with your ageing lab about their criteria for seasonal ring deposition.

Seasonal recruitment is coded to work smoothly with growth. If the recruitment occurring in each season is assigned the same growth pattern, then each seasonal cohort's growth trajectory is simply shifted along the age/time axis. At the end of the year, the early born cohorts will be larger, but all are growing with the same growth parameters, so all will converge in size as they approach their common maximum length (e.g., no seasonal effects on growth).

At the time of settlement, fish are assigned a size equal to the lower edge of the first population size bin and they grow linearly until they reach the age A1. SS generates a warning if the first population length bin is greater than 10 cm as this seems an unreasonably large value for a larval fish. A1 is in terms of real age elapsed since birth. All fish advance to the next integer age on January 1, regardless of birth season. For example, consider a 2 season model with some recruitment in each season and with each season's recruits coming from the same GP. At the end of the first year, the early born fish will be larger but both of the seasonal cohorts will advance to an integer age of 1 on Jan 1 of the next year. The full growth curve is still calculated below A1, but the size-at-age used by SS is the linear replacement. Because the linear growth trajectory can never go negative, there is no need for the additive constant to the standard deviation (necessary for the growth model used in SS2 V1.x), but the option to add a constant has been retained in the model.

Movement

Here the movement of fish between areas are defined. This is a box transfer with no explicit adjacency of areas, so fish can move from any area to any other area in each time step. While not incorporated yet, there is a desire for future versions of SS to have the capability to allow sex-specific movement, and also to allow some sort of mirroring so that sexes and growth patterns can share the same movement parameters if desired.

Typical Value	Description and Options
COND: only if areas > 1	
2	Enter Number of movement definitions.

Typical Value	Description and Options
1.0	First age that moves. This value is a real number, not an integer, to allow for an in-year start to movement in a multi-season model. It is the real age at the beginning of a season, even though movement does not occur until the end of the season. For example, in a setup with two 6-month seasons a value of 0.5 will cause the age 0 fish to not move when they complete their first 6 month season of life, and then to move at the end of their second season because they start movement capability when they reach the age of 0.5 years (6 months).
1 1 1 2 4 10 1 2 2 1 4 10	Movement definitions: season, growth pattern, source area, destination, age1, and age2. The example shown here has 1 growth patterns and 2 areas with fish moving between the two areas. The rate of movement will be controlled by the movement parameters later defined in the mortality-growth parameter section. Here the age1 and age2 specify the range over which the movement parameters are interpolated with movement constant below age1 and above age2.

Two parameters will be entered later for each growth pattern, area pair, and season.

- Movement is constant at the first parameter (P1) below the specified minimum age for movement change, constant at the second parameter (P2) above maximum age for movement change, and linearly interpolated for intermediate ages.
- A movement rate parameter can be set to use the same value as the corresponding parameter for the first defined movement pattern by entering a parameter value of -9999 and a negative phase value.
- For each source area, the implicit movement parameter value is 0.0 (movement within a single area). However, this default value can be replaced if the stay movement is selected to have an explicit pair of parameters (e.g., specify movement rate for area 1 to area 1) and will require additional parameter lines.
- A constant movement rate across all ages can be accomplished by either:
 - Setting both movement ages to 0, not estimating the first movement parameter, and using a second movement parameter to cover all ages from 0 to the maximum number of ages.
 - Setting movement ages to any value, estimating the first movement parameter, and setting the second movement parameter to have a value of -9998 with a negative phase.
- The parameter is exponentiated so that a movement parameter value of 0 becomes 1.0.

- For each source area, all movement rates are then summed and divided by this sum so that 100% of the fish are accounted for in the movement calculations.

$$\text{rate}_i = \frac{e^{p_i}}{\sum_{j=1}^N e^{p_j}} \quad (2)$$

- At least one movement parameter must be fixed so that all other movement parameters are estimated relative to it. This is achieved naturally by not specifying the stay rate parameter so it has a fixed value of 0.0.
- The resultant movement rates are multiplied by season duration in a seasonal model.

Time Blocks

Typical Value	Description and Options
3	Number of block patterns. These patterns can be referred to in the parameter sections to create a separate parameter value for each block.
COND > 0: 3 2 1	Following inputs are omitted if the number of block patterns equals 0. Blocks per pattern:
1975 1985 1986 1990 1995 2001	Beginning and ending years for blocks in design 1; years not assigned to a block period retain the baseline value for a parameter that uses this pattern.
1987 1990 1995 2001	Beginning and ending years for blocks in design 2.
1999 2002	Beginning and ending years for blocks in design 3.

Auto-generation

Auto-generation is a useful way to automatically have SS create the required short time-varying parameter lines which will be written in the control.ss_new file. These parameter lines can then be copied into the control file and modified as needed. As example, if you want to add a block to natural mortality, modify the block and block function entry of the mortality parameter line, ensure that auto-generation is set to 0 (for the biology section at least) and run the model without estimation.

The control.ss_new file will now show the required block parameter line specification for natural mortality and this line can be copied into the main control file. Note, that if auto-generation is on (set to 0), the model will not expect to read the time-varying parameters in that section of the control file and will error out if they are present

Typical Value	Description and Options
1	Environmental/Block/Deviation adjust method for all time-varying parameters. 1 = warning relative to base parameter bounds; and 3 = no bound check. Logistic bound check form from previous SS versions (e.g., SS v.3.24) is no longer an option.
1 1 1 1 1	Auto-generation of time-varying parameter lines. Five values control auto-generation for parameter block sections: 1-biology, 2-spawn-recruitment, 3-catchability, 4-tag (future), and 5-selectivity. The accepted values are: 0 = auto-generate all time-varying parameters (no time-varying parameters are expected); 1 = read each time-varying parameter line as exists in the control file; and 2 = read each line and auto-generate if read if the time-varying parameter value for LO = -12345. Useful to generate reasonable starting values.

Biology

Natural Mortality

Natural mortality (M) has some options that are referenced to integer age, and some to real age since settlement. So, if M varies by age, M will change by season and cohorts born early in the year will have different M than late born cohorts.

Lorenzen natural mortality is based on the concept that natural mortality varies over the life cycle of a fish, which is driven by physiological and ecological processes. So, natural mortality is scaled by the length of the fish.

Typical Value	Description and Options
1	Natural Mortality Options: 0 = A single parameter;

Typical Value	Description and Options
	1 = N breakpoints; 2 = Lorenzen; 3 = Read age specific M and do not do seasonal interpolation; and 4 = Read age specific and do seasonal interpolation, if appropriate.
COND = 0	No additional natural mortality controls.
COND = 1	
4	Number of breakpoints. Then read a vector of ages for these breakpoints. Later, per sex x GP, read N parameters for the natural mortality at each breakpoint.
2.5 4.5 9.0 15.0	Vector of age breakpoints.
COND = 2	
4	Reference age for Lorenzen natural mortality: read one additional integer value that is the reference age. Later read one parameter for each sex x growth pattern that will be the M at the reference age. Other ages will have an M scaled to its body size-at-age. However, if platoons are used, all will have the same M as their growth pattern. Lorenzen M calculation will be updated if the starting year growth parameters are active, but if growth parameters vary during the time-series, the M is not further updated. So be careful in using Lorenzen when there is time-varying growth.
COND = 3 or 4	Do not read any natural mortality parameters in the mortality growth parameter section. With option 3, these M values are held fixed for the integer age (no seasonality or birth season considerations). With option 4, there is seasonal interpolation based on real age, just as in options 1 and 2.
0.20 0.25 ... 0.20 0.23 ...	Age-specific M values where in a 2 sex model the first row is female and the second row is male. If there are multiple growth patterns female growth pattern 1-N is read first followed by males 1-N growth pattern.

Growth

Timing

When fish recruit at the real age of 0.0 at settlement, they have body size equal to the lower edge of the first population size bin. The fish then grow linearly until they reach a real age equal to the input value “growth-at-age for L1” and have a size equal to the parameter value for L1 (the minimum length parameter). As they age further, they grow according the selected growth equation. The growth curve is calibrated to go through the size L2 parameter when they reach the age of maximum length.

Maximum Length (Linf)

If “Growth at age for L2” is set equal to 999, then the size at the L2 parameter is used as Linf.

Schnute growth function

The Schnute implementation of a 3-parameter growth function is invoked by entering 2 in the growth type field. Then a fourth parameter is read after reading the von Bertalanffy growth coefficient parameter (K). When this fourth parameter has a value of 1.0, it is equivalent to the standard von Bertalanffy growth curve. When this function was first introduced in SS, it required that A0 be set to 0.0.

Mean size-at-maximum age

The mean size of fish in the max age age bin depends upon how close the growth curve is to Linf by the time it reaches max age and the mortality rate of fish after they reach max age. SS provides an option for the mortality rate to use in this calculation during the initial equilibrium year. This must be specified by the user and should be reasonably close to M plus initial F. In SS v.3.30, this uses the von Bertalanffy growth out to 3 times the maximum population age and decays the numbers at age by $\exp(-\text{value set here})$. For subsequent years of the time series, SS should update the size-at-maximum age according to the weighted average mean size of fish already at maximum age and the size of fish just graduating into maximum age. Unfortunately, this updating is only happening in years with time-varying growth. This will hopefully be fixed in a the future version.

Age-specific K

This option creates age-specific K multipliers for each age of a user-specified age range, with independent multiplicative factors for each age in the range and for each growth pattern / sex. The null value is 1.0 and each age's K is set to the next earlier age's K times the value of the current age's multiplier. Each of these multipliers is entered as a full parameter line, so inherits all time-varying capabilities of full parameters. The lower end of this age range cannot extend younger than the specified age for which the first growth parameter applies. This is a beta model feature, so examine output closely to assure you are getting the size-at-age pattern you expect. Beware of using this option in a model with seasons within year because the K deviations are indexed solely by integer age according to birth year. There is no offset for birth season timing effects, nor is there any seasonal interpolation of the age-varying K.

Growth cessation

A growth cessation model was developed for the application to tropical tuna species (Maunder et al., 2018). Growth cessation allows for a linear relationship between length and age, followed by a marked reduction of growth after the onset of sexual maturity by assuming linear growth for the youngest individuals and then a logistic function to model the decreasing growth rate at older ages.

Example growth specifications:

Typical Value	Description and Options
1	<p>Growth Model:</p> <p>1 = von Bertalanffy (2 parameters);</p> <p>2 = Schnute's generalized growth curve (aka Richards curve) with 3 parameters. Third parameter has null value of 1.0;</p> <p>3 = von Bertalanffy with age-specific K multipliers for specified range of ages, requires additional inputs below following the placeholder for future growth feature;</p> <p>4 = age specific K. Set base K as K for age = nages and working backwards and the age-specific K = K for the next older age * multiplier, requires additional inputs below following the placeholder for future growth feature;</p> <p>5 = age specific K. Set base K as K for nages and work backwards and the age-specific K = base K * multiplier, requires additional inputs below following the placeholder for future growth feature;</p> <p>6 = not implemented;</p> <p>7 = not implemented; and</p> <p>8 = growth cessation. Decreases the K for older fish. If implemented, the Amin and Amax parameters, the next two lines, need to be set at 0 and 999 respectively. The mortality-growth parameter section requires the base K parameter line which is interpreted as the steepness of the logistic function that models the reduction in the growth increment by age followed by a second parameter line which is the parameter related to the maximum growth rate.</p>
1.66	Growth Amin (A1): Reference age for first size-at-age (post-settlement) parameter.
25	Growth Amax (A2): Reference age for second size-at-age parameter (999 to use as L infinity).
0.20	<p>Exponential decay for growth above maximum age (plus group: fixed at 0.20 in SS v.3.24; should approximate initial Z). Alternative Options:</p> <p>-998 = Disable growth above maximum age (plus group) similar to earlier versions of SS (prior to SS v.3.24); and</p> <p>-999 = Replicate the simpler calculation done in SS v.3.24.</p>
0	Placeholder for future growth feature.

Typical Value	Description and Options
COND = 3	Growth model: age-specific K
2	Number of K multipliers to read;
5	Minimum age for age-specific K; and
7	Maximum age for age-specific K.
COND = 4 or 5	Growth model: age-specific K
2	Number of K multipliers to read;
7	Maximum age for age-specific K;
6	Second age for age-specific K; and
5	Minimum age for age-specific K.
0	Standard deviation added to length-at-age: Enter 0.10 to mimic SS2 V1.xx. Recommend using a value of 0.0.
1	<p>CV Pattern (cannot be time-varying)</p> <p>0: $CV=f(LAA)$, so the 2 parameters are in terms of CV of the distribution of length-at-age and the interpolation between these 2 parameters is a function of mean length-at-age;</p> <p>1: $CV=f(A)$, so interpolation is a function of age;</p> <p>2: $SD=f(LAA)$, so parameters define the standard deviations of length-at-age and interpolation is a function of mean length-at-age;</p> <p>3: $SD=f(A)$; and</p> <p>4: Lognormal distribution of size-at-age. Input parameters will specify the standard deviation of natural log size-at-age (e.g., entered values will typically be between 0.05 and 0.15). A bias adjustment is applied so the lognormal distribution of size-at-age will have the same mean size as when a normal distribution is used.</p>

Maturity-Fecundity

Typical Value	Description and Options
2	<p>Maturity Option:</p> <p>1 = length logistic;</p> <p>2 = age logistic;</p> <p>3 = read maturity-at-age for each female growth pattern;</p> <p>4 = read a fecundity x maturity-at-age vector for all ages;</p>

Typical Value	Description and Options
	<p>5 = disabled; and</p> <p>6 = read vector of length-based maturity values.</p> <p>Note: need to read 2 parameter lines (maturity at 50% and maturity slope) even if option 3 or 4 is selected.</p>
COND = 3 or 4 0 0.05 0.10 ...	<p>Maturity Option</p> <p>Vector of age-specific maturity or fecundity. One row of length Nages + 1 based on the maximum population age for each female growth pattern.</p>
COND = 6 0 0.05 0.10 ...	<p>Maturity Option</p> <p>Vector of length-specific maturity or fecundity, based on the population length bins. One row of length equal to the number of population length bins (defined in the data file) for each female growth pattern.</p>
1	<p>First Mature Age: all ages below the first mature age will have maturity set to zero. This value is overridden if maturity option is 3 or 4 or if empirical weight-at-age (wtatage.ss) is used, but still must exist here.</p>
1	<p>Fecundity Option (irrelevant if maturity option is 4 or wtatage.ss is used):</p> <p>1 = to interpret the 2 egg parameters as linear eggs/kg on body weight (current SS default), so fecundity = $wt * (a + b * wt)$, so value of a=1, b=0 causes eggs to be equivalent to spawning biomass;</p> <p>2 = to set fecundity= $a * L^b$;</p> <p>3 = to set fecundity= $a * W^b$, so values of a=1, b=1 causes fecundity to be equivalent to spawning biomass;</p> <p>4 = fecundity = $a + b * L$; and</p> <p>5 = eggs = $a + b * wt$.</p>

Hermaphroditism

Typical Value	Description and Options
0	<p>Hermaphroditism Option:</p> <p>0 = not used;</p> <p>1 = invoke female to male age-specific function; and</p> <p>-1 = invoke male to female age-specific function.</p> <p>Note: this creates the annual, age-specific fraction that change sex, it is not the fraction that is each sex.</p>
COND = 1 or COND = -1	<p>Read 2 lines below if hermaphroditism is selected. Also read 3 parameters after reading the male (option 1) or female (option -1) weight-length parameter.</p>
-1	<p>Hermaphroditism Season:</p> <p>-1 to do transition at the end of each season (after mortality and before movement); and</p> <p><positive integer> to select just one season.</p>
1	<p>Include males or females, dependent upon sex change direction, in spawning biomass;</p> <p>0 = no males or females, dependent upon sex change direction, in spawning biomass; and</p> <p>1 = simple addition of males to females.</p>

The hermaphroditism option requires three full parameter lines in the mortality growth section:

1. A parameter line for inflection point (in age) at when fish may begin to change sex.
2. The standard deviation.
3. The asymptote of the maximum proportion that will transition to the other sex

These parameter lines are entered directly after the weight-at-length parameters for males.

Parameter offset method

Typical Value	Description and Options
1	<p>Parameter Offset Method:</p> <p>1 = direct assignment;</p> <p>2 = for each growth pattern by sex, parameter defines offset from sex 1, offsets are in exponential terms, so for example: $M_{\text{old male}} = M_{\text{old female}} * \exp(M_{\text{old male}})$; and</p> <p>3 = for each growth pattern by sex, parameter defines offset from growth pattern 1 sex 1. For females, given that "natM option" is breakpoint and there are two breakpoints, parameter defines offset from early age (e.g., $M_{\text{old female}} = M_{\text{young female}} * \exp(M_{\text{old female}})$). For males, given that "natM option" is breakpoint and there are two breakpoints, parameter is defined as offset from females AND from early age (e.g., $M_{\text{old male}} = M_{\text{young female}} * \exp(M_{\text{young male}}) * \exp(M_{\text{old male}})$).</p>

Catch Multiplier

These parameter lines are only included in the control file if the catch multiplier field in the data file is set to 1 for a fleet. The model expected catch C_{exp} by fleet is estimated by:

$$C_{exp} = \frac{C_{obs}}{c_{mult}} \quad (3)$$

where C_{obs} is the input catch by fleet (observed catch) within the data file and c_{mult} is the estimated (or fixed) catch multiplier. It has year-specific, not season-specific, time-varying capabilities. In the catch likelihood calculation, expected catch is multiplied by the catch multiplier by year and fishery to get C_{obs} before being compared to the observed retained catch as modified by the c_{mult} .

Ageing Error Parameters

These parameters are only included in the control file if one of the ageing error definitions in the data file has requested this feature (by putting a negative value for the ageing error of the age zero fish of one ageing error definition). As of version 3.30.12, these parameters now have time-varying capability. Seven additional full parameter lines are required. The parameter lines specify:

1. Age at which the estimated pattern begins (just linear below this age), this is the start age.

2. Bias at start age (as additive offset from unbiased age).
3. Bias at maximum (as additive offset from unbiased age).
4. Power function coefficient for interpolating between those 2 values (value of 0.0 produces linear interpolation in the bias).
5. Standard deviation at start age.
6. Standard deviation at max age.
7. Power function coefficient for interpolating between those 2 values.

Code for implementing vectors of mean age and standard deviation of age within SS can be located in the code examples.

Sex ratio

The last line in the mortality-growth parameter section allows the user to fix or estimate the sex ratio between female and male fish at recruitment. The parameter is specified in the fraction of female fish and is applied at settlement. The default option is a sex ratio of 0.50 with this parameter not being estimated. Any composition data input as type = 3, both sexes, will be informative to the sex ratio because SS scales females and males together, not separately, for this data type input. Estimation of the sex ratio is a new feature within SS and should be done with care with the user checking that the answer is reflective of the data.

As of v.3.30.12, this parameter now has time-varying capability similar to other parameters in the mortality-growth section.

Read Biology Parameters

Next, SS reads the mortality-growth (MG) parameters in generally the following order (may vary based on selected options):

Parameter	Description
Females	Female natural mortality and growth parameters in the following order by growth pattern.

Parameter	Description
M	Natural mortality for female growth pattern 1, where the number of natural mortality parameters depends on the option selected.
COND if M option = 1	
N breakpoints	N-1 parameter lines as an exponential offsets from the previous reference age.
Lmin	Length at Amin (units in cm) for female, growth pattern 1.
Lmax	Length at Amax (units in cm) for female, growth pattern 1.
VBK	von Bertalanffy growth coefficient (units are per year) for females, growth pattern 1.
COND if growth type = 2	
Richards Coefficient	Only include this parameter if Richards growth function is used. If included, a parameter value of 1.0 will have a null effect and produce a growth curve identical to von Bertalanffy.
COND if growth type >=3	
N parameter lines equal to the number K deviations for the ages specified above.	Age-Specific K
CV young	Variability for size at age <= Amin for females, growth pattern 1. Note that CV cannot vary over time, so do not set up env-link or a deviation vector. Also, units are either as CV or as standard deviation, depending on assigned value of CV pattern.
CV old	Variability for size at age >= Amax for females, growth pattern 1. For intermediate ages, do a linear interpolation of CV on means size-at-age. Note that the units for CV will depend on the CV pattern and the value of mortality-growth parameter as offset. The CV value cannot vary over time.
WtLen scale	Coefficient to convert length in cm to weight in kg for females.
WtLen exp	Exponent in to convert length to weight for females.
Mat-50%	Maturity logistic inflection (in cm or years) where female maturity-at-length (or age) is a logistic function: $M_l = 1/(1 + \exp(\alpha * (l_a - \beta)))$. The α is the slope, l_a is the size-at-age, and β is the inflection of the maturity curve. Value ignored for maturity option 3, 4, and 6.
Mat-slope	Logistic slope (must have negative value). Value ignored for maturity option 3, 4, and 6.

Parameter	Description
Eggs-alpha	Two fecundity parameters; usage depends on the selected fecundity option. Must be included here even if vector is read in the control section above.
Eggs-beta	
COND: growth pattern > 1	Repeat female parameters in the above order for growth pattern 2.
Males	Male natural mortality and growth parameters in the following order by growth pattern.
M	Natural mortality for male GP1, where the number of natural mortality parameters depends on the option selected.
COND if M option = 1	
N breakpoints	N-1 parameter lines as an exponential offsets from the previous reference age.
Lmin	Length at Amin (units in cm) for male, growth pattern 1. In a two sex model, fixing the INIT value a 0 will assume the same Lmin as the female parameter value.
Lmax	Length at Amax (units in cm) for male, growth pattern 1. In a two sex model, fixing the INIT value a 0 will assume the same Lmax as the female parameter value.
VBK	von Bertalanffy growth coefficient (units are per year) for males, growth pattern 1. In a two sex model, fixing the INIT value a 0 will assume the same k as the female parameter value.
COND if growth type = 2	
Richards Coefficient	Only include this parameter if Richards growth function is used. If included, a parameter value of 1.0 will have a null effect and produce a growth curve identical to Bertalanffy.
COND if growth type = 3	Age-Specific K
N parameter lines equal to the number K deviations for the ages specified above.	
CV young	Variability for size at age <= Amin for males, GP1. Note that CV cannot vary over time, so do not set up env-link or a deviation vector. Also, units are either as CV or as standard deviation, depending on assigned value of CV pattern.

Parameter	Description
CV old	Variability for size at age $\geq A_{max}$ for males, growth pattern 1. For intermediate ages, do a linear interpolation of CV on means size-at-age. Note that the units for CV will depend on the CV pattern and the value of mortality-growth parameters as offset.
WtLen scale WtLen exp	Coefficient to convert length in cm to weight in kg for males. Exponent to convert length to weight for males.
COND: growth pattern > 1	Repeat male parameters in the above order for growth pattern 2.
COND: Hermaphroditism	3 parameters lines define a normal distribution for the transition rate of females to males (or vice versa).
Inflect Age	Hermaphrodite inflection age.
StDev	Hermaphrodite standard deviation (in age) .
Asmp Rate	Hermaphrodite asymptotic rate.
COND: Recruitment Distribution Method = 2 or 3	3 parameters lines defining recruitment distribution.
Recruitment Dist. GP	Recruitment apportionment by growth pattern, if multiple growth patterns, multiple entries required.
Recruitment Dist. Area	Recruitment apportionment by area, if multiple areas, multiple entries required.
Recruitment Dist. Month	Recruitment apportionment by month, if multiple months, multiple entries required.
Cohort growth deviation	Set equal to 1.0 and do not estimate; it is deviations from this base that matter.
2 x N selected movement pairs	Movement parameters
COND: The following lines are only	required when the associated features are turned on.
Ageing Error	Turned on in the data file.
Catch	For each fleet selected for this option in the data file.
Multiplier	
Fraction female	Fraction female at the time of recruitment by growth pattern, if multiple growth patterns, multiple entries required.

Example format for mortality-growth parameter section with 2 sexes, 2 areas. Parameters marked with COND are conditional on selecting that feature:

LO	HI	INIT	Prior Value	<other entries>	Block Fxn	Parameter Label
0	0.50	0.15	0.1	...	0	#NatM_p_1_Fem_GP_1
0	45	21	36	...	0	#L_at_Amin_Fem_GP_1
40	90	70	70	...	0	#L_at_Amax_Fem_GP_1
0	0.25	0.15	0.10	...	0	#VonBert_K_Fem_GP_1
0.10	0.25	0.15	0.20	...	0	#CV_young_Fem_GP_1
0.10	0.25	0.15	0.20	...	0	#CV_old_Fem_GP_1
-3	3	2e-6	0	...	0	#Wtlen_1_Fem
-3	4	3	3	...	0	#Wtlen_2_Fem
50	60	55	55	...	0	#Mat50%_Fem
-3	3	-0.2	-0.2	...	0	#Mat_slope_Fem
-5	5	0	0	...	0	#Eggs/kg_inter_Fem
-50	5	0	0	...	0	#Eggs/kg_slope_wt_Fem
0	0.50	0.15	0.1	...	0	#NatM_p_1_Mal_GP_1
0	45	21	36	...	0	#L_at_Amin_Mal_GP_1
40	90	70	70	...	0	#L_at_Amax_Mal_GP_1
0	0.25	0.15	0.10	...	0	#VonBert_K_Mal_GP_1
0.10	0.25	0.15	0.20	...	0	#CV_young_Mal_GP_1
0.10	0.25	0.15	0.20	...	0	#CV_old_Mal_GP_1
-3	3	2e-6	0	...	0	#Wtlen_1_Mal
-3	4	3	3	...	0	#Wtlen_2_Mal
0	0	0	0	...	0	#RecrDist_GP_1
0	0	0	0	...	0	#RecrDist_Area_1
0	0	0	0	...	0	#RecrDist_Area_2
0	0	0	0	...	0	#RecrDist_Settlement_1
0.2	5	1	1	...	0	#Cohort_Grow_Dev
-5	5	-4	1	...	0	#Move_A_seas1_GP1_from_1to2 (COND)
-5	5	-4	1	...	0	#Move_B_seas1_GP1_from_1to2 (COND)
-99	99	1	0	...	0	#AgeKeyParm1 (COND)
-99	99	0.288	0	...	0	#Age_Key_Parms 2 to 5 (COND)
-99	99	0.715	0	...	0	#Age_Key_Parm6 (COND)
0.2	3.0	1.0	0	...	0	#Catch_mult_fleet1 (COND)

LO	HI	INIT	Prior Value	<other entries>	Block Type	Parameter Label
0.001	0.999	0.5	0.5	...	0	#Frac_Female_GP_1

Setting Male Parameters Equal to Females

SS allows a short-cut for males to use the same parameter values as female fish for natural mortality, length minimum (Length at Amin), maximum length (Length at Amax), coefficient at younger ages (CV1), coefficient at older ages (CV2), and the growth coefficient (K). If the INIT parameter value for males is set equal to 0.0 and the phase set to negative, not estimated, each of these male parameters will use the corresponding female parameter value for the males.

Time-Varying Parameter Specification and Setup

Any of the parameters defined above can be made time-varying through linkage to an environmental data series, through time blocks or trend, or by setting up annual deviations. The basic options for making biology, spawner-recruitment, catchability, and selectivity parameters change over time is detailed here. A more detailed in depth description is available in the Using Time-Varying Parameter section. After SS reads the parameter line section, which include options to create environmental link, blocks, or deviation vectors, if time-varying parameters are specified then SS expects to read the time-varying short parameter lines. Time-varying parameters are an optional input and if no parameters are specified to have time-varying properties, this omitted from the control file.

Special consideration should be given regarding which parameter values are applied during forecast years. The model will default to use all base parameter values during the forecast period. However, there are controls in the forecast file which allow the user to specify specific parameter years to be applied during the forecast period for selectivity, relative F, and recruitment.

Specification of Time-Varying Parameters in Long Parameter Lines

Selecting time-varying properties for a parameter is specified using element 8 - 14 in the long parameter line setup. Each element and the options for selection are as described here:

- Environmental Link and Variance (element 8)
 - The environmental link and variance (env_var&link) matrix is populated with the read environmental data for columns 1-N environmental variables. Alternatively, the environmental link can be linked to derived quantities mapped to columns -1 to -4 where each value corresponds to the following quantities:

- * -1; for $\ln(\text{relative spawning biomass})$;
 - * -2; for recruitment deviation;
 - * -3; for $\ln(\text{relative summary biomass})$ (e.g., current year summary biomass divided by the unfished summary biomass);
 - * -4; for $\ln(\text{relative summary numbers})$.
- So, environmental input 103 would use link type 1 (specified in element 9) and apply it to environmental data column 3 and environmental input -103 would use link type 1 and apply it to the "-3" column which is $\ln(\text{relative summary biomass})$.
 - These four derived quantities are all calculated at the beginning of each year within the model, so they are available inside SS to use as the basis for time-varying parameter links without violating any order of operations rule.
- Deviation Link (element 9)
 - 1 = multiplicative ($P_y * = \exp(\text{dev}_y * \text{dev}_{se})$),
 - 2 = additive ($P_y + = \text{env}_y * \text{dev}_{se}$),
 - 3 = random walk options are now implemented by using rho in the objective function. SS now expects the estimated deviations to be normal in distribution and the deviation values are multiplied by the standard error parameter as they are used,
 - 4 = zero reverting random walk with rho. The deviation parameter is now multiplied by the standard error parameter, rather than deviations being penalized according to a specified standard error (the approach in SS v.3.24).
 - The option of applying the final model year deviation value into the forecast period was added in v. 3.30.13. This new option is specified by selecting the appropriate deviation link option (1, 2, 3, or 4) and appending a 2 at the front (21, 22, 23, or 24) which will use the final year deviation value for all forecast years.
 - Deviation Minimum Year (element 10)
 - Year for deviations to start for parameter
 - Deviation Maximum Year (element 11)
 - Year for deviations to end for parameter
 - Deviation Phase (element 12)

- integer, this available element in the long parameter line is now a deviation vector specific phase control
- Blocks (element 13)
 - >0: time block index for parameter.
 - -1: trend with final as offset from base parameter and offset values is in natural log space, also inflection year is in natural log space and the offset from $\ln(0.5)$. No additional parameter lines are required. Three parameters will be estimated; end trend parameter value logistic offset, inflection year logistic offset, and slope.
 - -2: trend with final as standalone value. No additional parameter lines are required. Three parameters will be estimated; end trend parameter value, inflection year, and slope.
 - -3 end value is a fraction of base parameter maximum - minimum; inflection year is fraction of end year - start year. No additional parameter lines are required. Three parameters will be estimated; end trend parameter value as a fraction, inflection year as a fraction, and slope.
- Block Functional Form (element 14)
 - 0: multiplicative parameter ($P_y = P_{base} * \exp(tv_para)$),
 - 1: additive parameter ($P_y = P_{base} + tv_para$),
 - 2: replace parameter ($P_y = tv_para$),
 - 3: random walk across blocks ($P_{block} = P_{block,-1} + tv_para$),
 - 4: mean reverting random walk

When using time-varying parameters through the application of time-blocks (specified using element 13 and 14), it is important to consider which parameters will affect which years of the time series.

There are three main situations:

1. Offset approach: One or more time blocks are created and cover all or a subset of the years. Each block gets a parameter that is used as an offset from the base parameter. In this situation you typically will allow SS to estimate the base parameter and each of the offset parameters. In years not covered by blocks, the base parameter alone is used. However, if blocks cover all the years, then the value of the block parameter is completely correlated with the mean of the block offsets, so model convergence and variance estimation could be affected. The recommended approach when using offsets is to not have all the years be covered by blocks when doing offsets, or to fix the base parameter value at a reasonable level when doing offsets for all years.

2. Replacement approach: Option A: Here time blocks are created which cover a subset of the years. The base parameter is used in the non-block years and the value of that base parameter is replaced by the block parameter in each respective block. In this situation, you typically allow SS to estimate the base parameter and each of the block parameters.
3. Replacement-Option B: Here replacement time blocks are created for all the years. In this case the base parameter is simply a placeholder that is always replaced by a block parameter. In this situation, do not allow SS to estimate the base parameter and only estimate the corresponding block replacement parameters, otherwise, the search algorithm will be attempting to estimate parameters that do not contribute to the log likelihood, so model convergence and variance estimation could be affected. Note however, that the minimum and maximum for the base parameter are used as checks on the minimum and maximum of the blocks.

If a time-varying parameter line is specified requiring an additional parameter a short parameter line will need to be created following the main parameter block section. An example short parameter block would consist of the following:

LO	HI	INIT	Prior Value	Prior SD	Prior Type	Phase	Parameter Label
COND: Only if MG parameters are time-varying							
-99	99	1	0	0.01	0	-1	#Wtlen_1_Fem_ENV_add
-99	99	1	0	0.01	0	-1	#Wtlen_2_Fem_ENV_add

Time-Varying Growth Considerations

When time-varying growth is used, there are some additional considerations to be aware of:

- Growth deviations propagate into the forecast. The user can select which growth parameters get used during the forecast by setting the end year of the last block. If the last block ends in the model's end year, then the growth parameters in effect during the forecast will revert to the "no-block" baseline level. By setting the end year of the last block to end year (endyr) + 1, the model will continue the last block's growth parameter levels throughout the forecast.
- The equilibrium benchmark quantities (MSY, F40%, etc.) previously used end year (endyr) body size-at-age, which is a disequilibrium vector. There is a capability to specify a range of years over which to average the size-at-age used in the benchmark calculations.

Seasonal Biology Parameters

Seasonal effects are available for weight-length parameters, maturity, fecundity, and for the growth parameter K. The seasonal parameter values adjust the base parameter value for that season.

$$P' = P * \exp(\text{seas_value}) \quad (4)$$

Control file continued:

Value	Description
0 0 0 0 0 0 0 0 0 0	Seasonality for selected biology parameters (not a conditional input). Read 10 integers to specify which biology parameters have seasonality: female-wtlen1, female-wtlen2, maturity1, maturity2, fecundity1, fecundity2, male-wtlen1, male-wtlen2, L1, K. Reading a positive value selects that factor for seasonality.
COND: If any factors have seasonality, then read N seasons parameters that define the seasonal offsets from the base parameter value.	
<short parameter line(s)>	Read N seasons short parameter lines for each factor selected for seasonality. The parameter values define an exponential offset from the base parameter value.

Spawner-Recruitment

The spawner-recruitment section starts by specification of the functional relationship that will be used.

Control file continued:

Value	Label	Description
3	Spawner- Recruitment Relationship	The options are: 2: Ricker: 2 parameters: $\ln(R_0)$ and steepness (h); 3: Standard Beverton-Holt, 2 parameters: $\ln(R_0)$ and steepness; 4: Ignore steepness and no bias adjustment. Use this in conjunction with very low emphasis on recruitment deviations to get CAGEAN-like unconstrained recruitment estimates, 2 parameters, but only uses the first one;

Value	Label	Description
		5: Hockey stick:3 parameters: $\ln(R_0)$, steepness, and R_{\min} for $\ln(R_0)$, fraction of virgin SSB at which inflection occurs, and the R level at SSB=0.0; 6: Beverton-Holt with flat-top beyond Bzero, 2 parameters: $\ln(R_0)$ and steepness; 7: Survivorship function: 3 parameters: $\ln(R_0)$, z_{frac} , and β , suitable for sharks and low fecundity stocks to assure recruits are \leq population production; 8: Shepherd re-parameterization: 3 parameters: $\ln(R_0)$, steepness, and shape parameter, c (added to version 3.30.11 and is in beta mode); and 9: Ricker re-parameterization: 3 parameters: $\ln(R_0)$, steepness, and Ricker power, γ (added to version 3.30.11 and is in beta mode).
1	Equilibrium recruitment	Use steepness in initial equilibrium recruitment calculation 0 = none; and 1 = use steepness (h).
0	Future Feature	Reserved for the future option to make realized σ_R a function of the stock-recruitment curve.

Spawner-Recruitment Functions

The number of age-0 fish is related to spawning biomass according to a stock-recruitment relationship. SS has the option of the Beverton-Holt, Ricker, Hockey-Stick, and a survival-based stock recruitment relationship.

Beverton-Holt

The Beverton-Holt Stock Recruitment curve is calculated as:

$$R_y = \frac{4hR_0SB_y}{SB_0(1-h) + SB_y(5h-1)} e^{-0.5b_y\sigma_R^2 + \tilde{R}_y} \quad \tilde{R}_y \sim N(0; \sigma_R^2) \quad (5)$$

where R_0 is the unfished equilibrium recruitment, SB_0 is the unfished equilibrium spawning biomass (corresponding to R_0), SB_y is the spawning biomass at the start of the spawning season during year y , h is the steepness parameter, b_y is the bias adjustment fraction applied during year y , σ_R is the standard deviation among recruitment deviations in natural log space, and \tilde{R}_y is the lognormal

recruitment deviation for year y . The bias-adjustment factor (Methot et al., 2011) ensures unbiased estimation of mean recruitment even during data-poor eras in which the maximum likelihood estimate of the recruitment deviation is near 0.0.

Ricker

The Ricker Stock Recruitment curve is calculated as:

$$R_y = \frac{R_0 S B_y}{S B_0} e^{h(1-SB_y/SB_0)} e^{-0.5b_y\sigma_R^2 + \tilde{R}_y} \quad \tilde{R}_y \sim N(0; \sigma_R^2) \quad (6)$$

where the stock recruitment parameters have the same meaning as described above for the Beverton-Holt.

Hockey-Stick

The hockey-stick recruitment curve is calculated as:

$$R_y = \text{join}(R_{\min} R_0 + R_0 \frac{S B_y}{h S B_0} (1 - R_{\min})) + R_0 (1 - \text{join}) e^{-0.5b_y\sigma_R^2 + \tilde{R}_y} \quad \tilde{R}_y \sim N(0; \sigma_R^2) \quad (7)$$

where R_{\min} is the minimum recruitment level predicted at a spawning size of zero and is set by the user in the control file, h is defined as the fraction of $S B_0$ below which recruitment declines linearly, and join is defined as:

$$\text{join} = \left[1 + e^{1000 * \frac{(S B_0 - h S B_0)}{S B_0}} \right]^{-1} \quad (8)$$

Survivorship

The survivorship stock recruitment relationship based on Taylor et al. (2013) is a stock-recruitment model that enables explicit modeling of survival between embryos and age 0 recruits, and allows the description of a wide range of pre-recruit survival curves. The model is especially useful for low fecundity species that produce relatively few offspring per litter and exhibit a more direct connection between spawning output and recruitment than species generating millions of eggs.

Survival-based recruitment is constrained so that the recruitment rate cannot exceed fecundity. The relationship between survival and spawning output is based on parameters which are on a natural log scale. These are:

$$z_0 = -\ln(S_0) \quad (9)$$

which is the negative of the natural log of the equilibrium survival S_0 , and can be thought of as a pre-recruit instantaneous mortality rate at equilibrium, and

$$z_{\min} = -\ln(S_{\max}) = z_0(1 - z_{\text{frac}}) \quad (10)$$

which is the negative of the natural log of the maximum pre-recruit survival rate (S_{\max} , the limit as spawning output approaches 0), and is parameterized as a function of z_{frac} (which represents the reduction in mortality as a fraction of z_0) so the expression is well defined over the parameter range $0 < z_{\text{frac}} < 1$.

Recruitment at age 0 for each year in the time series is calculated as:

$$R_y = SB_y e^{\left(-z_0 + (z_0 - z_{\min}) \left(1 - (SB_y/SB_0)^\beta\right)\right)} e^{\tilde{R}_y} \quad \tilde{R}_y \sim N(0; \sigma_R^2) \quad (11)$$

where SB_y is the spawning output in year y , β is a parameter controlling the shape of density-dependent relationship between relative spawning depletion SB_y/SB_0 and pre-recruit survival (with limit $\beta < 1$), \tilde{R}_y is the recruitment in year y , and σ_R is the standard deviation of recruitment in natural log space.

As implemented in Stock Synthesis, the parameters needed to apply the stock-recruitment relationship based on the pre-recruit survival are $\ln(R_0)$, z_{frac} , and β . The value of $\ln(R_0)$ defines the equilibrium quantities of SB_0 , S_0 , and z_0 for a given set of biological inputs (either estimated or fixed), regardless of the values of z_{frac} and β .

The interpretation of the quantity $z_0 = -\ln(S_0)$ as pre-recruit instantaneous mortality rate at unfished equilibrium is imperfect because the recruitment in a given year is calculated as a product of the survival fraction S_y and the spawning output SB_y for that same time period so that there is not a 1-year lag between quantification of eggs or pups and recruitment at age 0, which is when recruits are calculated within Stock Synthesis. However, if age 0 or some set of youngest ages is not selected by any fishery or survey, then density dependent survival may be assumed to occur anywhere before the first appearance of any cohort in the data or model expectations. In such cases, the upper limit on survival up to age a is given by $S_{\max} e^{-aM}$.

Nevertheless, interpreting z_0 as an instantaneous mortality helps with the understanding of z_{frac} . This parameter controls the magnitude of the density-dependent increase in survival associated with a reduction in spawning output. It represents the fraction by which this mortality-like rate is reduced as spawning output is reduced from SB_0 to 0. This is approximately equal to the increase in survival as a fraction of the maximum possible increase in survival. That is:

$$z_{\text{frac}} = \frac{\ln(S_{\max}) - \ln(S_0)}{-\ln(S_0)} \approx \frac{S_{\max} - S_0}{1 - S_0} \quad (12)$$

For example, if $S_0 = 0.4$, $z_{\text{frac}} = 0.8$, then the resulting fraction increase in survival is $(S_{\max} - S_0)/(1 - S_0) = 0.72$.

The parameter β controls the point where survival changes fastest as a function of spawning depletion. A value of $\beta = 1$ corresponds to a linear change in natural log survival and an approximately linear relationship between survival and spawning depletion. Values of $\beta < 1$ have survival increasing fastest at low spawning output (concave decreasing survival) whereas $\beta > 1$ has the

increase in survival occurring fastest closer to the unfished equilibrium (convex decreasing survival).

The steepness (h) of the spawner-recruit curve (defined as recruitment relative to R_0 at a spawning depletion level of 0.2) based on pre-recruit survival can be derived from the parameters discussed above according to the relationship and associated inequality:

$$h = 0.2e^{z_0 z_{\text{frac}}(1-0.2^\beta)} < 0.2e^{z_0} = \frac{1}{5S_0} = \frac{SB_0}{5R_0} \quad (13)$$

Unlike the Beverton-Holt stock-recruitment relationship, recruitment can increase above R_0 for stocks that are below SB_0 and thus the steepness is not fundamentally constrained below 1. However, in many cases, steepness will be limited well below 1 by the inequality above, which implies an inverse relationship between the maximum steepness and equilibrium survival. Specifically, the inequality above bounds steepness below 1 for all cases where $S_0 > 0.2$, which are those with the lowest fecundity, an intuitively reasonable result. For example, with $S_0 = 0.4$, the steepness is limited below 0.5, regardless of the choice of z_{frac} or β . This natural limit on steepness may be one of the most valuable aspects of this stock-recruitment relationship.

Code for the survival based recruitment can be found in the code examples section.

Shepherd

The Shepherd stock recruit curve is calculated as:

$$R_y = \left(\frac{SB_y}{SB_0} \right) \frac{5h_{adj}R_0(1-0.2^c)}{(1-5h_{adj}0.2^c) + (5h_{adj}-1)\left(\frac{SB_y}{SB_0}\right)^c} e^{-0.5b_y\sigma_R^2 + \tilde{R}_y} \quad \tilde{R}_y \sim N(0; \sigma_R^2) \quad (14)$$

where c is the shape parameter for the stock recruitment curve, and h_{adj} is the transformed steepness parameter defined as:

$$h_{adj} = 0.2 + \left(\frac{h-0.2}{0.8} \right) \left(\frac{1}{5*0.2^c} - 0.2 \right) \quad (15)$$

More details can be found in Punt and Cope (2019).

Ricker Re-parameterization

The Ricker stock recruit curve re-parameterized version. More details can be found in Punt and Cope (2019).

$$R_y = R_0 * (1 - temp) * e^{\ln(5h)(1-SB_y/SB_0)^\gamma / 0.8^\gamma} \quad (16)$$

where γ is the Ricker shape parameter and $temp$ is defined as:

$$temp = \begin{cases} 1 - SB_y/SB_0 & \text{if } 1 - SB_y/SB_0 > 0 \\ 0.001 & \text{if } 1 - SB_y/SB_0 \leq 0 \end{cases} \quad (17)$$

where $temp$ stabilizes recruitment at R_0 if $SB_y > SB_0$.

Spawner-Recruitment Parameter Setup

Read the required number of long parameter set-up lines (e.g. LO, HI, INIT, PRIOR, PRIOR TYPE, SD, PHASE, ..., and BLOCK TYPE). These parameters are:

Value	Label	Description
8.5	ln(R0)	Log of virgin recruitment level.
0.60	Steepness	Steepness of spawner recruitment relationship, bound by 0.2 and 1.0 for the Beverton-Holt.
COND:	3rd Parameter	If SRR = 5, 7, or 8 Optional depending on which spawner-recruitment relationship function is used.
0.60	σ_R	Standard deviation of natural log recruitment. This parameter has two related roles. It penalizes deviations from the spawner-recruitment curve, and it defines the offset between the arithmetic mean spawner-recruitment curve (as calculated from ln(R0) and steepness) and the expected geometric mean (which is the basis from which the deviations are calculated. Thus the value of σ_R must be selected to approximate the true average recruitment deviation.
0	Regime Parameter	This replaces the R1 offset parameter. It can have a block for the initial equilibrium year, so can fully replicate the functionality of the previous R1 offset approach. The SR regime parameter is intended to have a base value of 0.0 and not be estimated. Similar to cohort-growth deviation, it serves simply as a base for adding time-varying adjustments. This concept is similar to the old environment effect on deviates feature in SS v.3.24 and earlier.
0	Autocorrelation	Autocorrelation in recruitment.

Example set-up of the spawner-recruitment section:

LO	HI	INIT	PRIOR	<other entries>	Block Fxn	Parameter Label
----	----	------	-------	-----------------	-----------	-----------------

3	31	8.81	10.3	...	0	#SR_LN(R0)
0.2	1	0.61	0.70	...	0	#SR_BH_steep
0	2	0.60	0.80	...	0	#SR_sigmaR
-5	5	0	0	...	0	#SR_regime
-99	99	0	0	...	0	#SR_autocorr

The R0, steepness, and regime shift parameters can be time-varying by blocks, trends, environmental linkages, or random deviation. Details on how to specify time-varying parameters can be found in the Time-Varying Parameter Specification and Setup section. However, not all of these options make sense for all parameters and the following items should be considered:

- The σ_R and autocorrelation parameters can not be time-varying.
- The internal estimation of the autocorrelation parameter has been performance tested (Johnson et al., 2016) and was determined to not be well estimated internally. An external estimation for selecting an autocorrelation value is currently recommended.
- In SS v.3.24, the R0 and steepness parameters were used for both the virgin calculation and for the MSY (benchmark) calculations. In SS v.3.30, these usages are more explicitly defined. The value of R0 and steepness in the initial year are used for virgin and for calculation the denominator in depletion estimates. The average value of R0 and steepness in the range of years specified as the benchmark years 9 and 10 (see forecast.ss) is used for MSY type calculations. So, for example, a long-term climate effect could cause R0 to change over time and B_{MSY} could now be calculated for some future range of years.
- The regime shift parameter line allows for multi-year or environmentally driven deviations from R0 without changing R0 itself. Then the recruitment deviations are annual deviations from the current regime and these recruitment deviations can have autocorrelation. This option should be used judiciously because the same algebraic effect on the calculated recruitment can be achieved by different combinations of spawner-recruit parameter options.
- Preliminary recommendation: Use block, trend or environmental effects on R0 only for very long-term and highly persistent effects; use time-vary effects on regime shift for transitory, but multi-year deviations from R0.
- Note that setting a block for just the initial equilibrium year is equivalent to R1_offset. If R1 offset is being used in SS v.3.24, then sstrans.exe will automatically add a new block for the initial equilibrium year so that you can easily re-implement a R1 offset effect.

In SS v.3.24 and before, only a short parameter line was used for the spawner-recruitment section. SS v.3.30 now requires long parameter lines in the spawner-recruitment section because it now uses the same time-varying parameter approach as the biology and selectivity parameters. This generic time-varying approach replaces the spawner recruit envlink concept in SS v.3.24. Also, the R1 offset was effectively a block to implement a regime shift for the initial equilibrium year. Now in v.3.30, the R1 offset parameter is replaced by a parameter termed "regime shift". The spawner recruit regime parameter is intended to have a base value of 0.0 and not be estimated. Similar to the cohort-growth deviation, it serves simply as a base for adding time-varying adjustments. This regime shift parameter can have blocks, environmental links or random deviations.

If a SS v.3.24 model implemented the R1 offset approach this can be mimicked in v.3.30 by the stock recruit regime parameter. For any given year, including startyr - 1, the R0 is replaced by $R0 * \exp(SR_Regime_y)$. So to mimic the R1_offset, you need to put a block on SR_regime for y = startyr - 1. But if SR_regime has some change during your main time series, that change will be filtered through the stock-recruit relationship into an impact on the numbers at age for whatever years are impacted. This can also apply to the forecast. A block on SR_regime is also easier than some of the old dummy environmental variables that were created in the past to adjust recruitment up and down for long periods.

If the R0 or steepness parameters are time-varying, then SS will use the current year's parameters to calculate recruits as a function of the spawning biomass. If the spawner recruit regime parameter is time-varying, then SS applies this offset after calculating recruits as function of spawning biomass. The expected deviations can no longer be linked to "env": instead the environmental effect will be on the "regime".

Recruitment Deviation Setup

Control file continued:

Value	Label	Description
1	Do Recruitment Deviations	<p>This selects the way in which recruitment deviations are coded:</p> <p>0: None (so all recruitments come from spawner recruitment curve).</p> <p>1: Deviation vector (previously the only option): the deviations are encoded as a deviation vector, so ADMB enforces a sum-to-zero constraint.</p> <p>2: Simple deviations: the deviations do not have an explicit constraint to sum to zero, although they still should end up having close to a zero sum. The difference in model performance between options (1) and (2) has not been fully explored to date.</p>

Value	Label	Description
		<p>3: Deviation vector (added in v.3.30.13) where the estimated recruitment is equal to the R0 adjusted for blocks multiplied by a simple deviation vector of unconstrained deviations. The negative log likelihood from the deviation vector is equal to the natural log of the estimated recruitment divided by the expected recruitment by year adjusted for the spawner-recruit curve, regimes, environmental parameters, and bias-adjustment. The negative log likelihood between option 2 and 3 is approximately equal.</p> <p>4: Similar to option 3 but includes a penalty based on the sum of the deviations (added in v.3.30.13).</p>
1971	Main recruitment deviations begin year	If begin year is less than the model start year, then the early deviations are used to modify the initial age composition. However, if set to be more than the population maximum age before start year, it is changed to equal to the maximum age before start year.
1999	Main recruitment deviations end year	If recruitment deviations end year is later than retro year, it is reset to equal retro year.
3	Main recruitment deviations phase.	
1	Advanced Options	<p>0: Use default values for advanced options</p> <p>1: Read values for the 11 advanced options.</p>
COND = 1 Beginning of advanced options		
	1950	<p>Early Recruitment Deviation Start Year:</p> <p>0: skip (default);</p> <p>+year: absolute year (must be less than begin year of main recruitment deviations); and</p> <p>-integer: set relative to main recruitment deviations start year.</p> <p>Note: because this is a deviation vector, it should be long enough so that recruitment deviations for individual years are not unduly constrained.</p>
	6	<p>Early Recruitment Deviation Phase:</p> <p>Negative value: default value to not estimate early deviations.</p> <p>Users may want to set to a late phase if there is not much early data.</p>
	0	<p>Forecast Recruitment Phase:</p> <p>0 = Default value.</p>

Value	Label	Description
		Forecast recruitment deviations always begin in the first year after the end of the main recruitment deviations. Recruitment in the forecast period is deterministic derived from the specified stock-recruitment relationship. Setting their phase to 0 causes their phase to be set to max lambda phase +1 (so that they become active after rest of parameters have converged.). However, it is possible here to set an earlier phase for their estimation, or to set a negative phase to keep the forecast recruitment deviations at a constant level.
1		Forecast Recruitment Deviations Lambda: 1 = Default value. This lambda is for the log likelihood of the forecast recruitment deviations that occur before endyr + 1. Use a larger value here if solitary, noisy data at end of time series cause unruly recruitment deviation estimation.
1956		Last year with no bias adjustment.
1970		First year with full bias adjustment.
2001		Last year with full bias adjustment.
2002		First recent year with no bias adjustment. These four entries control how the bias adjustment is phased in and then phased back out when the model is searching for the maximum log likelihood. Bias adjustment is automatically turned off when in MCMC mode. For intervening years between the first and second years in this list, the amount of bias adjustment that will be applied is linearly phased in. The first year with full bias adjustment should be a few years into the data-rich period so that SS will apply the full bias-correction only to those recruitment deviations that have enough data to inform the model about the full range of recruitment variability. Defaults for the four year values: start year – 1000, start year – Nages, main recruitment deviation final year, end year + 1. See Recruitment Likelihood with Bias Adjustment for more information.
0.85		Max Bias Adjustment: > 0: value for the maximum bias adjustment during the MLE mode. -1: bias adjustment set to 1.0 for all years, including forecast, with estimated recruitment deviations (similar to MCMC). -2: bias adjustment set to 1.0 for all years from start to end model year, bias adjustment set to 0 for the forecast period. -3: bias adjustment set to 0 for all model and forecast years.

Value	Label	Description
0		Period For Recruitment Cycles: Use this when SS is configured to model seasons as years and there is a need to impose a periodicity to the expected recruitment level. If value is >0, then read that number of full parameter lines below define the recruitment cycle.
-5		Minimum Recruitment Deviation: Min value for recruitment deviation. Negative phase = Default value.
5		Maximum Recruitment Deviation: Max value for recruitment deviation. Late Phase = Default value (e.g., 5)
2		Number of Explicit Recruitment Deviations to Read: 0: Default, do not read any recruitment deviations; Integer: read this number of recruitment deviations.
End of advanced options		
COND = If N recruitment cycle is > 0, enter N full parameter lines below.		
	<parameter line>	Full parameter line for each of the N periods of recruitment cycle.
COND = If N explicit recruitment deviations is > 0, then enter N lines below.		
1977	3.0	Enter year and deviation.
1984	3.0	Two example recruitment deviations being read. Note: SS will rescale the entire vector of recruitment deviation after reading these deviations, so by reading two positive values, all other recruitment deviations will be scaled to a small negative value to achieve a sum to zero condition before starting model estimation.

Recruitment Eras

Conceptually, SS treats the early, data-poor period, the main data-rich period, and the recent/forecast time period as three eras along a continuum. The user has control of the break year between eras. Each era has its own vector. The early era is defined as a vector (prior to V3.10 this was a deviation vector) so it can have zeros during the earliest years not informed by data and then a few years with non-zero values without imposing a zero-centering on this collection of deviations. The main era can be a vector of simple deviations, or a deviation vector but it is normally implemented as a deviation vector so that the spawner-recruitment function is its central tendency. The last era does not force a zero-centered deviation vector so it can have zeros during the actual forecast and

non-zero values in last few years of the time series. The early and last eras are optional, but their use can help prevent SS from balancing a preponderance of negative deviations in early years against a preponderance of positive deviations in later years. When the 3 eras are used, it would be typically to turn on the main era during an early model phase, turn on the early era during a later phase, then have the last era turn on in the final phase.

Recruitment Likelihood with Bias Adjustment

For each year in the total recruitment deviation time series (early, mid, late/forecast) the contribution of that year to the log likelihood is equal to: $dev^2 / (2.0 * \sigma_R^2) + offset * \ln(\sigma_R)$ where offset is the recruitment bias adjustment between the arithmetic and geometric mean of expected recruitment for that year. With this approach, years with a zero or small offset value do not contribute to the second component. σ_R may be estimable when there is good data to establish the time series of recruitment deviations.

The recruitment bias adjustment implemented in SS is based upon the work documented in Methot et al. (2011) and following the work of Maunder and Deriso (2003). The concept is based upon the following logic. The σ_R represents the true variability of recruitment in the population. It provides the constraining penalty for the estimates of recruitment deviations and it is not affected by data. Where data that are informative about recruitment deviations are available, the total variability in recruitment, σ_R , is partitioned into a signal (the variability among the recruitment estimates) and the residual, the variance of each recruitment estimate calculated as:

$$SE(\hat{r}_y)^2 + SD(\hat{r})^2 = \left(\left(\frac{1}{\sigma_d^2} + \frac{1}{\sigma_R^2} \right)^{-1/2} \right)^2 + \left(\frac{\sigma_R^2}{(\sigma_R^2 + \sigma_d^2)^{1/2}} \right)^2 = \sigma_R^2 \quad (18)$$

Where there are no data, no signal can be estimated and the individual recruitment deviations collapse towards 0.0 and the variance of each recruitment deviation approaches σ_R . Conversely, where there highly informative data about the recruitment deviations, then the variability among the estimated recruitment deviations will approach σ_R and the variance of each recruitment deviation will approach zero. Perfect data will estimate the recruitment time series signal perfectly. Of course, we never have perfect data so we should always expect the estimated signal (variability among the recruitment deviations) to be less than the true population recruitment variability.

The correct offset (bias adjustment) to apply to the expected value for recruitment is based on the concept that a time series of estimated recruitments should be mean unbiased, not median unbiased, because the biomass of a stock depends upon the cumulative number of recruits, which is dominated by the large recruitments. The degree of offset depends upon the degree of recruitment signal that can be estimated. Where no recruitment signal can be estimated, the median recruitment is the same

as the mean recruitment, so no offset is applied. Where lognormal recruitment signal can be estimated, the mean recruitment will be greater than the median recruitment. The value:

$$b_y = \frac{E\left(SD(\hat{r}_y) \right)^2}{\sigma_R^2} = 1 - \frac{SE(\hat{r}_y)^2}{\sigma_R^2} \quad (19)$$

of the offset then depends upon the partitioning of σ_R into between and within recruitment variability. The most appropriate degree of bias adjustment can be approximated from the relationship among σ_R , recruitment variability (the signal), and recruitment residual error. Because the quantity and quality of data varies during a time series, SS allows the user to control the rate at which the offset is ramped in during the early, data-poor years, and then ramped back to zero for the forecast years. On output SS calculates the mean bias adjustment during the early and main eras and compares it to the RMSE of estimated recruitment deviations in the report.sso file. A warning is generated if the RMSE is small and the bias adjustment is larger than 2.0 times the ratio of $RMSE^2$ to σ_R^2 . Additional information on recruitment bias adjustment can be found here.

In MCMC mode, the model still draws recruitment deviations from the lognormal distribution, so the full offset is used such that the expected mean recruitment from this lognormal distribution will stay equal to the mean from the spawner-recruitment curve. When SS reaches the MCMC and MCEVAL phases, all bias adjustment values are set to 1.0 for all active recruitment deviations because the model is now re-sampling from the full lognormal distribution of each recruitment.

Recruitment Cycle

When SS is configured such that seasons are modeled as years, the concept of season within year disappears. However, there may be reason to still want to model a repeating pattern in expected recruitment to track an actual seasonal cycle in recruitment. If the recruitment cycle factor is set to a positive integer, this value is interpreted as the number of time units in the cycle and this number of full parameter lines will be read. The cyclic effect is modeled as an $\exp(p)$ factor times R_0 , so a parameter value of 0.0 has nil effect. In order to maintain the same number of total recruits over the duration of the cycle, a penalty is introduced so that the cumulative effect of the cycle produces the same number of recruits as $N_{cycles} * R_0$. Because the cyclic factor operates as an exponential, this penalty is different than a penalty that would cause the sum of the cyclic factors to be 0.0. This is done by adding a penalty to the parameter likelihood, where:

$$\begin{aligned} X &= \sum (e^p) \\ Y &= N_{cycle} \\ Penalty &= 100000 * (X - Y)^2 \end{aligned} \quad (20)$$

Initial Age Composition

A non-equilibrium initial age composition is achieved by setting the first year of the recruitment deviations before the model start year. These pre-start year recruitment deviations will be applied to the initial equilibrium age composition to adjust this composition before starting the time series. The model first applies the initial F level to an equilibrium age composition to get a preliminary N-at-age vector and the catch that comes from applying the F's to that vector, then it applies the recruitment deviations for the specified number of younger ages in this vector. If the number of estimated ages in the initial age composition is less than maximum age, then the older ages will retain their equilibrium levels. Because the older ages in the initial age composition will have progressively less information from which to estimate their true deviation, the start of the bias adjustment should be set accordingly.

Fishing Mortality Method

There are three methods available for calculation of fishing mortality (F): Pope's approximation, Baranov's continuous F with each F as a model parameter, and a hybrid F method that does a Pope's approximation to provide initial values for iterative adjustment of the Baranov continuous F values to closely approximate the observed catch. It is recommended in most cases to use the hybrid F approach. With the hybrid method, the final values are in terms of continuous F, but do not need to be specified as full parameters. In a 2 fishery, low F case (e.g., similar to natural mortality or lower), the hybrid method is just as fast as the Pope approximation and produces identical results.

However, when F is very high, the problem becomes quite computationally stiff for Pope's approximation and the hybrid method so convergence in ADMB may slow due to more sensitive gradients in the log likelihood. In these high F cases it may be better to use F option 2, continuous F as full parameters. It is also advisable to allow SS to start with good values for the F parameters. This can be done by specifying a later phase (>1) under the conditional input for F method = 2 where early phases will use the hybrid method, then switch to F as parameter in later phases and transfer the hybrid F values to the parameter initial values.

F as parameter is also preferred for situations where catch is known imprecisely and it is acceptable to have a solution in which the final F values do not reproduce the input catch levels exactly.

Control file continued:

Typical Value	Description and Options
0.2	<p>F ballpark</p> <p>This value is compared to the sum of the F's for the specified year (defined on the next line). The sum is over all seasons and areas. In older versions of SS, the lambda was automatically set to 0.0 in the final phase, the user now has control over whether to reduce the lambda in later phases.</p>

Typical Value	Description and Options
-1990	F ballpark year Negative value disable F ballpark.
3	F Method 1 = Pope's (discrete); 2 = Baranov (continuous) F as a parameter; and 3 = Hybrid F (recommended).
2.9	Maximum F This maximum is applied within each season and area. A value of 0.9 is recommended for F method 1, and a value of about 4 is recommended for F method 2 and 3.
COND: F method = 1, no additional input for Pope's approximation.	
COND: F method = 2:	
0.10 1 0	Initial F value, phase, and the number of F detail setup lines to read (example has 0). For phases prior to the phase of the F value becoming active, SS will use the hybrid option and the F values so calculated become the starting values for the F parameters when this phase is reached.
If F method = 2 and N for F detail is > 0	
1 1980 1 0.20 0.05 4	Fleet, year, season, F, SE, phase - these values override the catch standard error values in the data file and the overall starting F value and phase read just above.
COND: F method = 3	
4	Number of tuning iterations in hybrid method. A value of 2 or 3 is sufficient with a single fleet and low Fs. A value of 5 or so may be needed to match the catch near exactly when there are many fleets and high F.

Initial Fishing Mortality

Read a short parameter setup line for each fishery. The parameters are the fishing mortalities for the initial equilibrium. Do not try to estimate parameters for fisheries with zero initial equilibrium catch.

If there is catch, then give a starting value greater than zero and it generally is best to estimate the parameter in phase 1.

In SS3.30, the initial equilibrium year has explicit seasons, so the needed initial F values will also be by season. Initial F values are only needed for fleet/seasons that have catch. So, if no fleet/season combo has catch, then no parameters are needed.

It is possible to use the initial F method to achieve an estimate of the initial equilibrium Z in cases where the initial equilibrium catch is unknown. To do this requires 2 changes to the input files:

1. Data File: Include a positive value for the initial equilibrium catch for at least one fleet, often with a higher standard error depending upon the situation.
2. Control File: Add an initial F parameter line (short parameter line) for each fleet and season with initial equilibrium catch to be estimated immediately after the Fishing Mortality set-up. It will be influenced by the early age and size comps which should have some information about the early levels of Z.

An example setup with a single fishery fleet with initial equilibrium catches:

0.1	# F ballpark						
-2001	# F ballpark year (negative value to disable)						
3	# F method: 1=Pope; 2=Baranov; 3=Hybrid						
3	# Maximum F value						
4	# Number of iterations for tuning F in hybrid method						
# Initial F parameters							
LO	HI	INIT	Prior	Pr. SD	Pr. Type	Phase	Label
0	3	0.1	0	99	0	1	#InitF Fleet 1

Catchability

Catchability is the scaling factor that relates a model quantity to the expected value for some type of data (index). Typically this is used to converted selected numbers or biomass for a fleet into the expected value for a survey or CPUE by that fleet. In SS, the concept has been extended so that, for example, a time series of an environmental factor could be treated as a survey of the time series of

deviations for some parameter. This flexibility means that a family of link functions beyond simple proportionality is needed.

For each fishery and survey with an index, enter a row with the entries as described below:

1. Fleet Number

2. Link type or index of deviation vector: An assumed functional form between Q , the expected value, and the survey observation.

- (a) 1 = simple Q , proportional assumption about Q : $y = q * x$.
- (b) 2 = mirror simple Q - this will mirror the Q value from another fleet. Mirror in Q must refer to a lower number fleet relative to the fleet with the mirrored Q (example: fleet 3 mirror fleet 2). Requires a Q parameter line for the fleet but will not be used.
- (c) 3 = Q with power, 2 parameters establish a parameter for non-linearity in survey-abundance linkage. Assumes proportional with offset and power function: $y = qx^c$ where $q = \exp(\ln Q_{base})$ thus the c is not related to expected biomass but vulnerable biomass to Q . Therefore, $c < 0$ leads to hyper-stability and $c > 0$ leads to hyper-depletion.
- (d) 4 = mirror Q with offset (2 parameter lines required). The mirrored Q with offset for with be reported as base $Q + \text{offset value}$. Mirror in Q must refer to a lower number fleet relative to the fleet with the mirrored Q . See mirrored Q with offset below for example set up.
- (e) If the parameter is for an index of a deviation vector (index units = 35), use this column to enter the index of the deviation vector to which the index is related.

3. Extra input for link information (i.e., mirror fleet)

- (a) >0 = mirror the Q from another (lower numbered survey designated by referencing the fleet number)
- (b) If a depletion survey, option 34, approach is being applied the following values in this column determines how phases are adjusted:
 - 0 = add 1 to phases of all parameters. Only $R0$ active in new phase 1. Mimics the default option of previous model versions;
 - 1 = only $R0$ active in phase 1. Then finish with no other parameters becoming active; useful for data-limited draws of other fixed parameters. Essentially, this option allows SS to mimic DB-SRA; and

- 2 = no phase adjustments, can be used when profiling on fixed R0.

4. Do extra standard error

- (a) 0 = skip (typical); and
- (b) 1 = estimate a parameter that will contain an additive constant to be added to the input standard deviation of the survey variability.

5. Bias adjustment - adjusts for log-normal bias when using an informative prior on Q.

- (a) 0 = no bias adjustment applied; and
- (b) 1 = apply bias adjustment. Bias correction will be applied to the estimated Q value.

6. Q float

- (a) 0 = no float, parameter is estimated; and
- (b) 1 = float, analytical solution is used, but parameter line still required.
- (c) Additional information regarding the use of and appropriate application of float in Q can be found in the Float Q section below.

For a setup with a single survey, the Q setup matrix could be:

Fleet Num.	Link Type	Link Info	Extra SD	Bias Adjust	Float	Label
3	1	0	1	1	0	#Survey
-9999	0	0	0	0	0	#End Read

LO	HI	INIT	<other entries>	PHASE	<other entries>	Block Fxn	Parameter Label
-5	5	-0.12	...	1	...	0	#Survey1 LnQ base
0	0.5	0.1	...	-1	...	0	#Survey1 Extra SD

If the Q base parameter specifies that it is time-varying by the annual deviation method, short parameter lines to specify the specifications of the deviation vector come after all the base Q parameters.

Mirrored Q with offset

Below is an example setup for fleets with a mirrored Q and offset from another fleet (link type = 4):

For a setup with a single survey, the Q setup matrix could be:

Fleet Num.	Link Type	Link Info	Extra SD	Bias Adjust	Float	Label
1	1	0	1	0	0	#Fleet 1
2	4	1	0	0	0	#Fleet 2
-9999	0	0	0	0	0	#End Read

A long parameter line is expected for each link parameter (i.e. Q) and for the

LO	HI	INIT	<other entries>	PHASE	<other entries>	Block Fxn	Parameter Label
-7	5	0.51	...	1	...	0	#Fleet 1 LnQ base
0	0.5	0.1	...	-1	...	0	#Fleet 1 Extra SD
-7	5	-6	...	-1	...	0	#Fleet 2 LnQ base
-8	5	-7	...	-1	...	0	#Fleet 2 Mirror Q offset

Float Q

The use and development of float in Q has evolved over the life of SS. The original approach in earlier versions of SS (version 3.24 and before) is that with Q "float" the units of the survey or fishery CPUE were treated as dimensionless so the Q was adjusted within each model iteration to maintain a mean difference of 0.0 between observed and expected (usually in natural log space). In contrast, Q as a parameter (float = 0) one had the ability to interpret the absolute scaling of Q and put a prior on it to help guide the model solution. Also, with Q as a parameter the code allowed for Q to be time-varying.

Then midway through the evolution of the SS v.3.24 code lineage a new Q option was introduced based on user recommendations. This option allowed Q to float and to compare the resulting Q value to a prior, hence the information in that prior would pull the model solution in direction of a floated Q that came close to the prior.

Currently, in 3.30, that float with prior capability is fully embraced. All fleets that have any survey or CPUE options need to have a catchability specification and get a base Q parameter in the list. Any of these Q's can be either:

- Fixed: by not floating and not estimating.
- Floating: not estimated as an active parameter, set phase to negative, and not capable of being time-varying. Can have a prior, or not. Future versions of SS may allow Q to be time-varying and then rescaling that Q vector according to the float logic but this is not yet currently implemented.
- Estimated as active parameter: so not floating. Can be time-varying and can have a prior.

Q relates the units of the survey or CPUE to the population abundance, not the population density per unit area. But many surveys and most fishery CPUE is a proportional to mean fish density per unit area. This does not have any impact in a one area model because the role of area is absorbed into the value of Q. In a multi-area model, one may want to assert that the relative difference in CPUE between two areas is informative about the relative abundance between the areas. However, CPUE is a measure of fish density per unit area, so one may want to multiply CPUE by area before putting the data into SS so that asserting the same Q for the two areas will be informative about relative abundance.

In SS v.3.30.13, a new catchability option has been added that allows Q to be mirrored and to add an offset to $\ln(Q)$ of the primary area when calculating the $\ln(Q)$ for the dependent area. The offset is a parameter and, hence, can be estimated and have a prior. This option allows the CPUE data to stay in density units and the effect of relative stock area is contained in the value of the $\ln(Q)$ offset.

Q Conversion Issues Between SS v.3.24 and v.3.30

In SS v.3.24 it was common to use the deviation approach implemented as if it was survey specific blocks to create a time-varying Q for a single survey. In some cases, only one year's deviation was made active in order to implement, in effect, a block for Q. The transition executable (sstrans.exe) cannot convert this, but an analogous approach is available in SS v.3.30 because true blocks can now be used, as well as environmental links and annual deviations. Also note that deviations in SS v.3.24 were survey specific (so no parameter for years with no survey). In SS v.3.30, deviations are always year-specific, so you might have a deviation created for a year with no survey.

Selectivity and Discard

For each fleet and survey, read a definition line for size selectivity and retention.

Example Setup for Selectivity:

Size Selectivity:

Pattern	Discard	Male	Special	Label
1	2	0	0	#Fishery1
1	0	0	0	#Survey1
0	0	0	0	#Survey2

Age Selectivity:

Pattern	Discard	Male	Special	Label
11	0	0	0	#Fishery1
11	0	0	0	#Survey1
11	0	0	0	#Survey2

Pattern

Specify the size/age selectivity pattern. See the Selectivity Pattern section for user options.

Discard

Discard options:

- Option = 0: no discarding by fleet.
- Option = 1: program will read 4 retention parameters after reading the specified number of selectivity parameters and all discarded fish are assumed dead.
- Option = 2: program will read 4 retention parameters and 4 discard mortality parameters.
- Option = 3: no additional parameters are read and all fish are assumed discarded and dead.
- Option = 4: program will read 7 retention parameters for dome-shaped retention and 4 discard mortality parameters.
- Option = negative fleet number: will mirror the retention and discard mortality pattern of the lower numbered fleet.

Male

Male specific selectivity options:

- Option = 0: no male specific selectivity parameters required. Male and female selectivity will be the same.
- Option = 1: program will read 4 additional parameters to define the male selectivity relative to the female selectivity. Anytime the male selectivity is caused to be greater than 1.0; the entire male/female matrix of selectivity values is scaled by the max so that the realized max is 1.0. Note, this may cause gradient problems.
- Option 2: main selectivity parameters define male selectivity and female selectivity is estimated as an offset from male selectivity. This alternative is preferable if female selectivity is less than male selectivity.
- Option 3: only available if the selectivity pattern is 1, 20, or 24 and it causes the male selectivity parameters to be offset from the female parameters, rather than the male selectivity being an offset from the female selectivity.

Special

Special (0/value): This value is used in different ways depending on the context. If the selectivity type is to mirror another selectivity type, then put the index of that source fleet or survey here. It must refer to a lower numbered fleet/survey. If the selectivity type is 6 (linear segment), then put the number of segments here. If the selectivity type is 7, then put a 1 here to keep selectivity constant above the mean average size for old fish of morph 1.

Age Selectivity

For each fleet and survey, read a definition line for age selectivity. The 4 values to be read are the same as for the size-selectivity.

As of SS v. 3.30.15, for some selectivity patterns the user can specify the minimum age of selected fish. Most selectivity curves by default select age 0 fish (i.e., inherently specify the minimum age of selected fish as 0). However, it is fairly common for the age bins specified in the data file to start at age 1. This means that any age 0 fish selected are pooled up into the age 1' bin, which will have a detrimental effect on fitting age-composition data. In order to prevent the selection of age 0 (or older) fish, the user can specify the minimum selected age for some selectivity patterns (12, 13, 14, 16, 18, 26, or 27) in versions of SS 3.30.15 and later. For example, if the minimum selected age is 1 (so that age 0 fish are not selected), selectivity pattern type can be specified as 1XX, where XX is the selectivity pattern. A more specific example is if selectivity is age-logistic and the minimum selected age desired is 1, the selectivity pattern would be specified as 112 (the regular age-logistic selectivity pattern is option 12). The user can also select higher minimum selected ages, if desired; for example, 212 would be the age-logistic selectivity pattern with a minimum selected age of 2 (so that age 0 and 1 fish are not selected).

Reading the Selectivity and Retention Parameters

Read the required number of parameter setup lines as specified by the definition lines above. The complete order of the parameter setup lines is:

1. Size selectivity for fishery 1;
2. Retention for fishery 1 (if discard specified);
3. Discard Mortality for fishery 1 (if discard specified);
4. Male offsets for size selectivity for fishery 1 (if offsets used);
5. <repeat for additional fleets and surveys>;
6. Age selectivity for fishery 1;
7. Retention for fishery 1 (if discard specified);
8. Discard Mortality for fishery 1 (if discard specified);
9. Male offsets for age selectivity for fishery 1 (if offsets used); and
10. <repeat for additional fleets and surveys>.

The list of parameters to be read from the above setup would be:

LO	HI	INIT	PRIOR	<other entries>	Block Fxn	Parameter Label
19	80	53.5	50	...	0	#SizeSel p1 fishery 1
0.01	60	18.9	15	...	0	#SizeSel p2 fishery 1
20	70	38.6	40	...	0	#Retain p1 fishery 1
0.1	10	6.5	1	...	0	#Retain p2 fishery 1
0.001	1	0.98	1	...	0	#Retain p3 fishery 1
-10	10	1	0	...	0	#Retain p4 fishery 1
0.1	1	0.6	0.6	...	0	#DiscMort p1 fishery 1
-2	2	0	0	...	0	#DiscMort p2 fishery 1
20	70	40	40	...	0	#DiscMort p3 fishery 1
0.1	10	1	1	...	0	#DiscMort p4 fishery 1
19	80	53.5	50	...	0	#SizeSel p1 survey 1
0.01	60	18.9	15	...	0	#SizeSel p2 survey 1
0	40	0	5	...	0	#AgeSel p1 fishery 1
0	40	40	5	...	0	#AgeSel p2 fishery 1

LO	HI	INIT	PRIOR	<other entries>	Block Fxn	Parameter Label
0	40	0	5	...	0	#AgeSel p1 survey 1
0	40	40	5	...	0	#AgeSel p2 survey 1
0	40	0	5	...	0	#AgeSel p1 survey 2
0	40	0	5	...	0	#AgeSel p2 survey 2

Selectivity Patterns

The currently defined selectivity patterns, and corresponding required number of parameters, are:

SIZE BASED SELECTIVITY		
Pattern	N Parameters	Description
0	0	Selectivity = 1.0 for all sizes.
1	2	Logistic selectivity.
2	8	Discontinued: Double logistic with defined peak (uses IF joiners). Use pattern 8 instead.
4	0	Discontinued: Set size selectivity equal to female fecundity. Use special survey pattern 30 located in the index section within the data file instead.
5	2	Mirror another selectivity. The two parameters select bin range.
6	2 + special	Non-parametric
7	8	Discontinued: Double logistic with defined peak, uses smooth joiners; special = 1 causes constant selectivity above L_{inf} for morph 1. Use pattern 8.
8	8	Double logistic, with defined peak, uses smooth joiners; special=1 causes constant selectivity above L_{inf} for morph 1.
9	6	Simple double logistic with no defined peak.
15	0	Mirror another selectivity (same as for age selectivity).
21	2	Non-parametric size selectivity
22	4	Double normal; similar to CASAL.

Pattern	N Parameters	Description
23	6	Same as the double normal pattern 24 except the final selectivity is now directly interpreted as the terminal selectivity value. Cannot be used with Pope's F method because maximum selectivity may be greater than 1.
24	6	Double normal with defined initial and final selectivity level – Recommended option.
25	3	Exponential-logistic.
27	3 + 2*N nodes	Cubic spline with N nodes.
42	5 + 2*x	Selectivity pattern 27 but with 2 additional scaling parameters.
43	4 + x	Selectivity pattern 6 but with 2 additional scaling parameters.

AGE BASED SELECTIVITY

Pattern	N Parameters	Description
0	0	Selectivity = 1.0 for ages 0+.
10	0	Selectivity = 1.0 for all ages beginning at age 1. If it is desired that age-0 fish be selected, then use pattern #11 and set minimum age to 0.0.
11	2	Selectivity = 1.0 for a specified age range.
12	2	Logistic selectivity.
13	8	Double logistic, IF joiners. Use discouraged. Use pattern 18 instead.
14	nages + 1	Separate parameter for each age (empirical), value at age is $\frac{1}{1+exp(-x)}$.
15	0	Mirror another age-specific selectivity pattern.
16	2	Coleraine single Gaussian
17	nages + 1	Empirical as a random walk from previous age. For all ages in the population beginning with Amin = 1 for the fishery and 2 for the survey, there is a corresponding set of selectivity parameters for each fleet, p_a .
18	8	Double logistic, with defined peak, uses smooth joiners.

Pattern	N Parameters	Description
19	6	Simple double logistic with no defined peak.
20	6	Double normal with defined initial and final level. Recommended option.
26	3	Exponential logistic.
27	$3 + 2 \cdot N$ nodes	Cubic spline in age based on N nodes.
41	$2 + \text{nages} + 1$	Selectivity pattern 17 but with 2 additional scaling parameters.
42	$5 + 2 \cdot N$ nodes	Selectivity pattern 27 but with 2 additional scaling parameters.
44	$4 + \text{nages}$	Similar to age selectivity pattern 17 but with separate parameters for males and with revised controls.
45	$4 + \text{nages}$	Similar to age selectivity pattern 14 but with separate parameters for males and with revised controls.

Special Selectivity Options

Special selectivity options (type 30 in size based selectivity) are no longer specified within the control file. Specifying the use of one of these selectivity types is now done within the data file by selecting the survey "units" (see the section on Index units).

Selectivity Pattern Details

Pattern 1 (size) and 12 (age) - Simple Logistic

Within SS logistic selectivity for the primary sex (if selectivity varies by sex) is formulated as:

$$S_l = \frac{1.0}{1 + \exp(-\ln(19)(L_l - p1)/p2)} \quad (21)$$

where L_l is the length bin. If age based selectivity is selected then the length bin is replaced by the age vector. If sex specific selectivity is specified the non-primary sex the p1 and p2 parameters are estimated as offsets. Note that with a large p2 parameter, selectivity may not reach 1.0 at the largest size bin. The parameters are:

- p1 - size/age at inflection; and
- p2 - width for 95% selection; a negative width causes a descending curve.

Pattern 5 (size) - Mirror Selectivity

Two parameters select the min and max bin number (not min max size) of the source selectivity pattern. If first parameter has value ≤ 0 , then interpreted as a value of 1 (e.g., first bin). If second parameter has value ≤ 0 , then interpreted as maximum length bin (e.g., last bin specified in the data file). The mirrored selectivity pattern must have be from a lower fleet number (e.g., already specified before the mirrored fleet).

Pattern 6 (size) - Non-parametric Selectivity

Non-parametric size selectivity uses a set of linear segments. The first way point is at Length = p1 and the last way point is at Length = p2. The total number of way points is specified by the value of the Special factor in the selectivity set-up, so the N intervals is one less than the number of way points. Intermediate way points are located at equidistant intervals between p1 and p2. Parameters 3 to N are the selectivity values at the way points, entered as logistic, e.g. $1/(1 + \exp(-x))$. Ramps from 10 cm to p3 if $L < p1$. Constant at Np if $L > p2$. Note that prior to version 3.03 the way points were specified in terms of bin number, rather than length.

Pattern 8 (size) and 18 (age) - Double Logistic

Users are discouraged from using the double logistic selectivity. The double normal selectivity pattern (size pattern 24, age pattern 20) provides similar functionality but with only 6 parameters.

- p1 – peak: size (age) for peak. Should be an integer and should be at bin boundary and not estimated. But options 7 and 18 may allow estimation.

- p2 – initial: selectivity at length bin=1 (minimum length) or age=0.
- p3 – inflection: a logit transform $(1/(1 + \exp(-x)))$ is used so that the transformed value will be between 0 and 1. So a p1 value of -1.1 will be transformed to 0.25 and used to set the selectivity equal to 0.5 at a size (age) equal to 0.25 of the way between minimum length and peak.
- p4 – slope1: $\ln(\text{slope})$ of left side (ascending) selectivity.
- p5 – final: logit transform for selectivity at maximum length (or maximum age).
- p6 – inflection2: logit transform for size (age) at right side selectivity equal to half way between peak + peak width and maximum length (or maximum age).
- p7 – slop2: $\ln(\text{slope})$ of right side (descending) selectivity.
- p8 – peak width: width of flattop.

Pattern 14 (age) - Revise Age

Age-selectivity pattern 14 to allow selectivity-at-age to be the same as selectivity at the next younger age. When using this option, the range on each parameter should be approximately -5 to 9 to prevent the parameters from drifting into extreme values with nil gradient. SS calculates the age-based selectivity as where $a = 1$ to $a = Amax + 1$:

$$S_a = \frac{1}{1 + \exp(-(p_{a+1} + temp))} \quad (22)$$

$$temp = 9 - \max(p_a)$$

Pattern 17 (age) - Random Walk

This selectivity pattern provides for a random walk in $\ln(\text{selectivity})$. For each age $a \geq A_{\min}$, where A_{\min} is the minimum age for which selectivity is allowed to be non-zero, there is a selectivity parameter, p_a , controlling the changing selectivity from age $a - 1$ to age a .

The selectivity at age a is computed as

$$S_a = \exp(S'_a - S'_{\max}), \quad (23)$$

where

$$S'_a = \sum_{i=a_{\min}}^A p_i \quad (24)$$

and

$$S'_{\max} = \max\{S'_a\}. \quad (25)$$

Selectivity is fixed at $S_a = 0$ for $a < A_{\min}$.

This formulation has the properties that the maximum selectivity equals 1, positive values of p_a are associated with increasing selectivity between ages $a - 1$ and a , and negative values are associated with decreasing selectivity between those ages and $p_a = 0$ gives constant selectivity.

The condition that maximum selectivity equals 1 results in one fewer degree of freedom than the number of estimated p_a . Therefore, at least one parameter should be fixed at an arbitrary value, typically $p_{A_{\min}} = 0$.

In typical usage:

- First parameter (for age 0) could have a value of -1000 so that the age 0 fish would get a selectivity of 0.0;
- Second parameter (for age 1) could have a value of 0.0 and not be estimated, so age 1 is the reference age against which subsequent changes occur.
- Next parameters get estimated values. To assure that selectivity increases for the younger ages, the parameter min for these parameters could be set to 0.0 or a slightly negative value.
- If dome-shaped selectivity is expected, then the parameters for older ages could have a range with the max set to 0.0 so they cannot increase further.
- To keep selectivity at a particular age the same as selectivity at the next younger age, set its parameter value to 0.0 and not estimated. This allows for all older ages to have the same selectivity.
- To keep a constant rate of change in selectivity across a range of ages, use the -999 flag to keep the same rate of change in $\ln(\text{selectivity})$ as for the previous age.
- Code for implementing random walk selectivity within SS can be found in the code example section.

Pattern 9 (size) and 19 (age) - Simple Double Logistic with no defined peak

- p1 - inflection1: ascending inflection size (in cm).
- p2 – slope1: ascending slope.
- p3 – inflection2: descending inflection size (in cm).
- p4 – slope2: descending slope.

- p5 – first bin: bin number for the first bin with non-zero selectivity (must be an integer bin number, not a size).
- p6 – offset: enter 0 if p3 is independent of p1; enter 1 if p3 is an offset from p1.

Pattern 22 (size) - Double Normal with Plateau

- p1 – peak1: beginning size for the plateau (in cm).
- p2 – peak2: ending size for the plateau. Calculated as a fraction of the distance between peak1 and 99% of the lower edge of the last size bin in the model. Transformed as $(1/(1+\exp(-p2)))$. So a value of 0 results in peak2 being halfway between peak1 and 99% of the last bin.
- p3 – upslope: $\ln(\text{variance})$ on ascending side.
- p4 – downslope: $\ln(\text{variance})$ on descending side.

Pattern 23 (size), 24 (size), and 20 (age) - Double Normal Selectivity

- p1 – peak: beginning size (or age) for the plateau (in cm or year).
- p2 – top: width of plateau, as logistic between peak and maximum length (or age).
- p3 – ascending width: parameter value is $\ln(\text{width})$.
- p4 – descending width: parameter value is $\ln(\text{width})$.
- p5 – initial: selectivity at first bin, as logistic between 0 and 1.
- p6 – final: selectivity at last bin, as logistic between 0 and 1 (for pattern 24). For pattern 23 age selectivity at last bin, as absolute value, so can be > 1.0 . Warning: do not allow this value to go above 1.0 if the F_method uses Pope's approximation. It can above 1.0 when F is in exponential form. When this parameter is above 1.0, the overall selectivity pattern will have an intermediate plateau at 1.0 (according to peak and top), then will ascend further to the final value.

Notes for Double Normal Selectivity:

- For the initial selectivity parameter (parameter 5)
 - -999 or -1000: ignore the initial selectivity algorithm and simply decay the small fish selectivity according to p3,

- < -1000: ignore the initial selectivity algorithm as above and then set selectivity equal to $1.0e-06$ for size bins 1 through bin = -1001 –value. So a value of –1003 would set selectivity to a nil level for bins 1 through 2 and begin using the modeled selectivity in bin 3.
- For the final selectivity parameter (parameter 6)
 - -999 or –1000: ignore the final selectivity algorithm and simply decay the large fish selectivity according to parameter 4,
 - <-1000: set selectivity constant for bins greater than bin number = -1000 – value.

Pattern 15 (size or age) - Mirror

No parameters. Whole age range is mirrored from another fleet. The mirrored selectivity pattern must have be from a lower fleet number (e.g., already specified before the mirrored fleet).

Pattern 16 (age) - Gaussian (similar to Coleraine)

- p1 – age below which selectivity declines
- p2 – scaling factor for decline

Pattern 9 (size) and 19 (age) - Simple Double Logistic

- p1 – ascending inflection age/size
- p2 – ascending slope
- p3 – descending inflection age/size
- p4 – descending slope
- p5 – age or size at first selection; this is a specification parameter, so must not be estimated. Enter integer that is age for pattern 19 and is bin number for pattern 9
- p6 – (0/1) where a value of 0 causes the descending inflection to be a standalone parameter, and a value of 1 causes the descending inflection to be interpreted as an offset from the ascending inflection. This is a specification parameter, so must not be estimated.

A value of $1.0e-6$ is added to the selectivity for all ages, even those below the minimum age.

Pattern 25 (size) and 26 (age) - Exponential logistic

- p1 – ascending rate, min: 0.02, max: 1.0, reasonable start value: 0.1.
- p2 – peak, as fraction of way between min size and max size. Parameter min value: 0.01; max: 0.99; reasonable start value: 0.5.
- p3 – descending rate, min: 0.001, max: 0.5, reasonable start value: 0.01. A value of 0.001 provides a nearly asymptotic curve. Values above 0.2 provide strongly dome-shaped function in which the p3 and p1 parameters interact strongly.

The exponential logistic selectivity is calculated as:

$$peak = \min(L_l) + p2(\max(L_l) - \min(L_l)) \quad (26)$$

where L_l is the length bins at bin l (if age-based substitute with age bins) and:

$$S_l = \frac{e^{p3 \cdot p1 \cdot (peak - L_l)}}{1 - p3(1 - e^{p1 \cdot (peak - L_l)})} \quad (27)$$

Pattern 27 (size or age)- Cubic Spline

This selectivity pattern uses the ADMB implementation of the cubic spline function. This function requires input of the number of nodes, the positions of those nodes, the parameter values at those nodes, and the slope of the function at the first and last node. In SS, the number of nodes is specified in the "special" column of the selectivity set-up. The pattern number 27 is used to invoke cubic spline for size selectivity and for age selectivity; the input syntax is identical.

For a 3 node setup, the SS input parameters would be:

- p1 – code for initial set-up (0, 1 or 2) as explained below
- p2 – gradient at the first node (should be a small positive value)
- p3 – gradient at the last node (should be zero or a small negative value)
- p4-p6 – the nodes in units of cm; must be in rank order and inside of the range of the population length bins. These must be held constant (not estimated, e.g., negative phase value) during a model run.
- p7-p9 – the values at the nodes. Units are $\ln(\text{selectivity})$.

Notes:

- There must be at least 3 nodes.

- One of the node selectivity parameter values should be held constant so others are estimated relative to it.
- Selectivity is forced to be constant for sizes greater than the size at the last node.
- The overall selectivity curve is scaled to have a peak equal to 1.0.
- Terminal nodes cannot be at the min or max population length bins.

Code for implementing cubic spline selectivity within SS can be found in code examples.

Auto-Generation of Cubic Spline Control File Set-Up: A New SS feature pioneered with the cubic spline function is a capability to produce more specific parameter labels and to auto-generate selectivity parameter setup. The auto-generation feature is controlled by the first selectivity parameter value for each fleet that is specified to use the cubic spline. There are 3 possible values for this setup parameter:

- 0: no auto-generation, process parameter setup as read.
- 1: auto-generate the node locations based on the specified number of nodes and on the cumulative size distribution of the data for this fleet/survey.
- 2: auto-generate the nodes and also the min, max, prior, initial, and phase for each parameter.

With either the auto-generate option #1 or #2, it still is necessary to include in the parameter file placeholder rows of values so that the `init_matrix` command can input the current number of values because all selectivity parameter lines are read as a single matrix dimensioned as N parameters x 14 columns. The read values of min, max, initial, prior, prior type, prior stddev, and phase will be overwritten.

Cumulative size and age distribution is calculated for each fleet, summing across all samples and both sexes. These distributions are output in `echoinput.sso` and in a new `OVERALL_COMPS` section of `report.sso`.

When the nodes are auto-generated, the first node is placed at the size corresponding to the 2.5% percentile of the cumulative size distribution, the last is placed at the 97.5% percentile of the size distribution, and the remainder are placed at equally spaced percentiles along the cumulative size distribution. These calculated node values are output into `control.ss_new`. So, the user could extract these nodes from `control.ss_new`, edit them to desired values, then, insert them into the input control file. Remember to turn off auto-generation in the revised control file.

When the complete auto-generation is selected, the `control.ss_new` would look like the table below:

LO	HI	INIT	<other entries>	Block Fxn	Parameter Label
0	2	2.0	...	0	#SizeSpline Code
-0.001	1	0.13	...	0	#SizeSpline GradLo
-1	0.001	-0.03	...	0	#SizeSpline GradHi
11	95	38	...	0	#SizeSpline Knot1
11	95	59	...	0	#SizeSpline Knot2
11	95	74	...	0	#SizeSpline Knot3
-9	7	-3	...	0	#SizeSpline Value1
-9	7	-1	...	0	#SizeSpline Value2
-9	7	-0.78	...	0	#SizeSpline Value3

Pattern 41 (age) - Random Walk with User-Defined Scaling

Selectivity pattern 17 with two additional parameters. The two additional parameters are the bin numbers to define the range of bins for scaling. All of the selectivity values will be scaled (divided) by the mean value over this range. The low and high bin numbers are defined before the other selectivity parameters.

LO	HI	INIT	<other entries>	Block Fxn	Parameter Label
0	20	10	...	0	#AgeSel_ScaleAgeLo
0	20	20	...	0	#AgeSel_ScaleAgeHi

Pattern 42 (size or age) - Cubic Spline with User-Defined Scaling

Selectivity pattern 27 with two additional parameters. The two additional parameters are the bin numbers to define the range of bins for scaling. All of the selectivity values will be scaled (divided) by the mean value over this range. The low and high bin numbers are defined before the other selectivity parameters.

LO	HI	INIT	<other entries>	Block Fxn	Parameter Label
0	20	10	...	0	#AgeSpline_ScaleAgeLo
0	20	20	...	0	#AgeSpline_ScaleAgeHi

Pattern 43 (size) - Non-parametric with User-Defined Scaling

Selectivity pattern 6 with two additional parameters. The two additional parameters are the bin numbers to define the range of bins for scaling. All of the selectivity values will be scaled (divided) by the mean value over this range. The low and high bin numbers are defined before the other selectivity parameters.

LO	HI	INIT	<other entries>	Block Fxn	Parameter Label
1	80	50	...	0	#SizeSel_ScaleBinLo
1	80	70	...	0	#SizeSel_ScaleBinHi

Pattern 44 (age)

Similar to pattern 17 but with separate parameters for males and females. This selectivity pattern provides for a random walk in $\ln(\text{selectivity})$. In typical usage:

- p1 - First parameter (for age 0) could have a value of -1000 so that the age 0 fish would get a selectivity of 0.0.
- p2 - The first age for which mean selectivity = 1.
- p3 - The last age for which mean selectivity = 1.
- p4 - Male mean selectivity relative to the female selectivity mean entered as $\ln(\text{ratio})$ for the male relative female selectivity.
- p5-p_n - Additional parameter lines for the natural log of the selectivity change between ages corresponding to the user specified number of changes in the "special" column for the selectivity specification for each sex with females entered first then males.
- -999 input indicates to the model to keep the change unchanged from the previous age (keeps same rate of change).
- -1000 input indicates used only for male selectivity indicates to the model to set the change in male selectivity equal to the female change in selectivity.

An example specification and setup for this selectivity option where selectivity is dome-shaped, peaking at age 2 with female and male selectivity are equal with 4 change points per sex:

#Pattern	Discard	Male	Special
44	0	0	4

LO	HI	INIT	<other entries>	Block Fxn	Parameter Label
0	20	0	...	0	#first selex age
0	20	2	...	0	#first age peak selex (mean)
0	20	2	...	0	#last age peak selex (mean)
-1	2	-0.001	...	0	#male ln(ratio)
-10	10	3.01	...	0	#female ln(selex) change 1
-10	10	1.56	...	0	#female ln(selex) change 2
-10	10	-0.15	...	0	#female ln(selex) change 3
-10	10	-0.15	...	0	#female ln(selex) change 4
-1000	10	-1000	...	0	#male ln(selex) change 1
-1000	10	-1000	...	0	#male ln(selex) change 2
-1000	10	-1000	...	0	#male ln(selex) change 3
-1000	10	-1000	...	0	#male ln(selex) change 4

Pattern 45 (age) - Revise Age

Similar to pattern 14 but with separate parameters for males and females. Age-selectivity pattern 45 to allow selectivity-at-age to be the same as selectivity at the next younger age.

- p1 - is the first age with non-zero selectivity.
- p2 - The first age in mean for peak selectivity
- p3 - The last age in mean for peak selectivity
- p4 - The male mean selectivity relative to the female mean, entered as ln(ratio) for the male relative female selectivity
- -999 input indicates to the model to keep the change unchanged from the previous age (keeps same rate of change).
- -1000 input indicates used only for male selectivity indicates to the model to set the change in male selectivity equal to the female change in selectivity.

An example specification and setup for this selectivity option where selectivity is asymptotic, with female and male selectivity are equal with 4 change points per sex:

#Pattern	Discard	Male	Special
----------	---------	------	---------

45	0	0	3
----	---	---	---

LO	HI	INIT	<other entries>	Block Fxn	Parameter Label
0	20	2	...	0	#first selex age
0	20	5	...	0	#first age peak selex (mean)
0	20	5	...	0	#last age peak selex (mean)
-1	2	-0.001	...	0	#male ln(ratio)
-10	10	-8.1	...	0	#female ln(selex) change 1
-10	10	-4.1	...	0	#female ln(selex) change 2
-10	10	-1.8	...	0	#female ln(selex) change 3
-1000	10	-1000	...	0	#male ln(selex) change 1
-1000	10	-1000	...	0	#male ln(selex) change 2
-1000	10	-1000	...	0	#male ln(selex) change 3

Retention

Retention is defined as a logistic function of size or age. It does not apply to surveys. Four parameters (for asymptotic retention) or seven parameters (for dome-shaped retention) are used:

- Asymptotic (4 parameters):
 - p1 – ascending inflection,
 - p2 – ascending slope,
 - p3 – maximum retention controlling the height of the asymptote (smaller values result in lower asymptotes), often a time-varying quantity to match the observed amount of discard. As of v. 3.30.01, this parameter is now input in logit space ranging between -10 and 10. A fixed value of -999 would assume no retention of fish and a value of 999 would set asymptotic retention equal to 1.0,
 - p4 – male offset to ascending inflection (arithmetic, not multiplicative),
- Dome-shaped (add the following 3 parameters):

- p5 – descending inflection,
- p6 – descending slope, and
- p7 – male offset to descending inflection (arithmetic, not multiplicative).

$$\text{Retention}_l = \left(\frac{P3}{1 + e^{\frac{-(L_l - (P1 + P4 * male))}{P2}}} \right) * \left(1 - \frac{1}{1 + e^{\frac{-(L_l - (P5 + P7 * male))}{P6}}} \right) \quad (28)$$

Discard Mortality

Discard mortality is defined as a logistic function of size such that mortality declines from 1.0 to an asymptotic level as fish get larger. It does not apply to surveys and it does not affect the calculation of expected values for discard data. It is applied so that the total mortality rate is:

$$\text{dead fish} = \text{selectivity} * (\text{retain} + (1.0 - \text{retain}) * \text{discard mortality})$$

If discard mortality is 1.0, all selected fish are dead; if discard mortality is 0.0, only the retained fish are dead.

Four parameters are used:

- p1 – descending inflection
- p2 – descending slope
- p3 – maximum discard mortality
- p4 – male offset to descending inflection (arithmetic, not multiplicative)

Discard mortality is calculated as:

$$\text{Discard Mortality}_l = 1 - \frac{1 - P3}{1 + e^{\frac{-(L_l - (P1 + P4 * male))}{P2}}} \quad (29)$$

Sex-Specific Selectivity

There are two approaches to specifying sex specific selectivity. One approach allows male selectivity to be specified as a fraction of female selectivity (or vice versa). This first approach can be used for any selectivity pattern. The other option allows for separate selectivity parameters for each sex plus

an additional parameter to define the scaling of one sex's peak selectivity relative to the other sex's peak. This second approach has only been implemented for a few selectivity patterns.

Male Selectivity as a Fraction of Female Selectivity (or vice versa):

If the "domale" flag is set to 1, then the selectivity parameters define female selectivity and the offset defined below sets male selectivity relative to female selectivity. The two sexes switch roles if the "domale" flag is set to 2. Generally it is best to select the option so that the dependent sex has lower selectivity, thus obviating the need to rescale for selectivities that are greater than 1.0. Sex specific selectivity is done the same way for all size and age selectivity options.

- p1 – size (age) at which a dogleg occurs (set to an integer at a bin boundary and do not estimate).
- p2 – $\ln(\text{relative selectivity})$ at minimum length or age = 0. Typically this will be set to a value of 0.0 (for no offset) and not estimated. It would be a rare circumstance in which the youngest/smallest fish had sex-specific selectivity.
- p3 – $\ln(\text{relative selectivity})$ at the dogleg.
- p4 – $\ln(\text{relative selectivity})$ at maximum length or max age.

For intermediate ages, the natural log values are linearly interpolated on size (age).

If selectivity for the dependent sex is greater than the selectivity for the first sex (which always peaks at 1.0), then the male-female selectivity matrix is rescaled to have a maximum of 1.0.

Male Selectivity Estimated as Offsets from Female Selectivity (or vice versa):

A new sex selectivity option (3 or 4) has been implemented for size selectivity patterns 1 (logistic) and 23 and 24 (double normal) or age selectivity pattern 20 (double normal age). Rather than calculate male selectivity as an offset from female selectivity, here the male selectivity is calculated by making the male parameters an offset from the female parameters (option 3), or females are offset from males with option 4. The description below applies to option 3. If the size selectivity pattern is 1 (logistic), then read 3 parameters:

- male parameter 1 is added to the first selectivity parameter (inflection)
- male parameter 2 is added to the second selectivity parameter (width of curve)
- male parameter 3 is the asymptotic selectivity

If the size selectivity pattern is 20, 23 or 24 (double normal), then:

- male parameter 1 is added to the first selectivity parameter (peak)

- male parameter 2 is added to the third selectivity parameter (width of ascending side); then $\exp(\text{this sum})$ per previous transform
- male parameter 3 is added to the fourth selectivity parameter (width of descending side); then $\exp(\text{sum})$ per previous transform
- male parameter 4 is added to the sixth selectivity parameter (selectivity at final size bin); then $1/(1+\exp(-\text{sum}))$ per previous transform
- male parameter 5 is the apical selectivity for males

Note that the male selectivity offsets currently cannot be time-varying because they are offsets from female selectivity, they inherit the time-varying characteristics of the female selectivity.

Dirichlet Multinomial Error for Data Weighting

If the Dirichlet multinomial error distribution was selected in the data file for either length or age data weighting an additional parameter line is required immediately following the selectivity parameter block.

The list of parameters to be read from the above setup would be:

LO	HI	INIT	<other entries>	Block Fxn	Parameter Label
-5	10	0.5	...	0	#ln(DM theta) Age or Length 1
-5	10	0.5	...	0	#ln(DM theta) Age or Length 2

Selectivity Time-Varying Options

The setup required for specifying time-varying options across parameter sections within SS (i.e., biology, recruitment, catchability, or selectivity) is the same and the basic description can be found above. If time-varying selectivity parameters are invoked, short parameter lines are required to be read following the standard selectivity parameter section. If no selectivity parameters invoke any time-varying properties, this section is left blank (or completely commented out with #).

Two-Dimensional Auto-Regressive Selectivity

A new experimental feature added within SS v.3.30.03.02. Earlier versions do not have this feature and hence this input is not expected. This features allows for a matrix of auto-correlation selectivity

deviations by age or length and time as described in Xu et al. (2019). Unlike other options for time-varying selectivity in SS, these deviations are not in the selectivity parameters themselves, but exponential multipliers on whatever baseline selectivity has been chosen.

When using this option for age-based selectivity, if there are not too many ages, a good choice for the baseline selectivity might be random walk selectivity (pattern 17) because it provides the most flexibility, allowing the deviations to be used only for the annual deviations around this baseline rather than the account for misspecification of the underlying functional form. Otherwise, a simple parametric selectivity form like logistic, exponential logistic, or double normal would be a better choices. This option has not yet been explored adequately to provide guidance on best practices.

Typical Value			Description and Options							
1			Two-dimensional auto-regressive selectivity: 0 = Not used, 1 = Use 2D AR.							
COND = 1 Read the following long parameter lines:										
Fleet	Ymin	Ymax	Amin	Amax	Sigma Amax	Use Rho	Len(1)/ Age(2)	Phase	Before Range	After Range
1	1979	2015	2	10	1	1	2	5	0	0
continued:										
LO	HI	INIT	PRIOR	PRIOR SD	PRIOR TYPE	PHASE	LABEL			
0	4	1	1	0.1	6	-4	#Sigma selex			
-1	1	0	0	0.1	6	-4	#Rho year			
-1	1	0	0	0.1	6	-4	#Rho age			
#Terminator line of 11 in length indicates the end of parameter input lines										
-9999	1	1	1	1	1	1	1	1	1	1

Parameter Definitions:

- Fleet: Fleet number to which semi-parametric deviations should be added.
- Ymin: First year with deviations.
- Ymax: Last year with deviations.

- Amin: First integer age (or population length bin index) with deviations.
- Amax: Last integer age (or population length bin index) with deviations.
- Sigma Amax: the last age (or population length bin index) for which a separate sigma should be read. For simplicity, it is easiest to set Sigma Amax equal to the Amin value which requires only a single sigma line (otherwise $N = \text{Sigma Amax} - \text{Amin}$ input lines are required for the sigma parameters). Also, a value < 0 will allow for only one sigma parameter should be read and then used for all bins (this is option is available in versions 3.30.15 and higher).
- Use Rho: Use autocorrelation parameters.
- Len(1) / Age(2): 1 or 2 to specify whether the deviations should be applied to length- or age-based selectivity.
- Phase: Phase to begin estimation of the deviation parameters.
- Before Range: How should selectivity be modeled in the years prior to Ymin? Available options are (0) apply no deviations, (1) use deviations from the first year with deviations (Ymin), and (3) use average across all years with deviations (Ymin to Ymax).
- After Range: Similar to Before Range but defines how selectivity should be modeled after Ymax.

Tag Recapture Parameters

Specify if tagging data are being used:

Typical Value		Description and Options				
1		Tagging Data Present: 0 = No tagging data, or if tagging data is present in the data file, a value of 0 here will auto-generate the tag parameter section in the control.ss_new file. 1 = Read following lines of tagging data.				
COND = 1 Read the following long parameter lines:						
LO	HI	INIT	PRIOR	<other entries>	Block Fxn	Parameter Label
-10	10	9	9	...	0	#TG loss init 1
-10	10	9	9	...	0	#TG loss chronic 1
1	10	2	2	...	0	#TG loss overdispersion 1

LO	HI	INIT	PRIOR	<other entries>	Block Fxn	Parameter Label
-10	10	9	9	...	0	#TG report fleet 1
-4	0	0	0	...	0	#TG report decay 1

If there are multiple tagging groups or multiple fleets reporting tagged fish the additional needed parameter lines should be entered in order for each parameter type (i.e, TG loss init 1, TG loss init 2, TG loss chronic 1, TG loss chronic 2,..., TG report decay 1, TG report decay 2).

Five parameter types are required for tagging data. Both the tag loss parameters and the reporting rate parameters are input as any real number and a logistic transformation is used to keep the resulting rates between 0 and 1:

- Initial tag loss, representing fraction of tags that are shed or associated with tag-induced mortality immediately after tagging (1 parameter per tag group).
- Chronic tag loss, representing annual rate of tag loss (1 parameter per tag group).
- Overdispersion parameter for the negative binomial likelihood associated with the total number of tag recaptures per tag group (1 parameter per tag group). This value represents the ratio of the variance to the mean of the distribution. As parameterized in ADMB, the likelihood is only defined for overdispersion parameters > 1, but a value of 1.001 will produce results essentially equal to those associated with a Poisson likelihood in which the variance is equal to the mean. This parameter is not transformed like the other 4 types.
- Initial reporting rate (1 parameter per fleet), the reporting rate of tags for a specific fleet.
- Reporting rate decay representing annual reduction in reporting rate (1 parameter per fleet).

The tagging reporting rate parameter is transformed within SS during estimation to maintain a positive value and is reported according to the transformation:

$$\text{Tagging Reporting Rate} = \frac{e^{\text{input parameter}}}{1 + e^{\text{input parameter}}} \quad (30)$$

Currently, tag parameters cannot be time-varying.

A shortcoming was identified in the recapture calculations when using Pope's F Method and multiple seasons in SS versions prior to 3.30.14. The internal calculations were corrected in version 3.30.14. Now the Z-at-age is applied internally for calculations of fishing pressure on the population when using the Pope calculations.

Mirroring of Tagging Parameters

In version 3.30.14, the ability to mirror the tagging parameters from another tag group or fleet was added. With this approach, the user can have just one parameter value for each of the five tagging parameter types and mirror all other parameters. Note that parameter lines are still required for the mirrored parameters and only lower numbered parameters can be mirrored. Mirroring is evoked through the phase input in the tagging parameter section. The options are:

- No mirroring among tag groups or fleets: phase > -1000,
- Mirror the next lower (i.e., already specified) tag group or fleet: phase = -1000 and set other parameter values the same as next lower Tag Group or fleet,
- Mirror a lower (i.e., already specified) tag group of fleet x: phase = -100x and set other parameter values the same as the mirrored tag group or fleet(i.e., if you would like to mirror fleet 1 then the phase should -1001).

To avoid having to specify mirrored parameter lines, the tag parameters can be auto-generated. The control.ss_new file created after running this model will have a full set of tagging parameter lines to use in future SS runs.

Variance Adjustment Factors

When doing iterative re-weighting of the input variance factors, it is convenient to do this in the control file, rather than the data file. This section creates that capability.

Read variance adjustment factors to be applied:

Factor	Fleet	Value	Description
1	2	0.5	# Survey CV for survey/fleet 2
4	1	0.25	# Length data for fleet 1
4	2	0.75	# Length data for fleet 2
-9999	0	0	# End read

Additive Survey CV - Factor 1

The survey input variance (labeled survey CV) is actually the standard deviation of the $\ln(\text{survey})$. The variance adjustment is added directly to this standard deviation. Set to 0.0 for no effect. Negative values are OK, but will crash if adjusted standard deviation value becomes negative.

Additive Discard - Factor 2

The input variance is the CV of the observation. Because this will cause observations of near zero discard to appear overly precise, the variance adjustment is added to the discard standard deviation, not to the CV. Set to 0.0 for no effect.

Additive Mean Body Weight - Factor 3

The input variance is in terms of the CV of the observation. Because such data are typically not very noisy, the variance adjustment is added to the CV and then multiplied by the observation to get the adjusted standard deviation of the observation.

Multiplicative Length Composition - Factor 4

The input variance is in terms of an effective sample size. The variance adjustment is multiplied times this sample size. Set variance adjustment to 1.0 for no effect.

Multiplicative Age Composition - Factor 5

Age composition is treated the same way as length composition.

Multiplicative Size-at-Age - Factor 6

Size-at-age input variance is the sample size for the N observations at each age. The variance adjustment is multiplied by these N values. Set to 1.0 for no effect.

Multiplicative Generalized Size Composition - Factor 7

Generalized size composition input variance is the sample size for each observation. The variance adjustment for each fleet is multiplied by these sample sizes. Set to 1.0 for no effect.

Variance Adjustment Usage Notes

The report.sso output file contains information in the "FIT_LEN_COMPS" and "FIT_AGE_COMPS" useful for determining if an adjustment of these input values is warranted to better match the scale of the average residual to the input variance scale.

Because the actual input variance factors are modified, it is these modified variance factors that are used when creating parametric bootstrap data files. So, the control files used to analyze bootstrap generated data files should have the variance adjustment factors reset to null levels.

Lambdas (Emphasis Factors)

These values are multiplied by the corresponding likelihood component to calculate the overall negative log likelihood to be minimized.

Typical Value	Description and Options
4	Max lambda phase: read this number of lambda values for each element below. The last lambda value is used for all higher numbered phases.
1	SD offset: 0 = The $\ln(\text{like})$ to omit the $+\ln(s)$ term, 1 = The $\ln(\text{like})$ to include the $\ln(s)$ term for CPUE, discard, growth CV, mean body weight, recruitment deviations. If you are estimating any variance parameters, SD offset must be set to 1.

Lambda Usage Notes

If the CV for size-at-age is being estimated and the model contains mean size-at-age data, then the flag for inclusion of the $+ \ln(\text{stddev})$ term in the likelihood must be included. Otherwise, the model will always get a better fit to the mean size-at-age data by increasing the parameter for CV of size-at-age.

The reading of the lambda values has been substantially altered with SS v.3.30. Instead of reading a matrix containing all the needed lambda values, SS now just reads those elements that will be given a value other than 1.0. After reading the datafile, SS sets lambda equal to 0.0 if there are no data for a particular fleet/data type, and a value of 1.0 if data exist. So beware if your data files had data but you had set the lambda to 0.0 in a previous version of SS. First read an integer for the number of changes.

Read the lambda adjustments by fleet and data type:

Likelihood Component	Fleet	Phase	Lambda Value	SizeFreq Method
1	2	2	1.5	1
4	2	2	10	1
4	2	3	0.2	1
-9999	1	1	1	1

The codes for component are:

1 = survey	10 = recruitment deviations
2 = discard	11 = parameter priors
3 = mean weight	12 = parameter deviations
4 = length	13 = crash penalty
5 = age	14 = morph composition
6 = size frequency	15 = tag composition
7 = size-at-age	16 = tag negative binomial
8 = catch	17 = F ballpark
9 = initial equilibrium catch	18 = regime shift

Controls for Variance of Derived Quantities

Additional standard deviation reported may be selected:

Typical Value	Description and Options
0	0 = No additional std dev reporting; 1 = read values below; and 2 = read specification for reporting stdev for selectivity, size,numbers, and natural mortality.

COND = 1 or 2: Read the following lines (split into 3 rows for readability):

- 4 values related to selectivity:
 1. fleet number (or 0 if not used),
 2. type of selectivity to report: 1 = length, 2 = age, or 3 = combined (age with length),
 3. year, and
 4. number of bins to report, list entered below (note - combined will report in age bins).
- 2 values related to growth:
 1. growth pattern (or 0 if not used), and
 2. number of ages for reporting: list entered below (note - each sex will be reported).
- 3 values related to numbers-at-age:
 1. area (enter -1 for sum of all areas), or 0 for not used,
 2. year, and
 3. number of ages to report, list entered below (note - each sex will be reported and summed across growth patterns).

COND = 2, enter the above quantities plus (available in versions 3.30.15 and higher):

- 2 additional values related to natural mortality-at-age:
 1. growth pattern (or 0 for not used), and

2. number of ages for reporting, list entered below (note - each sex will be reported).

Depending upon the entered options above subsequent conditional inputs may be need.

- COND if selectivity specs number of bins to report > 0:
 - Vector of bin indexes or -1 in first bin will overwrite entered vector with auto-generated list to cover logical range
- COND if growth specs number of bins to report > 0 and empirical weight at age not used:
 - ... as for selectivity bins
- COND if numbers-at-age specs number of ages > 0:
 - ... as for selectivity bins
- COND == 2 and mortality-at-age specs number of ages > 0:
 - ... as for selectivity bins

Example Input:

1 0 = No additional std dev reporting;
 1 = read values below; and
 2 = read specification for reporting stdev for selectivity, size,numbers, and
 natural mortality.

1 1 -1 5 # Selectivity
 1 5 # Growth
 1 -1 5 # Numbers-at-age
 5 15 25 35 43 # Vector with selectivity std bins (-1 in first bin to self-generate)
 1 2 14 26 40 # Vector with growth std ages picks (-1 in first bin to self-generate)
 1 2 14 26 40 # Vector with numbers-at-age std ages (-1 in first bin to self-generate)

999 #End of the control file input

Optional Inputs

Empirical Weight-at-Age (wtatage.ss)

SS has the capability to read empirical body weight at age for the population and each fleet, in lieu of generating these weights internally from the growth parameters, weight-at-length, and size-selectivity. Selection of this option is done by setting an explicit switch near the top of the control file. The values are read from a separate file named, wtatage.ss. This file is only required to exist if this option is selected.

The first value read is a single integer for the maximum age used in reading this file. So if the maximum age is 40, there will be 41 columns of weight-at-age entries to read, with the first column being for age 0. If the number of ages specified in this table is greater than maximum age in the model, the extra weight-at-age values are ignored. If the number of ages in this table is less than maximum age in the model, the weight-at-age data for the number of ages in the file is filled in for all unread ages out to maximum age.

The format of this input file is:

40 Maximum Age								
			Growth		Birth			
Year	Season	Gender	Pattern	Season	Fleet	Age-0	Age-1	...
1971	1	1	1	1	-2	0	0	0.1003
1971	1	1	1	1	-1	0.0169	0.0864	0.2495
1971	1	1	1	1	0
1971	1	1	1	1	1
1971	1	1	1	1	2
-9999	1	1	1	1	0

where:

- Fleet = -2 is age-specific fecundity (e.g., fecundity-scalar*maturity*weight-at-age), so time-varying fecundity is possible to implement.
- Fleet = -1 is population weight-at-age at middle of the season.
- Fleet = 0 is population weight-at-age at the beginning of the season.

- There must be an entry for fecundity, population weight-at-age at the middle of the season, population weight-at-age at the beginning of the season, and weight-at-age for each fleet (even if these vectors are identical in some cases).
- Fleets that do not use biomass do not need to have weight-at-age assigned.
- Growth pattern and birth season probably will never be used, but are included for completeness.
- A negative value for year will fill the table from that year through the ending year of the forecast, overwriting anything that has already been read for those years.
- Judicious use of negative years in the right order will allow user to enter blocks without having to enter a row of info for each year
- There is no internal error checking to verify that weight-at-age has been read for every fleet and for every year.
- In the future, there could be an option to use another value of the control file switch to turn off all aspects of growth parameters and size selectivity.
- The values entered for $\text{endyr} + 1$ will be used for the benchmark calculations and for the forecast; this aspect needs a bit more checking.

Caveats

- SS will still calculate growth curves from the input parameters and can still calculate size-selectivity and can still examine size-composition data.
- However, there is no calculation of weight-at-age from the growth input, so there is no way to compare the input weight-at-age from the weight-at-age derived from the growth parameters.
- If weight-at-age is read and size-selectivity is used, a warning is generated.
- If weight-at-age is read and discard/retention is invoked, then a warning is generated because of untested consequences for the body weight of discarded fish.
- Age-0 fish seem to need to have $\text{weight}=0$ for spawning biomass calculation (code -2).

User Testing

- Setup a model with age-maturity (option 2) and only age selectivity.

- Take the output calculation of weight-at-age and fecundity-at-age from report.sso and put into wtatage.ss (as shown above).
- Re-run SS with this input weight-at-age (Maturity_Option 5) to see if identical results are produced relative to the run that had generated the weight-at-age from the growth parameters.

runnumber.ss

This file contains a single integer value. It is read when the program starts, incremented by 1, used when processing the profile value inputs (see below), used as an identifier in the batch output, then saved with the incremented value. Note that this incrementation may not occur if a run crashes.

profilevalues.ss

This file contains information for changing the value of selected parameters for each run in a batch. In the control file, each parameter that will be subject to modification by profilevalues.ss is designated by setting its phase to -9999 .

The first value in profilevalues.ss is the number of parameters to be batched. This value **MUST** match the number of parameters with phase set equal to -9999 in the control file. The program performs no checks for this equality. If the value is zero in the first field, then nothing else will be read. Otherwise, the model will read runnumber * Nparameters values and use the last Nparameters of these to replace the initial values of parameters designated with phase = -9999 in the control file.

Usage Note: If one of the batch runs crashes before saving the updated value of runnumber.ss, then the processing of the profilevalue.ss will not proceed as expected. Check the output carefully until a more robust procedure is developed. Also, this options was created before use of R became widespread. You probably can create a more flexible approach using R today.

Likelihood components

The objective function L is the weighted sum of the individual components indexed by kind of data i , and fishery/survey f as appropriate:

$$L = \sum_{i=1}^I \sum_{f=1}^{A_f} \omega_{i,f} L_{i,f} + \omega_R L_R + \sum_{\theta} \omega_{\theta} L_{\theta} + \sum_P \omega_P L_P \quad (31)$$

where L is the total objective function, i is the index of the objective function component, A_f is the number of fleets, $L_{i,f}$ is the objective function for data kind i for the fishery/survey f , $\omega_{i,f}$ is a weighting factor for each objective function component, ω is the parameter priors, and P is the random parameter deviations.

The components of the objective function based on the model set-up and data are:

Index	Source	Data/Parameter	Error structure
i	fishery/survey f	CPUE or Abundance index	user choice
i	fishery f	Discard biomass	user choice
i	fishery/survey f	Mean body weight	normal
i	fishery/survey f	Length composition	multinomial or Dirichlet-multinomial
i	fishery/survey f	Age composition	multinomial or Dirichlet-multinomial
i	fishery/survey f	Mean size-at-age	normal
i	fishery/survey f	Tag-recapture 1	multinomial
i	fishery/survey f	Tag-recapture 2	negative binomial
i	fishery f	Initial equilibrium catch	normal
R		Recruitment deviations	lognormal
P		Random parameter devs.	normal
θ		Parameter priors	user choice
F_B		F ballpark penalty	
C_P		Crash penalty	

Running SS

Command Line Interface

The name of the SS executable files often contains the phrase "safe" or "opt" (for optimized). The safe version includes checking for out of bounds values and should always be used whenever there is a change to the data file. The optimized version runs slightly faster but can result in data not being included in the model as intended if the safe version has not been run first. A file named "ss.exe" is typically the safe version unless the result of renaming by the user. In some situations, users may wish to rename the file they are using to ss.exe, but the longer file name can be used.

On Mac and Linux computers, the executable does not include an extension (like .exe on Windows). Running the executable on from the DOS command line in Windows simply require typing the executable name (without the .exe extension):

```
> ss
```

On Mac and Linux computers, the executable name must be preceded by a period and slash (unless it's location has been added to the user's PATH):

```
> ./ss
```

Additional ADMB commands can follow the executable name, such as "-nohess" to avoid calculating the Hessian matrix. To see a full list of options, add "-?" after the executable name (with a space in between).

On all operating systems, a copy of the SS executable can either be located in the same directory as the model input files or in a central location and referenced either by adding it to the PATH or by a script files. Further discussion on script files for Windows is below. Editing the PATH is not covered here.

Example of DOS batch input file

One file management approach is to put ss.exe in its own folder (example: C:\SS_model) and to put your input files in separate folder (example: C:\My Documents \SS_runs). Then a DOS batch file in the SS_runs folder can be run at the command line to start ss.exe. All output will appear in SS_runs folder.

A DOS batch file (e.g. SS.bat) might contain some explicit ADMB commands, some implicit commands, and some DOS commands:

```
c:\SS_model\ss.exe -cbs 5000000000 -gbs 50000000000 \%1 \%2 \%3 \%4
del ss.r0*
del ss.p0*
del ss.b0*
```

In this batch file, the `-cbs` and `-gbs` arguments allocate a large amount of memory for SS to use (you may need to edit these for your computer and SS configuration), and the `%1`, `%2` etc. allows passing of command line arguments such as `-nox` or `-nohess`. You add more items to the list of `%` arguments as needed.

An easy way to start a command line in your current directory (SS_runs) is to create a shortcut to the DOS command line prompt. The shortcut's target would be:

```
> %SystemRoot%\system32\cmd.exe
```

And it would start in:

```
> %CURRDIR%
```

An alternative shortcut is to have the executable within the model folder then use `Ctrl+Shift+Right Click` and then select either "Open powershell window here" or "Open command window here", depending upon your computer. From the command window the executable name can be typed along with additional inputs (e.g., `-nohess`) and the model run. If using the powershell type `cmd` and then hit enter prior to calling the model (ss).

Simple Batch

This first example relies upon having a set of prototype files that can be renamed to `starter.ss` and then used to direct the running of SS. The example also copies one of the output files to save it from being overwritten. This sequence is repeated 3 times here and can be repeated an unlimited number of times. The batch file can have any name with the `.bat` extension, and there is no particular limit to the DOS commands invoked. Note that brief output from each run will be appended to `cumreport.sso` (see below).

```
del ss.cor
del ss.std
copy starter.r01 starter.ss
c:\admodel\ss\ss.exe -sdonly
copy ss.std ss-std01.txt
copy starter.r01 starter.ss
c:\admodel\ss\ss.exe -sdonly
copy ss.std ss-std02.txt
```

Complicated Batch

This second example processes 25 dat files from a different directory, each time using the same ctl and nam file. The loop index is used in the file names, and the output is searched for particular keywords to accumulate a few key results into the file SUMMARY.TXT. Comparable batch processing can be accomplished by using R or other script processing programs.

```
del summary.txt
del ss-report.txt
copy /Y runnumber.zero runnumber.ss
FOR /L %%i IN (1,1,25) DO (
copy /Y ..\MakeData\A1-D1-%%i.dat Ase1.dat
del ss.std
del ss.cor
del ss.par
c:\admodel\ss\ss.exe
copy /Y ss.par A1-D1-A1-%%i.par
copy /Y ss.std A1-D1-A1-%%i.std
find "Number" A1-D1-A1-%%i.par >> Summary.txt
find "hessian" ss.cor >> Summary.txt)
```

Running Without Estimation

There may be time when users will want to run the model without parameter estimation. There are two suggested alternative approaches to do this with SS and ADMB.

The first approach requires the user to change the maximum phase value in the starter.ss file to 0 then running the model via the command widow as without calculating the hessian:

```
ss -nohess
```

The second approach is done all through the command window using the following commands:

```
ss -maxfn 0 -phase 50 -nohess
```

where maxfun specifies the number of function calls and phase is the maximum phase for the model to start estimation where the number should be greater than the maximum phase for estimating parameters within the model.

However, the approaches above differ in subtle ways. First, if the maximum phase is set to 0 in the starter file the total likelihood will differ by a small amount (0.25 likelihood units) compared to the

second approach which sets the maxfun and phase in the command window. This small difference is due a dummy parameter which is evaluated by the objective function when maximum phase in the starter is set to 0, resulting in a small contribution to the total likelihood of 0.25. However, all other likelihood components should not change.

The second difference between the two no estimation approaches is the reported number of "Active_count" of parameters in the Report file. If the command line approach is used (ss -maxfn 0 -phase 50 -nohess) then the active number of parameters will equal the number of parameters with positive phases, but because the model is started in a phase greater than the maximum phase in the model, these parameters do not move from the initial values in the control file (or the par file). The first approach where the maximum phase is changed in the starter file will report the number of "Active_count" parameters as 0.

The final thing to consider when running a model without estimation is whether you are starting from the par file or the control file. If you start from the par file (specified in the starter file: 1=use ss.par) then all parameters, including parameter deviations, will be fixed at the estimated values. However, if the model is not run with the par file, any parameter deviations (e.g., recruitment deviations) will not be included in the model run (a user could paste in the estimated recruitment deviations into the control file).

Running Parameter Profiles

Users will often want to run profiles over specific parameter to evaluate the information in the model to estimate the parameter based on changes in the log likelihood. There are two ways this can be done.

The first option is the use functions within `r4ss` to run the profile, summarize quantities across runs, and plot the output. The `SS_profile()` function will run the profile based on function inputs, `SSgetoutput()` will read quantities from each run Report file, `SSsummarize()` will summarize key model quantities, and the `SSplotProfile()` and `PinerPlot()` functions can be used to visualize results. Additional information regarding `r4ss` can be found on page 171.

The second way is to create and run a batch file to profile over parameters. This example will run a profile on natural mortality and spawner-recruitment steepness, of course. Edit the control file so that the natural mortality parameter and steepness parameter lines have the phase set to -9999. Edit `STARTER.SS` to refer to this control file and the appropriate data file.

Create a `PROFILEVALUES.SS` file

2 # number of parameters using profile feature

0.16 # value for first selected parameter when runnumber equals 1

0.35 # value for second selected parameter when runnumber equals 1

0.16 # value for first selected parameter when runnumber equals 2

0.40 # value for second selected parameter when runnumber equals 2
0.18 # value for first selected parameter when runnumber equals 3
0.40 # value for second selected parameter when runnumber equals 3
etc.; make it as long as you like.

Create a batch file that looks something like this. Or make it more complicated as in the example above.

```
del cumreport.sso
copy /Y runnumber.zero runnumber.ss % so you will start with runnumber=0
C:\SS330\ss.exe
C:\SS330\ss.exe
C:\SS330\ss.exe
```

Repeat as many times as you have set up conditions in the PROFILEVALUES.SS file. The summary results will all be collected in the cumreport.sso file. Each step of the profile will have a unique run number and its output will include the values of the natural mortality and steepness parameters for that run.

Re-Starting a Run

SS model runs can be restarted from a previously estimated set of parameter values. In the starter.ss file, enter a value of 1 on the first numeric input line. This will cause SS to read the file ss.par and use these parameter values in place of the initial values in the control file. This option only works if the number of parameters to be estimated in the new run is the same as the number of parameters in the previous run because only actively estimated parameters are saved to the file ss.par. The file ss.par can be edited with a text editor, so values can be changed and rows can be added or deleted. However, if the resulting number of elements does not match the setup in the control file, then unpredictable results will occur. Because ss.par is a text file, the values stored in it will not give exactly the same initial results as the run just completed. To achieve greater numerical accuracy, SS can also restart from ss.bar which is the binary version of ss.par. In order to do this, the user must make the change described above to the starter.ss file and must also enter `-binp ss.bar` as one of the command line options.

Debugging Tips

When SS input files are causing the program to crash or fail to produce sensible results, there are a few steps that can be taken to diagnose the problem. Before trying the steps below, examine the echoinput.sso file. It is highly annotated, so you should be able to see if SS is interpreting your input

files as you intended. Additionally, users should check the warning.sso file when attempting to debug a non-running model.

1. Set the turn_off_phase switch to 0 in the starter.ss file. This will cause the mode to not attempt to adjust any parameters and simply converges a dummy parameter. It will still produce a Report.sso file, which can be examined to see what has been calculated from the initial parameter values.
2. Turn the verbosity level to 2 in the starter.ss file. This will cause the program to display the value of each likelihood component to the screen on each iteration. So if the program is creating an illegal computation (e.g. divide by zero), it may show you which likelihood component contains the problematic calculation. If the program is producing a Report.sso file, you may then see which observation is causing the illegal calculation.
3. Run the program with the command `ss »SSpipe.txt`. This will cause all screen display to go to the specified text file (note, delete this file before running because it will be appended to). Examination of this file will show detailed statements produced during the reading and preprocessing of input files.
4. If SS fails to achieve a proper Hessian it exits without writing the detailed outputs in the FINAL_SECTION. If this happens, you can do a run with the `-nohess` option so you can view the Report.sso to attempt to diagnose the problem.

Keyboard Tips

Typing "N" during a run will cause ADMB to immediately advance to the next phase of estimation.

Typing "Q" during a run will cause ADMB to immediately go to the final phase. This bypasses estimation of the Hessian and will produce all of the SS outputs, which are coded in the FINAL_SECTION.

Running MCMC

Run SS v3.30

- This gives maximum posterior density estimates, report file, Hessian matrix and the .cor file.
- Look for parameters stuck on bounds which will degrade efficiency of MCMC implementation.

Run SS v3.30 with arguments `-mcmc xxxx -mcsave yyyy`

- Where: xxxx is the number of iterations for the chain, and yyyy is the thinning interval (1000 is a good place to start).
- MCMC chain starts at the MPD values.
- Recommended: Remove existing .psv files in run directory to generate a new chain.
- Recommended: Set DOS run detail switch in starter file to 0; reporting to screen will dramatically slow MCMC progress.
- Optional: Add `-nohess` to use the existing Hessian file without re-estimating.
- Optional: To start the MCMC chain from specific values change the par file; run the model with estimation, adjust the par file to the values that the chain should start from, change within the starter file for the model to begin from the par file, and call the MCMC function using `ss -mcmc xxxx -mcsave yyyy -nohess -noest`.
- Optional: Add `-noest -nohess` and modify starter file so that run will now start from the converged (or modified) parameter estimates in "ss.par".

Run SS v3.30 with argument `-mceval`

- This generates the posterior output files.
- Optional: Modify starter file entries to add a burn-in and thinning interval above and beyond the ADMB thinning interval applied at run time.
- Recommended: MCMC always begins with the maximum posterior density values and so a burn-in >0 should always be used.
- This step can be repeated for alternate forecast options (e.g., catch levels) without repeating step 2.

Optional: Run SS v3.30 with arguments `-mcr -mcmc xxxx -mcsave yyyy ...`

- This restarts and extends an uninterrupted chain previously completed (note that any intermediate runs without the `-mcr` command in the same directory will break this option).

Note: when SS switches to MCMC or MCEVAL mode, it sets all the bias adjustment factors to 1.0 for any years with recruitment deviations defined. SS does not create a report file after completing MCMC because it would show values based on the last MCMC step.

Output Files

Standard ADMB output files

Standard ADMB files are created by SS. These are:

ss.par – This file has the final parameter values. They are listed in the order they are declared in SS. This file can be read back into SS to restart a run with these values (see running SS on page 150).

ss.std – This file has the parameter values and their estimated standard deviation for those parameters that were active during the model run. It also contains the derived quantities declared as standard deviation report variables. All of this information is also report in the covar.sso. Also, the parameter section of Report.sso lists all the parameters with their SS generated names, denotes which were active in the reported run, displays the parameter standard deviations, then displays the derived quantities with their standard deviations.

ss.rep – This report file is created between phases so, unlike Report.sso, will be created even if the hessian fails. It does not contain as much output as shown in Report.sso.

ss.cor – This is the standard ADMB report for parameter and standard deviation report correlations. It is in matrix form and challenging to interpret. This same information is reported in covar.sso.

SS Summary

The ss_summary.sso file (available for versions 3.30.08.03 and later) is designed to put key model outputs all in one concise place. It is organized as a list. At the top of the file are some descriptors, followed by the likelihoods for each component, then the parameters and their standard errors, then the derived quantities and their standard errors. The total biomass, summary biomass, and catch were added to the quantities reported in this file in version 3.30.11 and later. This output was created to make it easy to compare the results between different versions of the executable, however, this file could be useful for numerous other uses.

SIS table

The SIS_table.sso file contains model output formatted for reading into the NMFS Species Information System (SIS). This file includes an assessment summary for categories of information

(abundance, recruitment, spawners, catch estimates) that are input into the SIS database. A time-series of estimated quantities which aggregates estimates across multiple areas and seasons are provided to summarize model results. Access to the SIS database is granted to all NOAA employees.

Derived Quantities

Before listing the derived quantities reported to the standard deviation report, there are a couple of topics that deserve further explanation.

Virgin Spawning Biomass vs Unfished Spawning Biomass

Unfished is the condition for which reference points (benchmark) are calculated. Virgin Spawning Biomass (B0) is the initial condition on which the start of the time-series depends. If biology or spawner-recruitment parameters are time-varying, then the benchmark year input in the forecast file tells SS which years to average in order to calculate "unfished". In this case, virgin recruitment and/or the virgin spawning biomass will differ from their unfished counterparts. Virgin recruitment and spawning biomass are reported in the mgmt_quant portion of the sd_report and are now labeled as "unfished" for clarity. Note that if $\ln(R0)$ is time-varying, then this will cause unfished to differ from virgin. However, if regime shift parameter is time-varying, then unfished will remain the same as virgin because the regime shift is treated as a temporary offset from virgin. Virgin spawning biomass is denoted as SPB_virgin and spawning biomass unfished is denoted as SPB_unf in the report file.

Virgin Spawning Biomass (B0) is used in four ways within SS:

1. Anchor for the spawner-recruitment relationship as virgin spawning biomass.
2. Basis for the initial equilibrium abundance.
3. Basis against which annual depletion is calculated.
4. Benchmark calculations.

However, if there is time-varying biology, then the 4th usage can have a different B0 calculation compared to the other usages.

Metric for Fishing Mortality

A generic single metric of annual fishing mortality is difficult to define in a generalized model that admits multiple areas, multiple biological cohorts, dome-shaped selectivity in size and age for each of many fleets. Several separate indices are provided and others could be calculated by a user from the detailed information in Report.sso.

Equilibrium SPR

This index focuses on the effect of fishing on the spawning potential of the stock. It is calculated as the ratio of the equilibrium reproductive output per recruit that would occur with the current year's F intensities and biology, to the equilibrium reproductive output per recruit that would occur with the current year's biology and no fishing. Thus it internalizes all seasonality, movement, weird selectivity patterns, and other factors. Because this index moves in the opposite direction than F intensity itself, it is usually reported as $1-SPR$. A benefit of this index is that it is a direct measure of common proxies used for F_{MSY} , such as $F_{40\%}$. A shortcoming of this index is that it does not directly demonstrate the fraction of the stock that is caught each year. The SPR value is also calculated in the benchmarks (see below).

The derived quantities report shows an annual SPR statistic. The options, as specified in the starter.ss file, are:

- 0 = skip
- 1 = $(1-SPR)/(1-SPR_{TGT})$
- 2 = $(1-SPR)/(1-SPR_{MSY})$
- 3 = $(1-SPR)/(1-SPR_{Btarget})$
- 4 = raw SPR

The SPR approach to measuring fishing intensity was implemented because the concept of a single annual F does not exist in SS because F varies by age, sex, and growth morph and season and area. There is no single F value that is applied to all ages unless you create a very simple SS setup with knife-edge selectivity. So, what you see in the SS options are various ways to calculate annual fishing intensity. They can be broken down into three categories. One is exploitation rate calculated simply as total catch divided by biomass from a defined age range. Another is SPR, which is a single measure of the equilibrium effect of fishing according to the F . The third category are various ways to calculate an average F . Some measures of fishing intensity will be misleading if applied inappropriately. For example, the sum of the apical F 's will be misleading if different fleets have very different selectivities or, worse, if they occur in different areas. The $F=Z-M$ approach to getting fishing intensity is a way

to have a single F that represents a number's weighted value across multiple areas, sexes, morphs, ages. An important distinction is that the exploitation rate and F -based approaches directly relate to the fraction of the population removed each year by fishing; whereas the SPR approach represents the cumulative effect of fishing so it's equivalent in F -space depends on M .

F std

This index provides a direct measure of fishing mortality. The options are:

- 0 = skip
- 1 = exploitation(Bio)
- 2 = exploitation(Num)
- 3 = sum(Frates)

The exploitation rates are calculated as the ratio of the total annual catch (in either biomass or numbers as specified) to the summary biomass or summary numbers on January 1. The sum of the F rates is simply the sum of all the apical F s. This makes sense if the F method is in terms of instantaneous F (not Pope's approximation) and if there are not fleets with widely different size/age at peak selectivity, and if there is no seasonality, and especially if there is only one area. In the derived quantities, there is an annual statistic that is the ratio of the can be annual F_std value to the corresponding benchmark statistic. The available options for the denominator are:

- 0 = raw
- 1 = F/F_{SPR}
- 2 = F/F_{MSY}
- 3 = $F/F_{Btarget}$

F-at-Age

Because the annual F is so difficult to interpret as a sum of individual F components, an indirect calculation of F -at-age is reported at the end of the report.sso file. This section of the report calculates Z -at-age simply as $\ln(N_{a+1,t+1}/N_{a,t})$. This is done on an annual basis and summed over all areas. It is done once using the fishing intensities as estimated (to get Z), and once with the F intensities set to 0.0 to get M -at-age. This latter sequence also provides a measure of dynamic $Bzero$. The user can then subtract the table of M -at-age/year from the table of Z -at-age/year to get a table of

F-at-age/year. From this apical F, average F over a range of ages, or other user-desired statistics could be calculated. Further work within SS with this table of values is anticipated.

MSY and other Benchmark Items

The following quantities are included in the `sdreport` vector `mgmt_quantities`, so obtain estimates of variance. Some additional quantities can be found in the benchmarks section of the `forecast_report.sso`.

Benchmark Item	Description
SSB_Unfished	Unfished reproductive potential (SSB is commonly female mature spawning biomass).
TotBio_Unfished	Total age 0+ biomass on January 1.
SmryBio_Unfished	Biomass for ages at or above the summary age on January 1.
Recr_Unfished	Unfished recruitment.
SSB_Btgt	SSB at user specified SSB target.
SPR_Btgt	Spawner potential ratio (SPR) at F intensity that produces user specified SSB target.
Fstd_Btgt	F statistic at F intensity that produces user specified SSB target.
TotYield_Btgt	Total yield at F intensity that produces user specified SSB target.
SSB_SPRtgt	SSB at user specified SPR target (but taking into account the spawner-recruitment relationship).
Fstd_SPRtgt	F intensity that produces user specified SPR target.
TotYield_SPRtgt	Total yield at F intensity that produces user specified SPR target.
SSB_MSY	SSB at F intensity that is associated with MSY; this F intensity may be directly calculated to produce MSY, or can be mapped to F_SPR or F_Btgt.
SPR_MSY	Spawner potential ratio (SPR) at F intensity associated with MSY.
Fstd_MSY	F statistic at F intensity associated with MSY.
TotYield_MSY	Total yield (biomass) at MSY.
RetYield_MSY	Retained yield (biomass) at MSY.

Brief cumulative output

Cum_Report.sso: contains a brief version of the run output, which is appended to current content of file so results of several runs can be collected together. This is especially useful when a batch of runs is being processed. Unless this file is deleted, it will contain a cumulative record of all runs done in that subdirectory. The first column contains the run number.

Bootstrap Data Files

It is possible to create bootstrap data files for SS where an internal parametric bootstrap function generates a simulated data set by parametric bootstrap sampling the expected values given the input observation error. The number of bootstrap data files to create is specified in the starter file where a value of 3 or greater will create a single bootstrap data file in the data.ss_new file. The first output provides the unaltered input data file (with annotations added). The second provides the expected values for only the data elements used in the model run. The third and subsequent outputs provide parametric bootstraps around the expected values. The bootstrap data file can be parsed using r4ss function SS_splitdat into individual data files and re-run with the model.

The bootstrapping procedure within SS is done via the following steps:

- Expected values of all input data are calculated (these are also used in the likelihood which compares observed to expected values for all data). The calculation of these expected values is described in detail under the "Observation Model" section of the appendix to Methot and Wetzel (2013).
- Parametric bootstrap data are calculated for each observation by sampling from a probability distribution corresponding to the likelihood for that data type using the expected values noted above. Examples of how this happens include the following:
 - Indices of abundance are sampled from the distribution used in the estimation model (as set in the "Errtype" column of the index configuration, most commonly lognormal but could be normal or T-distribution). The variability of the distribution from which the random sample is drawn is based on the combination of the input uncertainty and any estimated "Extra SD" parameter and any "add_to_survey_CV" value included under the input variance adjustments factors.
 - Length and age compositions are sampled from multinomial distributions with expected proportions in each bin based on the expected values and sample size equal to the

adjusted input sample size (input sample size multiplied by any inputs for "mult_by_lencomp_N" or "mult_by_agecomp_N" under the input variance adjustments factors).

- Discard data (fractions or absolute amounts) are generated from the chosen distribution (T-distribution, normal, log-normal, truncated-normal).
- Tagging data is generated using a negative binomial distribution to get the total number of recaptures for each tag group and a multinomial distribution to allocate those recaptures among fleets.

Given this, there are some assumptions implicit in the bootstrapping procedure (as implemented as of v.3.30.14) that users should be aware of:

- This procedure is strictly an observation error approach (i.e., process error in recruitment or any other time-varying parameter is not added). For simulation analyses, a common approach has been to input a new time series of recruitment deviations for each bootstrap data set.
- The sample size for conditional age-at-length data matches the inputs for each length bin. If stratified sampling is used, this may be appropriate, but if the ages represent a random subset of the selected population, this may result in less variability than if the associated length distribution were resampled.
- Currently, the aging error matrix is multiplied by the expected distribution of proportions at age, while the more correct order of operations would be to sample true ages, and then sample the observed age including aging error (it is possible these are mathematically identical).
- Often there is need to explore the removal (not include in the model fitting) of specific years in a data set which can be done by specifying a negative fleet number. If bootstrapping a data file, note that specifying a negative fleet in the data inputs for indices, length composition, or age composition will include the "observation" in the model (hence generating predicted values and bootstrap data sets for the data), but not in the negative log likelihood. The "observation values" used with negative fleet do not influence the predicted values, except when using tail compression with length or age composition. Non-zero values greater than the minimum tail compression should be used for the observation values when tail compression is being used, as using zeros or values smaller than the minimum tail compression can cause the predicted values to be reported as zero and shift predictions to other bins.

Forecast and Reference Points (Forecast-report.sso)

The Forecast-report file contains output of fishery reference points and forecasts. It is designed to meet the needs of the Pacific Fishery Management Council's Groundfish Fishery Management Plan, but it should be quite feasible to develop other regionally specific variants of this output.

The vector of forecast recruitment deviations is estimated during an additional model estimation phase. This vector includes any years after the end of the recruitment deviation time series and before or at the end year. When this vector starts before the ending year of the time series, then the estimates of these recruitments will be influenced by the data in these final years. This is problematic, because the original reason for not estimating these recruitments at the end of the time series was the poor signal/noise ratio in the available data. It is not that these data are worse than data from earlier in the time series, but the low amount of data accumulated for each cohort allows an individual datum to dominate the model's fit. Thus, an additional control is provided so that forecast recruitment deviations during these years can receive an extra weighting in order to counter-balance the influence of noisy data at the end of the time series.

An additional control is provided for the fraction of the log-bias adjustment to apply to the forecast recruitments. Recall that R is the expected mean level of recruitment for a particular year as specified by the spawner-recruitment curve and R' is the geometric mean recruitment level calculated by discounting R with the log-bias correction factor $e^{-0.5s^2}$. Thus a lognormal distribution of recruitment deviations centered on R' will produce a mean level of recruitment equal to R . During the modeled time series, the virgin recruitment level and any recruitments prior to the first year of recruitment deviations are set at the level of R , and the lognormal recruitment deviations are centered on the R' level. For the forecast recruitments, the fraction control can be set to 1.0 so that 100% of the log-bias correction is applied and the forecast recruitment deviations will be based on the R' level. This is certainly the configuration to use when the model is in MCMC mode. Setting the fraction to 0.0 during maximum likelihood forecasts would center the recruitment deviations, which all have a value of 0.0 in maximum likelihood mode, on R . Thus would provide a mean forecast that would be more comparable to the mean of the ensemble of forecasts produced in MCMC mode. Further work on this topic is underway.

Note:

- Cohorts continue growing according to their specific growth parameters in the forecast period rather than staying static at the end year values.
- Environmental data entered for future years can be used to adjust expected recruitment levels. However, environmental data will not affect growth or selectivity parameters in the forecast.

The top of the Forecast-report file shows the search for F_{SPR} and the search for F_{MSY} , allowing the user to verify convergence. Note: if the STD file shows aberrant results, such as all the standard deviations being the same value for all recruitments, then check the F_{MSY} search for convergence. The

F_{MSY} can be calculated, or set equal to one of the other F reference points per the selection made in starter.ss.

Main Output File, Report.sso

This is the primary output file. Its major sections are listed below.

The sections of the output file are:

- SS version number with date compiled. Time and date of model run. This info appears at the top of all output files.
- Comments
 - Input file lines starting with #C are echoed here.
- Keywords
 - List of keywords used in searching for output sections.
- Fleet Names
 - List of fishing fleet and survey names assigned in the data file.
- Likelihood
 - Final values of the negative log(likelihood) are presented.
- Input Variance Adjustments
 - The matrix of input variance adjustments is output here because these values affect the log likelihood calculations.
- Parameters
 - The parameters are listed here. For the estimated parameters, the display shows: Num (count of parameters), Label (as internally generated by SS), Value, Active_Cnt, Phase, Min, Max, Init, Prior, Prior_type, Prior_SD, Prior_Like, Parm_StD (standard deviation of parameter as calculated from inverse Hessian), Status (e.g. near bound), and Pr_atMin (value of prior penalty if parameter was near bound). The Active_Cnt entry is a count of the parameters in the same order they appear in the ss.cor file.
- Derived Quantities

- This section starts by showing the options selected from the starter.ss and forecast.ss input files:
 - * SPR ratio basis
 - * F report basis
 - * B ratio denominator

Then the time series of output, with standard deviation of estimates, are produced with internally generated labels. Note that these time series extend through the forecast era. The order of the output is: spawning biomass, recruitment, SPRratio, Fratio, Bratio, management quantities, forecast catch (as a target level), forecast catch as a limit level (OFL), Selex_std, Grow_std, NatAge_std. For the three "ratio" quantities, there is an additional column of output showing a Z-score calculation of the probability that the ratio differs from 1.0. The "management quantities" section is designed to meet the terms of reference for west coast groundfish assessments; other formats could be made available upon request. The standard deviation quantities at the end are set up according to specifications at the end of the control input file. In some cases, a user may specify that no derived quantity output of a certain type be produced. In those cases, SS substitutes a repeat output of the virgin spawning biomass so that vectors of null length are not created.

ADMB note, while vectors of null length are very useful for controlling optional model inputs, they cannot be used with current version of ADMB for sdreport quantities.

- Mortality-growth parameter by year after adjustments
 - This block shows the time series of mortality-growth parameters by year after adjustment by environmental links, blocks and deviations.
- Selectivity parameters (size) by year after adjustments
 - This block shows the size selectivity parameters, after adjustment, for each year in which a change occurs.
- Selectivity parameters (age) by year after adjustments
 - This block shows the age selectivity parameters, after adjustment, for each year in which a change occurs.
- Recruitment Distribution
 - This block shows the distribution of recruitment across growth patterns, genders, birth seasons, and areas in the end year of the model.
- Platoon Indexing

- This block shows the internal index values for various quantities. It can be a useful reference for complex model setups. The vocabulary is: Bio_Pattern refers to a collection of cohorts with the same defined growth and natural mortality parameters; sex is the next main index. If recruitment occurs in multiple seasons, then birth season is the index for that factor. The index labeled "Platoon" is used as a continuous index across all the other factor-specific indices. If sub-platoons are used, they are nested within the Bio_Pattern x Sex x Birth Season platoon. However, some of the output tables use the column label "platoon" as a continuous index across platoons and sub-platoons. Note that there is no index here for area. Each of the cohorts is distributed across areas and they retain their biological characteristics as they move among areas.
- Size Frequency Translation
 - If the generalized size frequency approach is used, this block shows the translation probabilities between population length bins and the units of the defined size frequency method. If the method uses body weight as the accumulator, then output is in corresponding units.
- Movement
 - This block shows movement rate between areas in a multi-area model.
- Exploitation
 - This block shows the time series of the selected F_std unit and the F multiplier for each fleet in terms of harvest rate (if Pope's approximation is used) or fully selected F.
- Index 2
 - This section reports the observed and expected values for each index. All are reported in one list with index number included as a selection field. At the bottom of this section, the root mean squared error of the fit to each index is compared to the mean input error level to assist the user in gauging the goodness-of-fit and potentially adjusting the input level of imprecision.
- Index 3
 - This section shows the parameter number assigned to each parameter used in this section.
- Discard
 - This is the list of observed and expected values for the amount (or fraction) discard.
- Mean Body Wt

- This is the list of observed and expected values for the mean body weight.
- Fit Len Comps
 - This is the list of the goodness of fit to the length compositions. The input and output levels of effective sample size are shown as a guide to adjusting the input levels to better match the model's ability to replicate these observations.
- Fit Age Comps
 - This has the same format as the length composition section.
- Fit Size Comps
 - This has the same format as the length composition section and is used for the generalized size composition summary.
- Len Selex
 - Here is the length selectivity and other length specific quantities for each fishery and survey.
- Age Selex
 - Here is reported the time series of age selectivity and other age-related quantities for each fishery and survey. Some are directly computed in terms of age, and others are derived from the combination of a length-based factor and the distribution of size-at-age.
- Environmental Data
 - The input values of environmental data are echoed here. In the future, the summary biomass in the previous year will be mirrored into environmental column –2 and that the age zero recruitment deviation into environmental column –1. Once so mirrored, they may enable density-dependent effects on model parameters.
- Numbers at Age
 - The output (in thousands of fish) is shown for each cohort tracked in the model.
- Numbers at Length
 - The output is shown for each cohort tracked in the model.
- Catch at Age

- The output is shown for each fleet. It is not necessary to show by area because each fleet operates in only one area.
- Biology
 - The first biology section shows the length-specific quantities in the ending year of the time series only. The derived quantity spawn is the product of female body weight, maturity and fecundity per weight. The second section shows natural mortality.
- Growth Parameters
 - This section shows the growth parameters, and associated derived quantities, for each year in which a change is estimated.
- Biology at Age
 - This section shows derived size-at-age and other quantities. It is the basis for the Bio report page of the Excel output processor.
- Mean Body Wt (begin)
 - This section reports the time series of mean body weight for each platoon. Values shown are for the beginning of each season of each year.
- Mean Size Time series
 - This section shows the time series of mean length-at-age for each platoon. At the bottom is the average mean size as the weighted average across all platoons for each gender.
- Age Length Key
 - This is reported for the midpoint of each season in the ending year.
- Age Age Key
 - This is the calculated distribution of observed ages for each true age for each of the defined ageing keys.
- Selectivity Database
 - This section contains the selectivities organized as a database, rather than as a set of vectors.
- Spawning Biomass Report 2, etc.

- The section shows annual total spawning biomass, then numbers-at-age at the beginning of each year for each Bio_Pattern and Sex as summed over sub-platoons and areas. Then Z-at-age is reported simply as $\ln(N_{t+1,a+1}N_{t,a})$. Then the Report_1 section loops back through the time series with all F values set to zero so that a dynamic Bzero, N-at-age, and M-at-age can be reported. The difference between Report_1 and Report_2 can be used to create an aggregate F-at-age.
- Composition Database
 - This section is reported to a separate file, compreport.sso, and contains the length composition, age composition, and mean size-at-age observed and expected values. It is arranged in a database format, rather than an array of vectors. Software to filter the output allows display of subsets of the database.

Using R To View Model Output (r4ss)

A collection of functions developed as a package, `r4ss`, for the statistical software R has been created to explore SS model output. The functions include tools for summarizing and plotting results, manipulating files, visualizing model parameterizations, and various other tasks. Currently, information on the code, including installation instructions, can be found on Github. The software package is under constant development to maintain compatibility with new versions of SS and to improve functionality.

Two of the most commonly used functions for model diagnostics are `SS_output` and `SS_plots`. After running a model using SS, the report can be read into R by the `SS_output` function which stores quantities in a list with named objects. This list can then be passed to the `SS_plots` function which creates a series of over 100 plots that are useful to visualize output such as model fit to the data and time series of quantities of interest.

The latest `r4ss` version on CRAN can be installed using a command like:

```
> install.packages("r4ss")
```

However, more frequent enhancements and bug fixes are posted to the GitHub project. The latest version of `r4ss` can be installed directly from GitHub at any time via the `devtools` package in R with the following commands:

```
> install.packages("devtools")
> devtools::install_github("r4ss/r4ss")
```

Note: `devtools` will give this message: *"WARNING: Rtools is required to build R packages, but is not currently installed."*. However, `Rtools` is not required for installing `r4ss` via `devtools`, so ignore this warning.

Once you have installed the `r4ss` package, it can be loaded in the regular manner:

```
> library(r4ss)
```

The results from a model run can be read in and plots created using the following commands:

```
> setwd("C:\\directory where model was run")
> base.model = SS_output(getwd())
> SS_plots(base.model)
```

The `SS_output()` function reads the Report file and stores this information in a list. The function by default will print select items to the screen that can be useful for evaluating a model (e.g., likelihood, estimated parameter with warnings if they hit bounds, parameters with high gradients).

This list created by `SS_output()` can be passed to other functions in the `r4ss` package to allow visualization of results.

The `SS_plots()` function will create plots from the list created by the `SS_output()` function that can be used to visualize model results. The plots by default will be written inside a "plots" folder inside the folder with the model files read by the `SS_output()` function. Additionally, an html file is written which allows for visualization of all plots created.

The functions included in `r4ss` ranging from general use to functions developed for specific model applications:

Core Functions	
<code>SS_output</code>	A function to create a list object for the output from Stock Synthesis
<code>SS_plots</code>	Plot many quantities related to output from Stock Synthesis
Plot functions called by <code>SS_plots</code> :	
<code>SSplotBiology</code>	Plot biology related quantities from Stock Synthesis model output, including mean weight, maturity, fecundity, and spawning output.
<code>SSplotCatch</code>	Plot catch related quantities
<code>SSplotCohorts</code>	Plot cumulative catch by cohort
<code>SSplotComps</code>	Plot composition data and fits
<code>SSplotData</code>	Timeline of presence/absence data by type, year, and fleet
<code>SSplotDiscard</code>	Plot fit to discard fraction
<code>SSplotIndices</code>	Plot indices of abundance and associated quantities
<code>SSplotMnwt</code>	Plot mean weight data and fits
<code>SSplotMovementMap</code>	Show movement rates on a map
<code>SSplotMovementRates</code>	Show movement rates on a map
<code>SSplotNumbers</code>	Plot numbers-at-age related data and fits
<code>SSplotRecdevs</code>	Plot recruitment deviations
<code>SSplotRecdist</code>	Plot of recruitment distribution among areas and seasons
<code>SSplotSelex</code>	Plot selectivity
<code>SSplotSexRatio</code>	Plot sex ratios

SSplotSummaryF	Plot time series summary of F (or harvest rate)
SSplotSpawnrecruit	Plot spawner-recruit curve
SSplotSPR	Plot SPR quantities
SSplotTags	Plot tagging data and fits
SSplotTimeseries	Plot time series data
SSplotYield	Plot yield and surplus production
SS_html	Create HTML files to view figures in browser
SS_fitbiasramp	Estimate bias adjustment for recruitment deviates
Model Comparisons and other diagnostics:	
SSplotPars	Plot distributions of priors, posteriors, and estimates
SSplotProfile	Plot likelihood profile results
PinerPlot	Plot fleet-specific contributions to likelihood profile
SSplotRetroRecruits	Make retrospective pattern of recruitment estimates (a.k.a. squid plot) as seen in Pacific hake assessments
Functions related to MCMC diagnostics:	
mcmc_nuisance	Summarize nuisance MCMC output
mcmc_out	Summarize, analyze, and plot key MCMC output
SSgetMCMC	Read MCMC output
SSplotMCMC_ExtraSelex	Plot uncertainty around chosen selectivity ogive from MCMC
Interactive tools for exploring functional forms:	
movepars	Explore movement parameterization
selfit	A function to visualize parameterization of double normal and double logistic selectivity in SS
selfit_spline	Visualize parameterization of cubic spline selectivity in SS
sel_line	A function for drawing selectivity curves
File manipulation for inputs:	
SS_readdat	Read data file

SS_readforecast	Read forecast file
SS_readstarter	Read starter file
SS_writedat	Write data file
SS_writeforecast	Write forecast file
SS_writestarter	Write starter file
SS_makedatlist	Make a list for SS data
SS_parlines	Get parameter lines from SS control file
SS_changepars	Change parameters in the control file
SSmakeMmatrix	Create inputs for entering a matrix of natural mortality by age and year
SS_profile	Run a likelihood profile in SS (incomplete)
NegLogInt_Fn	Calculated variances of time-varying parameters using SS implementation of the Laplace Approximation
File manipulations for outputs:	
SS_recdevs	Insert a vector of recruitment deviations into the control file
SS_splitdat	Split apart bootstrap data to make input file
Minor plotting functions:	
bubble3	Create a bubble plot
make_multifig	Create multi-figure plots
Make_multifig_sexratio	Create multi-figure plots of sex ratios
plotCI	Plot points with confidence intervals
rich_colors_short	Make a vector of colors
stackpoly	Plot stacked polygons
mountains	Make shaded polygons with a mountain-like appearance
Really specialized functions:	
DoProjectPlots	Make plots from Rebuilder program
IOTCmove	Make a map of movement for a 5-area Indian Ocean model

SSFishGraph	A function for converting SS output to the format used by FishGraph
TSCplot	Create a plot for the TSC report

Advanced SS Configuration Settings and Advice

Using Time-Varying Parameters

The approach to allowing parameters to have time-varying values has been completely overhauled in the transition from SS v.3.24 to SS v.3.30. Fortunately, the transition executable (sstrans.exe) will do the conversion for you, but you should review the new control file closely before simply running with it, especially for time-varying catchability parameters.

Time-Varying Parameter Change from Earlier SS Versions

In SS v.3.24, the group of biology parameters (termed mortality-growth parameters) and the selectivity parameters used the same long parameter line approach, but it was implemented with entirely different code, and hence was inefficient. The spawner-recruitment parameters used short parameter lines and a different approach for linkage to an environmental variable and the R1 offset provided a limited type of block. The catchability parameters also used short parameter lines and had its own approach to doing environmental linkage and random deviations, but not blocks. Then finally, the tagging parameters had long parameter lines, but there was no code to interpret any time-varying info in those lines. The situation was begging for a more modular approach.

Code Flow Version SS v.3.30

In SS v.3.30, mortality-growth, selectivity, stock recruitment relationship, catchability, and tag (soon but not as of v.3.30.12) base parameters all use long parameter lines and invoke blocks, trends, environmental linkages, and random deviations using identical syntax. As SS v.3.30 executes the SS_readcontrol code, it calls a function in SS_global called "create_timevary" whenever a base parameter has any one of the 4 types of time-varying options. In fact, block/trend, env and devs all can be applied to the same base parameter. Only blocks and trends are mutually exclusive, but any combined effect could be used together judiciously. "Create_timevary" creates all needed information to describe and index a list of time varying parameter specifications. in fact, if the auto-generation switch has been set to zero, then you will omit all of the needed parameters for implementing the time-varying effect and SS will auto-generate and use the needed parameters and write them out in the control.ss_new file. Then as SS gets into iterative parameter updating it starts by calling a function in SS_timevaryparm that processes each time-varying parameter specification (each of which can contain any combination of block/trend, env and dev specification) and creates a time-series of the parameter value that are used as SS subsequently loops through the years.

Parameter Order

The order of parameters has changed and the re-ordering is handled by the transition executable (sstrans.exe). Previously, for each of mortality-growth and selectivity parameters all environmental

link parameters were listed first, then block/trend parameters and then deviation parameters. In SS v.3.30, these parameters are re-organized such that all parameters that affect a base parameter are clustered together with block/trend first, then environmental, then deviation. So, if mortality-growth (MG) base parameters 3 and 7 had time varying changes, the order would look like:

MG base parameter 3	Block parameter 3-1
	Block parameter 3-2
	Environmental link parameter 3-1
	Deviation se parameter 3
	Deviation rho parameter 3
MG base parameter 7	Block parameter 7-1
	Deviation se parameter 7
	Deviation rho parameter 7

Link Functions

The functional form by which a time-varying parameter, Q , changes a base parameter, P , is a link function: $P'_y = f(P, Q)$. Typically, this is additive or multiplicative function, but the parameter mirroring feature is essentially a link that takes no parameter. Another type of link in SS is between a model state variable, such as available biomass, and the expected value for a survey. Typically, this is a simple proportional link taking one parameter, q , but the q power feature is essentially a 2 parameter link function. So, a parameter link function can change q over time, and a survey link function then uses the annual value of q to link the annual value of a state variable to the expected value for a survey. In SS v.3.24, various usages of positive and negative codes and other conventions were used to invoke additive vs multiplication links and other options. But as SS v.3.30 builds capability to allow an environment index to be a “survey” of a parameter deviation, we need a larger family of link functions such as logistic and even dome-shaped.

The link specifications in SS v.3.30 has been updated from SS v.3.24. Take special note of the environmental linkage specification where two bits of information are coded into one number. The new specification has the environmental link function denoted by the first environmental index to use specified by two additional. (e.g., environmental link specification of 204 is parsed by SS to use link type 2 using environmental variable 4).

Specification of Time-Varying Parameters in Long Parameter Lines

Selecting time-varying properties for a parameter is specified using element 8 - 14 in the long parameter line setup. Each element and the options for selection are as described here:

- Environmental Link and Variance (element 8)
 - The environmental link and variance (env_var&link) matrix is populated with the read environmental data for columns 1-N environmental variables. Alternatively, the environmental link can be linked to derived quantities mapped to columns -1 to -4 where each value corresponds to the following quantities:
 - * -1; for $\ln(\text{relative spawning biomass})$;
 - * -2; for recruitment deviation;
 - * -3; for $\ln(\text{relative summary biomass})$ (e.g., current year summary biomass divided by the unfished summary biomass);
 - * -4; for $\ln(\text{relative summary numbers})$.
 - So, environmental input 103 would use link type 1 (specified in element 9) and apply it to environmental data column 3 and environmental input -103 would use link type 1 and apply it to the "-3" column which is $\ln(\text{relative summary biomass})$.
 - These four derived quantities are all calculated at the beginning of each year within the model, so they are available inside SS to use as the basis for time-varying parameter links without violating any order of operations rule.
- Deviation Link (element 9)
 - 1 = multiplicative ($P_y * = \exp(\text{dev}_y * \text{dev}_{se})$),
 - 2 = additive ($P_y + = \text{env}_y * \text{dev}_{se}$),
 - 3 = random walk options are now implemented by using rho in the objective function. SS now expects the estimated deviations to be normal in distribution and the deviation values are multiplied by the standard error parameter as they are used,
 - 4 = zero reverting random walk with rho. The deviation parameter is now multiplied by the standard error parameter, rather than deviations being penalized according to a specified standard error (the approach in SS v.3.24).
 - The option of applying the final model year deviation value into the forecast period was added in v. 3.30.13. This new option is specified by selecting the appropriate deviation link option (1, 2, 3, or 4) and appending a 2 at the front (21, 22, 23, or 24) which will use the final year deviation value for all forecast years.
- Deviation Minimum Year (element 10)

- Year for deviations to start for parameter
- Deviation Maximum Year (element 11)
 - Year for deviations to end for parameter
- Deviation Phase (element 12)
 - integer, this available element in the long parameter line is now a deviation vector specific phase control
- Blocks (element 13)
 - >0: time block index for parameter.
 - -1: trend with final as offset from base parameter and offset values is in natural log space, also inflection year is in natural log space and the offset from $\ln(0.5)$. No additional parameter lines are required. Three parameters will be estimated; end trend parameter value logistic offset, inflection year logistic offset, and slope.
 - -2: trend with final as standalone value. No additional parameter lines are required. Three parameters will be estimated; end trend parameter value, inflection year, and slope.
 - -3 end value is a fraction of base parameter maximum - minimum; inflection year is fraction of end year - start year. No additional parameter lines are required. Three parameters will be estimated; end trend parameter value as a fraction, inflection year as a fraction, and slope.
- Block Functional Form (element 14)
 - 0: multiplicative parameter ($P_y = P_{base} * \exp(tv_para)$),
 - 1: additive parameter ($P_y = P_{base} + tv_para$),
 - 2: replace parameter ($P_y = tv_para$),
 - 3: random walk across blocks ($P_{block} = P_{block,-1} + tv_para$),
 - 4: mean reverting random walk

Block Trends

Additional information regarding the options for applying blocks (element 13):

- -1: Trend bounded by base parameter minimum maximum and parameters in transformed units (use with caution),
 - Logistic approach to trend as offset from base parameter
 - Transform the base parameter from the parameter section:

$$\text{temp} = -0.5 * \ln \left(\frac{\text{Parm}_{p, LB} - p, \text{Parm}_{UB} + 0.0000002}{\text{Parm}_p - \text{Parm}_{p, LB} + 0.0000001} - 1 \right) \quad (32)$$

- Add the offset. Note, that offset values in in the transform space.
- Back transform

$$\text{temp1} = \text{Parm}_{p, LB} + \frac{\text{Parm}_{p, UB} - \text{Parm}_{p, LB}}{1 + e^{-2 * \text{temp}}} \quad (33)$$

Parameterizing the Two-Dimensional Autoregressive Selectivity

When the two-dimensional autoregressive selectivity feature is turned on for a fleet, the selectivity is calculated as a product of the assumed selectivity pattern and a non-parametric deviation term deviating from this assumed pattern:

$$\hat{S}_{a,t} = S_a \exp^{\epsilon_{a,t}} \quad (34)$$

where S_a is specified in the corresponding age/length selectivity types section and it can be either parametric (recommended) or non-parametric (including any of the existing selectivity options in SS); $\epsilon_{a,t}$ is simulated as a two-dimensional first-order autoregressive (2D AR1) process:

$$\text{vec}(\epsilon) \sim MVN(\mathbf{0}, \sigma_s^2 \mathbf{R}_{\text{total}}) \quad (35)$$

where ϵ is the two-dimensional deviation matrix and $\sigma_s^2 \mathbf{R}_{\text{total}}$ is the covariance matrix for the 2D AR1 process. More specifically, σ_s^2 quantifies the variance in selectivity deviations and $\mathbf{R}_{\text{total}}$ is equal to the kronecker product (\otimes) of the two correlation matrices for the among-age and among-year AR1 processes:

$$\mathbf{R}_{\text{total}} = \mathbf{R} \otimes \tilde{\mathbf{R}} \quad (36)$$

$$\mathbf{R}_{a,\tilde{a}} = \rho_a^{|a-\tilde{a}|} \quad (37)$$

$$\tilde{\mathbf{R}}_{t,\tilde{t}} = \rho_t^{|t-\tilde{t}|} \quad (38)$$

where ρ_a and ρ_t are the among age and among year AR1 coefficients, respectively. When both of them are zero, \mathbf{R} and $\tilde{\mathbf{R}}$ are two identity matrices and their Kronecker product, $\mathbf{R}_{\text{total}}$, is also an identity matrix. In this case selectivity deviations are essentially identical and mutually independent:

$$\epsilon_{a,t} \sim N(0, \sigma_s^2) \quad (39)$$

Using the Two-Dimensional Autoregressive Selectivity

Note, Xu et al. (2019) has additional information on tuning the 2D AR selectivity parameters. First, fix the two AR1 coefficients (ρ_a and ρ_t) at 0 and tune σ_s iteratively to match the relationship:

$$\sigma_s^2 = SD(\epsilon)^2 + \frac{1}{(a_{max} - a_{min} + 1)(t_{max} - t_{min} + 1)} \sum_{a=a_{min}}^{a_{max}} \sum_{t=t_{min}}^{t_{max}} SE(\epsilon_{a,t})^2 \quad (40)$$

The minimal and maximal ages/lengths and years for the 2D AR1 process can be freely specified by users in the control file. However, we recommend specifying the minimal and maximal ages and years to cover the relatively "data-rich" age/length and year ranges only. Particularly we introduce:

$$b = 1 - \frac{\frac{1}{(a_{max}-a_{min}+1)(t_{max}-t_{min}+1)} \sum_{a=a_{min}}^{a_{max}} \sum_{t=t_{min}}^{t_{max}} SE(\epsilon_{a,t})^2}{\sigma_s^2} \quad (41)$$

as a measure of how rich the composition data is regarding estimating selectivity deviations. We also recommend using the Dirichlet-Multinomial method to "weight" the corresponding composition data while σ_s is interactively tuned in this step.

Second, fix σ_s at the value iteratively tuned in the previous step and estimate $\epsilon_{a,t}$. Plot both Pearson residuals and $\epsilon_{a,t}$ out on the age-year surface to check their 2D dimensions. If their distributions seems to be not random but rather be autocorrelated (deviation estimates have the same sign several ages and/or years in a row), users should consider estimating and then including the autocorrelations in $\epsilon_{a,t}$.

Third, extract the estimated selectivity deviation samples from the previous step for estimating ρ_a and ρ_t externally by fitting the samples to a stand-alone model written in Template-Model Builder (TMB). In this model, both ρ_a and ρ_t are bounded between 0 and 1 via applying a logic

transformation. If at least one of the two AR1 coefficients are notably different from 0, SS should be run one more time by fixing the two AR1 coefficients at their values externally estimated from deviation samples. The Pearson residuals and $\epsilon_{a,t}$ from this run are expected to distribute more randomly as the autocorrelations in selectivity deviations can be at least partially included in the 2D AR1 process.

Continuous seasonal recruitment

It is awkward in SS to set up a seasonal model such that recruitment can occur with similar and independent probability in any season of any year. Consequently, some users have attempted to setup SS so that each quarter appears as a year. They have set up all the data and parameters to treat quarters as if they were years (i.e., each still has a duration of 1.0 time step). This can work, but requires that all rate parameters be re-scaled to be correct for the quarters being treated as years.

Another option is available. If there is one season per year and the season duration is set to 3 (rather than the normal 12), then the season duration is calculated to be 3/12 or 0.25. This means that the rate parameters can stay in their normal per year scaling and this shorter season duration makes the necessary adjustments internally. Some other adjustments to make when doing quarters as years include:

- Re-index all "year seas" inputs to be in terms of quarter-year because all are now season 1; increase end year (endyr) value in sync with this.
- Increase max age because age is now in quarters.
- In the age error definitions, increase the number of entries in accord with new max age
- In the age error definitions, recode so that each quarter-age gets assigned to the correct age bin. This is because the age data are still in terms of age bins; i.e., the first 4 entries for quarter-ages 1 through 4 will all be assigned to age bin 1.5; the next four to age bin 2.5; you cannot accomplish the same result by editing the age bin values because the standard deviation of ageing error is in terms of age bin.
- In the control file, multiple the natural mortality age breakpoints and growth A_{min} and A_{max} values by 1/season duration.
- Decrease the R_0 parameter starting value because it is now the average number of recruitments per quarter year.
- Edit the recruitment deviation (rec_dev) start and end years to be in terms of quarter year.

- Edit any age selectivity parameters that refer to age to now refer to quarter age.
- If there needs to be some degree of seasonality to recruitment or some parameter, then you could create a cyclic pattern in the environmental input and make recruitment or some other parameter a function of this cyclic pattern.

A good test showing comparability of the 3 approaches to setting up a quarterly model should be done.

Detailed Information on SS Processes

The processes and calculations with SS can be complex and not transparent based on the model input files. Here additional information is provided to users to assist in understanding some of these processes.

Jitter

The jitter function has been updated with SS v.3.30. The following steps are now performed to determine the jittered starting parameter values (illustrated in Figure 1):

1. A normal distribution is calculated such that the $\text{pr}(P_{\text{MIN}}) = 0.1\%$ and the $\text{pr}(P_{\text{MAX}}) = 99.9\%$.
2. A jitter shift value, termed "K", is calculated from the distribution equal to $\text{pr}(P_{\text{CURRENT}})$.
3. A random value is drawn, "J", from the range of $K - \text{jitter}$ to $K + \text{jitter}$.
4. Any value which falls outside the 0-1 range (in the cumulative normal space) is mapped back from the bound to a point one-tenth of the way from the bound to the initial value.
5. J is a new cumulative normal probability value.
6. Calculate a new parameter value, P_{JITTERED} , such that $\text{pr}(P_{\text{JITTERED}}) = J$.

In SS, the jitter fraction defines a uniform distribution in cumulative normal space +/- the jitter fraction from the initial value (in cumulative normal space). The normal distribution for each parameter, for this purpose, is defined such that the minimum bound is at 0.001, and the maximum at 0.999 of the cumulative distribution. If the jitter fraction and original initial value are such that a portion of the uniform distribution goes beyond 0.001 or 0.999 of the cumulative normal, the new value is set to one-tenth of the way from the bound to the original initial value.

Therefore $\text{sigma} = (\text{max} - \text{min}) / 6.18$. For parameters that are on the log-scale, sigma may be the correct measure of variation for jitters, for real-space parameters, CV (= $\text{sigma} / \text{original initial value}$) may be a better measure.

If the original initial value is at or near the middle of the min-max range, then for each 0.1 of jitter, the range of jitters extends about 0.25 sigmas to either side of the original value (though as the total jitter increases the relationship varies more than this), and the average absolute jitter is about half of that. For values far from the middle of the min-max range, the resulting jitter is skewed in parameter space, and may hit the bound, invoking the resetting mentioned above.




Figure 1: Illustration of the jitter algorithm

To evaluate the jittering, the bounds, and the original initial values, a jitter_info table is available from r4ss, including sigma, CV and InitLocation columns (the latter referring to location within the cumulative normal – too close to 0 or 1 indicates a potential issue).

Note: parameters with $\text{min} \leq -99$ or $\text{max} \geq 999$ are not jittered to avoid unreasonable values (a warning is produced to indicate this).

Parameter Priors

Priors on parameters fulfill two roles in SS. First, for parameters provided with an informative prior, SS is receiving additional information about the true value of the parameter. This information works with the information in the data through the overall log likelihood function to arrive at the final parameter estimate. Second, diffuse priors provide only weak information about the value of a prior and serve to manage model performance during execution. For example, some selectivity parameters may become unimportant depending upon the values of other parameters of that selectivity function. In the double normal selectivity function, the parameters controlling the width of the peak and the slope of the descending side become redundant if the parameter controlling the final selectivity moves to a value indicating asymptotic selectivity. The width and slope parameters would no longer have any effect on the log likelihood, so they would have no gradient in the log likelihood and would

drift aimlessly. A diffuse prior would then steer them towards a central value and avoid them crashing into the bounds. Another benefit of diffuse priors is the control of parameters that are given unnaturally wide bounds. When a parameter is given too broad of a bound, then early in a model run it could drift into this tail and potentially get into a situation where the gradient with respect to that parameter approaches zero even though it is not at its global best value. Here the diffuse prior helps move the parameter back towards the middle of its range where it presumably will be more influential and estimable.

The options for parameter priors are described as a function of P_{val} , the value of the parameter for which a prior is being calculated, as well as the parameter bounds in the case of the beta distribution (P_{max} and P_{min}), and the input values for P_{prior} and P_{rSD} , which in some cases are the mean and standard deviation, but interpretation depends on the prior type. The Prior Likelihoods below represent the negative log likelihood in all cases.

Prior Types

Note that the numbering in SS v.3.30 is different from that used in SS v.3.24 (where confusingly -1 indicated no prior and 0 indicated a normal prior). The calculation of the negative log likelihood is provided below for each prior types, as a function of the following inputs:

P_{init}	The value of the parameter for which a prior is being calculated where init can either be the initial un-estimated value or the estimated value (3rd column in control or control.ss_new file)
P_{LB}	The lower bound of the parameter (1st column in control file)
P_{UB}	The upper bound of the parameter (2nd column in control file)
P_{PR}	The input value for the prior input (4th column in control file)
P_{PRSD}	The standard deviation input value for the prior (5th column in control file)

- **Prior Type = 0 = No prior applied**

In a Bayesian context this is equivalent to a uniform prior between the parameter bounds.

- **Prior Type = 1 = Symmetric beta prior**

The symmetric beta is scaled between parameter bounds, imposing a larger penalty near the bounds. Prior standard deviation of 0.05 is very diffuse and a value of 5.0 provides a smooth U-shaped prior. The prior input is ignored for this prior type.

$$\mu = -P_{PRSD} \cdot \ln \left(\frac{P_{UB} + P_{LB}}{2} - P_{LB} \right) - P_{PRSD} \cdot \ln(0.5) \quad (42)$$

$$\begin{aligned} \text{Prior Likelihood} = & -\mu - P_{PRSD} \cdot \ln(P_{init} - P_{LB} + 0.0001) \\ & - P_{PRSD} \cdot \ln \left(1 - \frac{P_{init} - P_{LB} - 0.0001}{P_{UB} - P_{LB}} \right) \end{aligned} \quad (43)$$




Figure 2: The shape of the symmetric beta prior across alternative standard deviation values and the change in the negative log-likelihood.

- **Prior Type = 2 = Beta prior**

The definition of μ is consistent with CASAL's formulation with the β_{PR} and α_{PR} corresponding to the m and n parameters.

$$\mu = \frac{P_{\text{PR}} - P_{\text{LB}}}{P_{\text{UB}} - P_{\text{LB}}} \quad (44)$$

$$\tau = \frac{(P_{\text{PR}} - P_{\text{LB}})(P_{\text{UB}} - P_{\text{PR}})}{P_{\text{PRSD}}^2} - 1 \quad (45)$$

$$\beta_{\text{PR}} = \tau \cdot \mu \quad (46)$$

$$\alpha_{\text{PR}} = \tau(1 - \mu) \quad (47)$$

$$\begin{aligned}
\text{Prior Likelihood} = & (1 - \beta_{\text{PR}}) \cdot \ln(0.0001 + P_{\text{init}} - P_{\text{LB}}) \\
& + (1 - \alpha_{\text{PR}}) \cdot \ln(0.0001 + P_{\text{UB}} - P_{\text{init}}) \\
& - (1 - \beta_{\text{PR}}) \cdot \ln(0.0001 + P_{\text{PR}} - P_{\text{LB}}) \\
& - (1 - \alpha_{\text{PR}}) \cdot \ln(0.0001 + P_{\text{UB}} - P_{\text{PR}})
\end{aligned} \tag{48}$$

- **Prior Type 3 = Lognormal prior**

Note that this is undefined for $p \leq 0$ so the lower bound on the parameter must be > 0 . The prior value is input into the parameter line in natural log space while the initial parameter value is defined in normal space (e.g., init = 0.20, prior = -1.609438).

$$\text{Prior Likelihood} = \frac{1}{2} \left(\frac{\ln(P_{\text{init}}) - P_{\text{PR}}}{P_{\text{PRSD}}} \right)^2 \tag{49}$$

- **Prior Type 4 = Lognormal prior with bias correction**

This option allows the prior mean value to be entered as the $\ln(\text{mean})$. Note that this is undefined for $p \leq 0$ so the lower bound on the parameter must be > 0 .

$$\text{Prior Likelihood} = \frac{1}{2} \left(\frac{\ln(P_{\text{init}}) - P_{\text{PR}} + \frac{1}{2} P_{\text{PRSD}}^2}{P_{\text{PRSD}}} \right)^2 \tag{50}$$

- **Prior Type 5 = Gamma prior**

The lower bound should be 0 or greater.

$$\text{scale} = \frac{P_{\text{PRSD}}^2}{P_{\text{PR}}} \tag{51}$$

$$\text{shape} = \frac{P_{\text{PR}}}{\text{scale}} \tag{52}$$

$$\text{Prior Likelihood} = -\text{shape} \cdot \ln(\text{scale}) - \ln(\Gamma(\text{shape})) + (\text{shape} - 1) \cdot \ln(P_{\text{init}}) - \frac{P_{\text{init}}}{\text{scale}} \tag{53}$$

- **Prior Type 6 = Normal prior**

Note that this function is independent of the parameter bounds.

$$\text{Prior Likelihood} = \frac{1}{2} \left(\frac{P_{\text{init}} - P_{\text{PR}}}{P_{\text{PRSD}}} \right)^2 \tag{54}$$

Forecast Module: Benchmark and Forecasting Calculations

SS version 3.20 introduced substantial upgrades to the benchmark and forecast module. The general intent was to make the forecast outputs more consistent with the requirement to set catch limits that have a known probability of exceeding the overfishing limit. In addition, this upgrade addressed several inadequacies with the previous module, including:

- The average selectivity and relative F was the same for the benchmark and the forecast calculations;
- The biology-at-age in $\text{endyr}+1$ was used as the biology for the benchmark, but biology-at-age propagated forward in the forecast if there was time-varying growth;
- The forecast module had an inefficient approach to calculation of overfishing limit (OFL) conditioned on previously catching ABC;
- The forecast module implementation of catch caps was incomplete and applied some caps on a seasonally, rather than the more logical annual basis;
- The F_{mult} scalar for fishing intensity presented a confusing concept for many users;
- No provision for specification of catch allocation among fleets;
- The forecast allowed for a blend of fixed input catches and catches calculated from target F ; this is not optimal for calculation of the variance of F conditioned on a catch policy that sets annual catch limits (ACLs).

The V3.20 module addressed these issues by:

- Providing for unique specification of a range of years from which to calculate average selectivity for benchmark, average selectivity for forecast, relative F for benchmark, and relative F for forecast;
- Create a new specification for the range of years over which to average size-at-age and fecundity-at-age for the benchmark calculation. In a setup with time-varying growth, it may make sense to do this over the entire range of years in the time series. Note that some additional quantities still use their endyr values, notably the migration rates and the allocation of recruitments among areas. This will be addressed shortly;
- Create a multiple pass approach that rectifies the OFL calculation;

- Improve the specification of catch caps and implement specification of catch allocations so that there can be an annual cap for each fleet, an annual cap for each area, and an annual allocation among groups of fleets (e.g. all recreational fleets vs. all commercial fleets);
- Introduce capability to have implementation error in the forecast catch (single value applied to all fleets in all seasons of the year).

Multiple Pass Forecast

The most complicated aspect of the changes is with regard to the multiple pass aspect of the forecast. This multiple pass approach is needed to calculate both OFL and ABC in a single model run. More importantly, the multiple passes are needed in order to mimic the actual sequence of assessment-management action – catch over a multi-year period. The first pass calculates OFL based on catching OFL each year, so presents the absolute maximum upper limit to catches. The second pass forecasts a catch based on a harvest policy, then applies catch caps and allocations, then updates the F 's to match these catches. In the third pass, stochastic recruitment and catch implementation error are implemented and SS calculates the F that would be needed in order to catch the adjusted catch amount previously calculated in the second pass. With this approach, SS is able to produce improved estimates of the probability that F would exceed the overfishing F . In effect it is the complement of the P^* approach. Rather than the P^* approach that calculates the stream of annual catches that would have an annual probability of $F > F_{limit}$, SS calculates the expected time series of P^* that would result from a specified harvest policy implemented as a buffer between F_{target} and F_{limit} .

The sequence of multiple forecast passes is as follows:

1. Pass 1 (a.k.a. Fcast_Loop1)
 - (a) Loop Years
 - i. SubLoop (a.k.a. ABC_Loop) = 1
 - A. $R = f(SSB)$ with no deviations
 - B. $F = F_{limit}$
 - C. Fixed input catch amounts ignored
 - D. No catch adjustments (caps and allocations)
 - E. No implementation error
 - F. Result: OFL conditioned on catching OFL each year
2. Pass 2

(a) Loop Years

i. SubLoop = 1

- A. $R=f(SSB)$ with no deviations
- B. $F=F_{limit}$
- C. Fixed input catch amounts ignored
- D. No catch adjustments (caps and allocations)
- E. No implementation error
- F. Result: OFL conditioned on catching ABC previous year. Stored in `std_vector`.

ii. SubLoop = 2

- A. $R=f(SSB)$ with no deviations
- B. $F=F_{target}$ to get catch for each fleet in each season
- C. Fixed input catch amounts replace catch from step 2
- D. Catch adjustments (caps and allocations) applied on annual basis (after looping through seasons and areas within this year). These adjustments utilize the logistic joiner approach common in SS so the overall results remain completely differentiable.
- E. No implementation error
- F. Result: ABC as adjusted for caps and allocations

iii. SubLoop = 3

- A. $R=f(SSB)$ with no deviations
- B. Catches from Pass 2 multiplied by the random term for implementation error
- C. F =adjusted to match the catch*error while taking into account the random recruitments. This is most easily visualized in a MCMC context where the recruitment deviation and the implementation error deviations take on non-zero values in each instance. In MLE, because the forecast recruitments and implementation error are estimated parameters with variance, their variance still propagates to the derived quantities in the forecast.

D. Result: Values for F, SSB, Recruitment, Catch are stored in std-vectors

- In addition, the ratios F/F_{limit} and SSB/SSB_{limit} or SSB/SSB_{target} are also stored in std_vectors.
- Estimated variance in these ratios allows calculation of annual probability that $F > F_{limit}$ or $B < B_{limit}$. This is essentially the realized P^* conditioned on the specified harvest policy.

Example Effects on Correlations

An example that illustrates the above process was conducted. The situation was a low M, late-maturing species, so changes are not dramatic. The example conducted a 10 year forecast and examined correlations with derived quantities in the last year of the forecast. This was done once with the full set of 3 passes as described above, and again with only 2 passes and stochastic recruitment occurring in pass 2, rather than 3. This alternative setup is more similar to forecasts done using previous model versions.

2 Forecast Passes with F from ABC and random recruitment			2 Forecast Passes with catch from target F and equilibrium recruitment		
Factor X	Factor Y	Corr	Factor X	Factor Y	Corr
A1 F 2011	RecrDev 2002	-0.126	A2 F 2011	RecrDev 2002	0.090
B1 F 2011	Recr 2002	0.312	B2 F 2011	Recr 2002	0.518
C1 ForeCatch 2011	RecrDev 2002	0.000	C2 ForeCatch 2011	RecrDev 2002	0.129
D1 ForeCatch 2011	Recr 2002	0.455	D2 Forecatch 2011	Recr 2002	0.555

Correlation A2 shows a small positive correlation between the recruitment deviation in 2002 and the F in 2011. This is probably due to the fact that a positive deviation in recruitment in 2002 will reduce the chances that the biomass in 2011 will be below the inflection point in the control rule. This occurs because in calculating catch from F, the model effectively "knows" the future recruitments. I predict that this B1 correlation would be near zero if there was no inflection in the control rule.

Correlation A1 shows this turning into a negative correlation. This is because the future catches are first calculated from equilibrium recruitment, then when random recruitments are implemented, a positive recruitment deviation will cause a negative deviation in the F needed to catch that now "fixed" amount of future catch.

Correlations B1 and B2 are in terms of absolute recruitment, not recruitment deviation. Now overall model conditions that cause a higher absolute recruitment level will also result in a higher forecast level. No surprise there, and the correlation is stronger when variance is based on catch is calculated from F (B2).

Correlation C2 shows a positive correlation between recruitment deviation in 2002 and forecast catch in 2011. However, correlation C1 is 0.0 because the forecast catch in 2011 is set based on equilibrium recruitment and is not influenced by the recruitment deviations.

Future Work

- More testing with high M , rapid turnover conditions
- Testing without inflection in control rule
- Consider separating implementation error into a pass #4 so results will more clearly show effect of assessment uncertainty separate from implementation uncertainty
- Consider adding a random "assessment" error which essentially is a random variable that scales population abundance before passing into the forecast stage. Complication is figuring out how to link it to the correlated error in the benchmark quantities
- Because all of these calculations occur only in the standard deviation phase (sdphase) or the MCMC evaluation (mceval) phase, it would be feasible for mceval calls to add an additional pass that is implemented many times and in which random forecast recruitment draws are made.
- Factors like selectivity and fleet relative F levels are calculated as an average of these values during the time series. This is internally consistent if these factors do not vary during the time series (although clearly this is a stiff model that will underestimate process variance). However, if these factors do vary over time, then the average used for the forecast will under-represent the variance. A better approach would be to set up the parameters of selectivity as a random process that extends throughout the forecast period, and to update estimated selectivity in each year of the forecast based upon the random realization of these parameters.

Fishing Mortality in Stock Synthesis

The implementation and reporting of fishing mortality rate, F , in SS has some aspects that can be confusing. This description provides an overview of the ways in which F is calculated, used, and reported.

Rationale

Fishery management systems expect to have a measure of annual fishing mortality that describes the intensity of the fishery such that an optimal level of F can be articulated and accountability measures can be invoked if F is too high, e.g., overfishing. This concept is simple and straightforward if the model is a simple biomass dynamics such that a single annual F value operates on the entirety of a non-age structured population. It also is simple for age-structured models that have a single fishing fleet and knife-edge selectivity beginning at some specified age.

The simplicity of F disappears quickly as models invoke a variety of realistic complexities such as: allowing the F to differ among ages or to be based on size; using a collection of fleets with different F levels and different age patterns for F ; spreading the population across areas and allowing different fleets with different F among the areas. An unambiguous measure of annual fishing intensity that represents the cumulative effect of all that complexity has not been defined. This problem has not been solved with SS, but some logical alternatives have been made available.

Nomenclature

The nomenclature below ignores sex, morphs and areas for simplicity. The quantities associated with F calculations are defined as:

f is fleet.

t is a time step; continuous across years y and seasons s ; equivalent to year if only 1 season.

a is age.

$C_{t,f}$ is fleet-specific catch in a time step.

$B_{t,f}$ is fleet specific available biomass, e.g., total biomass filtered by fleet-specific age selectivity, $s_{t,f,a}$.

$s_{t,f,a}$ is age-specific selectivity for a fleet. If selectivity is length-specific, then age-specific selectivity is calculated as the dot product across length bins of length selectivity and the normal (or lognormal) distribution of length-at-age. If selectivity is both length- and age-based, which is an entirely normal concept in SS, then age selectivity due to length selectivity is calculated first, then multiplied by the direct age selectivity. This compound age selectivity is used in the mortality calculations and is reported as $asel2$ in SS output. See appendix to Methot and Wetzel (2013) for more detail on this.

$F'_{t,f}$ is the apical fishing mortality for a fleet. This means that it is the rate for the age that has selectivity equal to 1.0. If your model is using F' 's as parameters, then the parameter values are for F' .

$F_{t,f,a}$ is age and fleet-specific fishing mortality rate equal to $F'_{t,f} * s_{t,f,a}$. Note that it is possible for no age to have a selectivity equal to 1.0. In this case, F' is still the rate for the hypothetical age that has selectivity equal to 1.0. The reported F' values are not rescaled to be an F for the age with peak selectivity. Users need to take this into account if they are comparing reported F' values to reported vector of $F_{t,f,a}$ values.

$annF_y$ is a measure of the total fishing intensity for a year, based on one of several user-specified options (see below).

$Fstd_y$ is a standardized measure of the total fishing intensity for a year and is reported in the derived quantities, so variance is calculated for this quantity. See below for how it relates to $annF$.

Terminology and reporting of $annF$ and $Fstd$ has been slightly revised for clarity in 3.30.15.00 and the description here follows the new conventions.

F Calculation

SS allows for three approaches to estimate the F' that will match the input values for retained catch. Note that SS is calculating the F' to match the retained catch conditional on the fraction of total catch that is retained, e.g., the total catch is partitioned into retained and discarded portions.

1. Pope's method decays the numbers-at-age to the middle of the season, calculates a harvest rate for each fleet, $H_{t,f}$, that is the ratio of $C_{t,f}$ to $B_{t,f}$, then decays the survivors to the end of the season. SS then calculates the total mortality, $Z_{t,a}$, from the ratio of survivors to initial numbers. The Z is subsequently used for in-season interpolation to get expected values for observations.
2. F as parameters uses the standard Baranov catch equation and lets ADMB find the F' parameter values that produce the lowest negative log-likelihood, which includes fit to the input catch data. F as parameters method tends to work better than Pope's or hybrid in high F situations because it allows for some lack of fit to catch levels in early iterations and can later improve this fit as it closes in on the best solution.
3. Hybrid F starts by calculating a harvest rate, H , using Pope's, then converts these H values, which have units of fractional harvest rate, into an approximate of F' in exponential units, then SS tunes these F' values over a few iterations to get a better match to each fleet's catch.

Items to note:

- SS includes a permutation on the F as parameters method. In the first few phases, SS uses hybrid, then between phases it converts these directly calculated F' values into parameters and proceeds in subsequent phases and MCMC to use the parameter approach. This variation on the parameter method is the recommend approach in high F situations.

- With Pope's method, the H values are fraction caught, so duration of the season does not matter. Parameter and hybrid treat F' identically and multiply the F' values by season duration (which has units of fraction of a year) as it is used. Each of the F methods ends up with a $Z_{t,f}$ that is used for in-season interpolation.

Relative F and F_{mult}

The F' is fleet-specific, so it is useful to have a concept of relative F , $\text{rel}F_f$, among fleets. In SS, $\text{rel}F_f = F'_{t,f} / \sum_f F'_{t,f}$ for a single time period t . In the benchmark and forecast routines, SS can calculate $\text{rel}F_f$ using $F'_{t,f}$ over a range of years, or the user can input custom $\text{rel}F$ values for benchmark and forecast in the forecast.ss file. Note that in a multi-season model setup, $\text{rel}F_f$ is implemented as $\text{rel}F_{s,f}$ where s is the season. These get multiplied by season duration as they are used.

In the benchmark section of the code, SS searches for an F_{mult} to achieve various management reference points (often referred to as benchmarks). In this search, SS calculates a benchmark F as $F'_{\text{ben},f} = F_{\text{mult}} * \text{rel}F_f$, then calculates equilibrium yield and spawning biomass per recruit (SPR). SS searches for the F_{mult} that satisfies the search conditions, first for user-specified SPR, then for user-specified spawning biomass at a management target (B_{TGT} or $F_{0.1}$), then for MSY. The resultant benchmark quantities are reported in the derived quantities, but F_{mult} and $F'_{\text{ben},f}$ are only reported in the Forecast_report.sso file. SS stores the benchmark F_{mult} values so that user can invoke them for the forecast.

Annual F

The $\text{ann}F$ is a single annual value across all fleets and areas according to $F_{\text{report_units}}$, which is specified by users in the starter file. If there are many fleets, across several areas and with very different selectivity patterns, $\text{ann}F$ can have a complicated relationship to apical F . The $F_{\text{report_units}}$ specification in the starter.ss file, see example line below, allows user to calculate it using F' directly, use exploitation rate, or be derived from Z -at-age.

Example F reporting unit specification in the starter.ss file:

5	# F_report_units:
	0 = skip;
	1 = exploitation(Bio);
	2 = exploitation(Num);
	3 = sum(Frates);
	4 = true F for range of ages;
	5 = unweighted avg. F for range of ages.

3 7	# min and max age over which average F will be calculated
-----	---

For options 4 and 5 of `F_report_units`, the F is calculated as $Z - M$ where Z is calculated as $\ln(N_{t+1,a+1}/N_{t,a})$, thus Z subsumes the effect of F .

The $\text{ann}F$ is calculated for each year of the estimated time series and of the forecast. Additionally, an $\text{ann}F$ is calculated in the benchmark calculations to provide equilibrium values that have the same units as $\text{ann}F$ from the time series. In versions previous to 3.30.15, it was labeled inaccurately as F_{std} in the output, not $\text{ann}F$. For example, in the Management Quantities section of derived quantities prior to 3.30.15, there is a quantity labeled `Fstd_Btgt`. This is more accurately labeled as the annual F associated with the biomass target, `ann_F_Btgt`, in 3.30.15.

F_{std}

F_{std} is a single annual value based on $\text{ann}F$ and the relationship to $\text{ann}F$ is specified by `F_report_basis` in the `starter.ss` file. The benchmark $\text{ann}F$ may be used to rescale the time series of $\text{ann}F$'s to become a time series of standardized values representing the intensity of fishing, F_{std} . The report basis is selected in the starter file as:

0	# <code>F_report_basis</code> :
	0 = raw F report;
	1 = F / F_{SPR} ;
	2 = F / F_{MSY} ;
	3 = F / F_{BTGT} .

For example, if user selects option 1, F / F_{SPR} , SS will take the time series of $\text{ann}F$ and divide each value by the $\text{ann}F$ calculated in benchmark.

Units for SS inputs related to F

Below is a list of items to consider in terms of units for F in SS:

- If `F_ballpark` is specified in the `control.ss` file, its units are the same as $\text{ann}F$, so is not fleet-specific.
- F as parameter values has units of fleet-specific apical F' .
- In the `forecast.ss` file there is an option to input a vector of $\text{rel}F$ values. These are dimensionless and will be rescaled by SS to sum to 1.0.

- In the forecast.ss file there is an option to specify an F scalar for the forecast. The units of F scalar are the same as the F_{mult} values calculated in benchmark. There are a full set of options for forecast F scalar that can be selected in the forecast file. If the forecast F scalar is set as F_{SPR} , then SS will use SPR_Fmult calculated in benchmark and reported in Forecast-report.sso. If user selects the option to input an annual F scalar, option 5, then the value is input on a following line. Whichever method the user selects for forecast F scalar (F_{mult}), SS will start the forecast by creating a fleet-specific vector of apical F values from $F_{\text{mult}} \cdot \text{rel}F_f$.
- Also in the forecast.ss file, the last section of inputs allows for input of time and fleet specific apical $F'_{t,f}$ values that override the basic forecast F specification described above.

Guidance on Population Dynamics Modeling

There are numerous assessment related topics that have arisen among SS users. In an attempt to capture information on these topics and how they may be related to the use of SS a collection of general information is provided here.

Data Weighting

In 2015 there was a CAPAM workshop dedicated to data-weighting. Description of the workshop can be found on the CAPAM website. The presentations from the workshop are available through that website and many of them were included in a special issue of Fisheries Research.

Currently, there are three main methods for weighting length and data applied for U.S. West Coast assessments using Stock Synthesis.

1. **McAllister - Ianelli:** Effective sample size is calculated from fit of observed to expected length or age compositions. Tuning algorithm is intended to make the arithmetic mean of the input sample size equal to the harmonic mean of the effective sample size (McAllister and Ianelli, 1997).
2. **Francis:** Based on variability in the observed mean length or age by year, where the sample sizes are adjusted such that the fit of the expected mean length or age should fit within the uncertainty intervals at a rate which is consistent with variability expected based on the adjusted sample sizes (Method "TA1.8") (Francis and Hilborn, 2011)
3. **Dirichlet-Multinomial:** A new likelihood (as opposed to the standard multinomial) which includes an estimable parameter (theta) which scales the input sample size. In this case, the term "Effective sample size" has a different interpretation than in the McAllister-Ianelli approach (Thorson et al., 2017).

Applying the methods

McAllister-Ianelli

The "Length_Comp_Fit_Summary" and "Age_Comp_Fit_Summary" sections in the Report file include information on the harmonic mean of the effective sample size and arithmetic mean of the input sample size used in this tuning method. In the r4ss package, these tables are returned by the SS_output function as `$Length_comp_Eff_N_tuning_check` and `$Age_comp_Eff_N_tuning_check`.

A convenient way to process these values into the format required by the control file is to use the function:

```
SS_tune_comps(replist, option = "MI")
```

where the input "replist" is the object created by SS_output. This function will return a table and also write a matching file called "suggested_tuning.ss" to the directory where the model was run.

For models using SS version 3.30, the table created by SS_tune_comps can be pasted into the bottom of the control file in the section labeled "Input variance adjustments", followed by the terminator line which indicates the end of the section.

Also see the help page for the r4ss function SS_varadjust which can be used to automatically write a new control file if you want to streamline the process of applying multiple iterations of this tuning method.

If the tuning has been implemented, the green lines in the figure below would approximately intersect at a point which is on the black 1-to-1 diagonal line in this figure created by the r4ss function SS_plots.




Figure 3: The relationship between the observed sample size (the input sample number) versus the effective sample size where the effective sample size is the product of the input sample size and the data weighting applied to the data set.

There are a couple of challenges posed by the McAllister-Ianelli data-weighting approach:

1. Subjective choice of how many iterations to take to achieve adequate convergence. Often just one iteration is applied.
2. Takes time to implement so tuning is rarely repeated during retrospective or sensitivity analyses.

Francis

Implementation: recommended adjustments are calculated by the `r4ss` functions `SSMethod.TA1.8` and `SSMethod.Cond.TA1.8`. These functions are rarely used alone but are called by the `SS_plots` function when making plots like the one below. For `SS v.3.30` models, the simplest way to get the adjustments in the format required by the control file is to use the `SS_tune_comps` function (described above under the McAllister-Ianelli method), but with a different option specified:

```
SS_tune_comps(replist, option = "Francis")
```

The figure below shows the estimated 95% intervals around the observed mean length by year based on the input sample size (thick lines) and the increase in that uncertainty which would occur if the sample sizes were adjusted according to the proposed multiplier.




Figure 4: The mean length of the length samples for each year from Fleet 1 (called the MexCal S1 NSP fleet) with 95% confidence intervals based on current samples sizes using the Francis data weighting method (referred to as TA1.8). Thinner intervals with capped ends show result of further adjusting sample sizes based on suggested multiplier, 0.2739, with 95% interval (0.1661-0.6305) for length data from Fleet 1. The blue line shows the expected variation of the mean length across years.

There are a several of challenges posed by the Francis data-weighting approach:

1. Subjective choice of how many iterations to take to achieve adequate convergence. Often just one iteration is applied.
2. Takes time to implement so tuning is rarely repeated during retrospective or sensitivity analyses.
3. Recommended adjustment can be sensitive to outliers (remove a few years of anomalous composition data can lead to large change in recommended adjustment).

Dirichlet-Multinomial

Change the choice of likelihood and set parameter choices in the data file:

- In the specification of the length and/or age data, change "CompError" column in age and length comp specifications from 0 to 1 and "ParmSelect" from 0 to a sequence of numbers from 1 to N where N is the total number of combinations of fleet and age/length.

- Resulting input should look similar to:

```
#_mintailcomp addtocomp combM+F CompressBins CompError ParmSelect minsamplesize
-1 0.001 0 0 1 1 0.01 #_fleet:1_Fishery
-1 0.001 0 0 1 2 0.01 #_fleet:2_ASHOP
-1 0.001 0 0 0 0 0.01 #_fleet:3_Foreign
-1 0.001 0 0 1 3 0.01 #_fleet:4_POP
-1 0.001 0 0 1 4 0.01 #_fleet:5_Triennial
```

- The ParmSelect column can also have repeated values so that multiple fleets share the same log(theta) parameter.
- If you have both length and age data, the ParmSelect should have separate numbers for each, e.g. 1 and 2 for the length comps and 3 and 4 for the age comps for the same two fleets.

Add parameter lines to the control file:

- Add as many parameter lines as the maximum numbers in the ParmSelect column. The new parameter lines go after the main selectivity parameters but before any time-varying selectivity parameters
- Jim Thorson initially recommended bounds of -5 to 20, with a starting value of 0 (which corresponds to a weight of about 50% of the input sample size). However, parameter estimates above 5 are associated with 99-100% weight with little information in the likelihood about the parameter value. Therefore, an upper bound of 5 may help identify cases that otherwise would have convergence issues.
- In consultation with Jim Thorson, Ian Taylor proposed a normal $N(0, 1.813)$ prior for the $\ln(DM_parm)$ parameters which to counteract the effect of the logistic transformation

between this parameter and the data weighting. The 1.813 value was calculated as the standard deviation of the distribution of $\log(\theta)$ values derived from starting with a uniform distribution on the weights, $weight = \theta/(1 + \theta) U(0, 1)$, and solving for $\log(\theta)$. This prior was applied to the 2020 Pacific Hake stock assessment (Grandin et al., 2020) and led to better MCMC convergence with relatively little impact on the maximum likelihood estimates.

- Example parameter lines are below (note that the labels have changed in SS version 3.30.15 from `ln(EffN_mult)` to `ln(DM_parm)`):

```
# Dirichlet parameters
#_LO HI INIT PRIOR PR_SD PR_type PHASE ... Blk_Fxn # Parm_name
-5 20 -0.2 0 50 0 2 ... 0 # ln(EffN_mult)_fishery - Dirichlet
-5 20 -0.2 0 50 0 2 ... 0 # ln(EffN_mult)_ashop - Dirichlet
-5 20 -0.2 0 50 0 2 ... 0 # ln(EffN_mult)_pop - Dirichlet
-5 20 -0.2 0 50 0 2 ... 0 # ln(EffN_mult)_tri - Dirichlet
-5 20 -0.2 0 50 0 2 ... 0 # ln(EffN_mult)_afsc - Dirichlet
```

- Reset any existing variance adjustments factors that might have been used for the McAllister-Ianelli or Francis tuning methods. In 3.24 this means setting the values to 1, in SS version 3.30, you can simply delete or comment-out the rows with the adjustments.

The `SS_output` function in `r4ss` returns table like the following:

```
$Dirichlet_Multinomial_pars
      Value Phase Min Max      Theta Theta/(1+Theta)
ln(EffN_mult)_1 12.7638000      2 -5 15 3.493397e+05 0.9999971
ln(EffN_mult)_2  2.5263600      2 -5 15 1.250789e+01 0.9259692
ln(EffN_mult)_3 11.6519000      2 -5 15 1.149095e+05 0.9999913
ln(EffN_mult)_4  1.9373900      2 -5 15 6.940612e+00 0.8740651
ln(EffN_mult)_5 10.2865000      2 -5 15 2.933393e+04 0.9999659
ln(EffN_mult)_6 11.5556000      2 -5 15 1.043598e+05 0.9999904
```

The ratio shown in the final column is the estimated multiplier which can be compared to the sample size adjustment estimated in the other tuning methods above (the `New_Var_adj` column in the table produced by the `SS_tune_comps` function in `r4ss`).

If the reported $\theta/(1 + \theta)$ ratio is close to 1.0, that indicates that the model is trying to tune the sample size as high as possible. In this case, the $\ln(\theta)$ parameters should be fixed at a high value, like the upper bound of 20, which will result in 100% weight being applied to the input sample sizes. An alternative would be to manually change the input sample sizes to a higher value so that the estimated weighting will be less than 100

There are a several of challenges posed by the Dirichlet-Multinomial data-weighting approach:

1. Not yet widely used so limited guidance is available.

2. Does not allow weights above 100% (by design) so it is not yet clear how best to deal with the situation when the estimated weight is close to 100%.
3. Parameterization has potential to cause convergence issues or inefficient MCMC sampling when weights are close to 100% if no prior is applied as discussed above.

Recruitment Variability and Bias Correction

Recruitments in SS are defined as lognormal deviates around a log-bias adjusted spawner-recruitment curve. The magnitude of the log-bias adjustment is calculated from the level of σ_R , which is the standard deviation of the recruitment deviations (in log-space). There are 5 segments of the time series in which to consider the effect of the log-bias adjustment: virgin; initial equilibrium; early data-poor period; data-rich period; very-recent/forecast. The choice of break points between these segments need not correspond directly with the settings for the bias adjustment, although some alignment might be desired. Methot et al. (2011) provide more detailed discussion of the bias adjustment than what is provided below but do not address the separation of time periods into separate segments. The approach is illustrated with figures associated the assessment for darkblotched rockfish (Gertseva and Thorson, James T, 2013).




Figure 5: Spawner-recruitment relationship for darkblotched rockfish (Gertseva and Thorson, James T, 2013). Red points represent estimated recruitments, the solid black line is the stock-recruit relationship and the green line represents the adjustment to this relationship after adjustment to account for the lognormal distribution associated with each year. The "+" symbol labeled 1915 near the right side represents both the virgin and initial equilibrium of the model. The numerous red points close to the initial conditions correspond to the early years of the model with low harvest rates.




Figure 6: Time series of log recruitment deviations for darkblotched rockfish with 95% uncertainty intervals. The start year of the model is 1915, but recruitment deviations are estimated starting in 1870. The 45 deviation estimates for 1870–1914 inform the age structure used in the start year. The black color for the years 1960–2011 indicates the "main" recruitment deviation vector, while the blue color for the years 1870–1959 and 2012–2024 indicates the "early" and "late/forecast" recruitment deviation vectors, respectively.

Virgin Biomass

The R_0 level of recruitment is a parameter of the spawner-recruitment curve. This recruitment and the corresponding spawning biomass are expected to represent the long-term arithmetic mean.

Initial Equilibrium

The level of recruitment is typically maintained at the R_0 level even though the initial equilibrium catch will reduce the spawning biomass below the virgin level. If steepness is moderately low or the initial F is high, then the lack of response in recruitment level may appear paradoxical. The logic is that building in the spawner-recruitment response to initial F would significantly complicate the calculations and would imply that the initial equilibrium catch level had been going on for multiple generations. If the lack of response is considered to be problematic in a particular application, then start the model at an earlier year and with a lower initial equilibrium catch so that the dynamics of the spawner-recruitment response get captured in the early period, rather than getting lost in the initial equilibrium.

Early Data-Poor Period

This is the early part of the time series where the only data typically are landed catch. There are no data to inform the model about the specific year-to-year fluctuations in recruitment, although the ending years of this period will begin to be influenced by the data. The "early time period" is not a formal concept. It is up to the user to decide whether to start estimating recruitment deviations beginning with the first year of the model, or to delay such estimation until the data become more informative. Modeling recruitment deviations in this period may lead to a more realistic portrayal of the uncertainty in depletion, but can also lead to spurious patterns in estimated recruitments that may be driven by the fit to index data or other sources that would not be expected to have accurate information on recruitment.

- Option A: Do not estimate recruitment deviations during this early period. During years prior to the first year of recruitment deviations, the model will set the recruitment equal to the level of the spawner-recruitment curve. Thus, it is a mean-based level of recruitment. Because these annual parameters are fixed to the level of the spawner-recruitment curve, they have no additional uncertainty and make no contribution to the variance of the model. This approach may produce relatively large, or small, magnitude deviations at the very beginning of the subsequent period, as the model "catches up" to any slight signal that could not be captured through estimated deviations in the early data-poor period. There may be some effect on the estimate of R_0 as a result of lack of model flexibility in balancing early period removals with signal in the early portion of the data-rich period.
- Option B: Estimate recruitment deviations for all the early years. Each of these recruitment deviations is now a deviation parameter so will have a variance that contributes to the total model variance. The estimated standard deviation of each of these deviation parameters should be similar to σ_R because σ_R is the only constraint on these parameters (however, the last few in the sequence will begin to feel the effect of the data so may have lower standard deviations).

Data-Rich Period

Here the length and or age data inform the model on the year-to-year level of recruitment. These fluctuations in recruitment are assumed to have a lognormal distribution around the log-bias adjusted spawner-recruitment curve. The level of σ_R input to the model should match this level of fluctuation to a reasonable degree. Because the recruitments are lognormal, they produce a mean biomass level that is comparable to the virgin biomass and thus the depletion level can be calculated without bias. However, if the early period has recruitment deviations estimated by maximum posterior density, then the depletion levels during the early part of the data-rich period may have some lingering effect of negative bias during the early time period. The level of σ_R should be at least as large as the level of variability in these estimated recruitments. If too high a level of σ_R is used, then a bias can occur in the estimate of spawner-recruitment steepness, which determines the trend in recruitment. This occurs when the early recruitments are taken directly from the spawner-recruitment curve, so are mean unbiased, then the later recruitments are estimated as deviations from the log-bias adjusted

curve. If σ_R is too large, then the bias-adjustment is too large, and the model may compensate by increasing steepness to keep the mean level of recent recruitments at the correct level.

Recent Years/Forecast

Here the situation is very similar to the early time period in that there are no data to inform the model about the year-to-year pattern in recruitment fluctuations so all deviations will be pulled to a zero level in the maximum posterior density. The structure of SS creates no sharp dividing line between the estimation period and the forecast period. In many cases one or more recruitments at the end of the time series will lack appreciable signal in the data and should therefore be treated as forecast recruit deviations. To the degree that some variability is observed in these recruitments, partial or full bias correction may be desirable for these deviations separate from the purely forecast deviations, there is therefore an additional control for the level of bias correction applied to forecast deviations occurring prior to end year + 1.




Figure 7: Timeseries of standard error estimates for the log recruitment deviations for darkblotched rockfish with 95% uncertainty intervals. As in Figure 6, the black color indicates the main recruitment period. This period with lower standard error is associated with higher variability among deviations (Figure 6). The red line at 0.75 indicates the σ_R value in this model.




Figure 8: Transformation of the standard error estimates (shown in Figure 7) for darkblotched rockfish following the approach suggested by Methot et al. (2011). These values were used to set the 5 values controlling the degree of bias adjustment (as a fraction of $\sigma_R/2$) to account for differences in the mean and median of the lognormal distribution from which the recruitment deviations are drawn. The red line indicates a bias adjustment of 0 up to the 1960.75, ramping up to a maximum adjustment level of 0.877 for the period 1990.4–2008.98, and reducing back to 0 starting in 2013.08. Note that these values controlling the bias adjustment need not be integer year values. Also the break points in the bias adjustment function need not match the break points between early, main, and late/forecast recruitment deviation vectors (indicated by blue and black colors in Figures 6 and 7). The blue line indicates a functional form that minimizes the sum of squared differences between the bias adjustment function and the transformed standard error values. The subtle differences between red and blue lines are unlikely to have any appreciable effect on the model results.




Figure 9: Comparison of timeseries of spawning depletion for darkblotched rockfish models with early recruitment deviations (starting in 1870) and without early deviations (only main recruitment deviations starting in 1960). The point estimates are similar, but the 95% uncertainty intervals are substantially different. With no recruitment deviations for the early period, the estimates of spawning depletion in the early years are very precise and uncertainty increases as the stock moves into the data rich period. In contrast, the addition of the early recruitment deviations results in a large uncertainty in spawning depletion for the early years and an increase in precision as the stock moves into the data rich period. In this application, the uncertainty associated with the recent years is independent of the assumptions about early recruitments.

Issues with Including Environmental Effects

The expected level of recruitment is a function of spawning biomass, an environmental time series, and a log-bias adjustment.

$$E(Recruitment) = f(SpBio) * \exp(\beta * envdata) * \exp(-0.5 * \sigma_R^2) \quad (55)$$

σ_R is the variability of the deviations, so it is in addition to the variance "created" by the environmental effect. So, as more of the total recruitment variability is explained by the environmental effect, the residual σ_R should be decreased. The model does not do this automatically.

The environmental effect is inherently lognormal. So when an environmental effect is included in the model, the arithmetic mean recruitment level will be increased above the level predicted by $f(SpBio)$

alone. The consequences of this have not yet been thoroughly investigated, but there probably should be another bias correction based on the variability of the environmental data as scaled by the estimated linkage parameter, β . It is also problematic that the environmental effect time series used as input is assumed to be measured without error.

The preferred approach to including environmental effects on recruitment is not to use the environmental effect in the direct calculation of the expected level of recruitment. Instead, the environmental data would be used as if it was a survey observation of the recruitment deviation. This approach is similar to using the environmental index as if it was a survey of age 0 recruitment abundance because by focusing on the fit to the deviations it removes the effect of SpBio on recruitment. In this alternative, the σ_R would not be changed by the environmental data; instead the environmental data would be used to explain some of the total variability represented by σ_R . This approach may also allow greater uncertainty in forecasts, as the variability in projected recruitments would reflect both the uncertainty in the environmental observations themselves and the model fit to these observations.

Initial Age Composition

If the first year with recruitment deviations is set less than the start year of the model, then these early deviations will modify the initial age composition. The amount of information on historical recruitment variability certainly will degrade as the model attempts to estimate deviations for older age groups in the initial equilibrium. So the degree of bias correction is reduced linearly in proportion to age so that the correction disappears when maximum age is reached. The initial age composition approach normally produces a result that is indistinguishable from a configuration that starts earlier in the time series and estimates a longer time series of recruitments. However, because the initial equilibrium is calculated from a recruitment level unaffected by spawner-recruitment steepness and initial age composition adjustments are applied after the initial equilibrium is calculated, it is possible that the initial age composition approach will produce a slightly different result than if the time series was started earlier and the deviations were being applied to the recruitment levels predicted from the spawner-recruitment curve.

Code Examples

Ageing Error Estimation

```
/* SS_Label_FUNCTION 45 get_age_age */
FUNCTION void get_age_age(const int Keynum, const int AgeKey_StartAge, const int AgeKey_Linear1, const int AgeKey_Linear2)
{
  // FUTURE: calculate adjustment to oldest age based on continued ageing of old fish
  age_age(Keynum).initialize();
  dvariable age;
  dvar_vector age_err_parm(1,7);
  dvariable temp;

  if(Keynum==Use_AgeKeyZero)
  {
    // SS_Label_45.1 set age_err_parm to mgp_adj, so can be time-varying according to MGparm options
    for (a=1;a<=7;a++)
    {age_err_parm(a)=mgp_adj(AgeKeyParm-1+a);}
    age_err(Use_AgeKeyZero,1)(0, AgeKey_StartAge)=r_ages(0, AgeKey_StartAge)+0.5;
    age_err(Use_AgeKeyZero,2)(0, AgeKey_StartAge)=age_err_parm(5)*(r_ages(0, AgeKey_StartAge)+0.5)/
    (age_err_parm(1)+0.5);
    // SS_Label_45.3 calc ageing bias
    if(AgeKey_Linear1==0)
    {
      age_err(Use_AgeKeyZero,1)(AgeKey_StartAge,nages)=0.5 + r_ages(AgeKey_StartAge,nages) +
      age_err_parm(2)+(age_err_parm(3)-age_err_parm(2))*(1.0-mfexp(-age_err_parm(4)*
      (r_ages(AgeKey_StartAge,nages)-age_err_parm(1)))) / (1.0-mfexp(-age_err_parm(4)*
      (r_ages(nages)-age_err_parm(1))));
    }
    else
    {
      age_err(Use_AgeKeyZero,1)(AgeKey_StartAge,nages)=0.5 + r_ages(AgeKey_StartAge,nages) +
      age_err_parm(2)+(age_err_parm(3)-age_err_parm(2))*
      (r_ages(AgeKey_StartAge,nages)-age_err_parm(1))/(r_ages(nages)-age_err_parm(1));
    }
    // SS_Label_45.4 calc ageing variance
    if(AgeKey_Linear2==0)
    {
      age_err(Use_AgeKeyZero,2)(AgeKey_StartAge,nages)=age_err_parm(5)+(age_err_parm(6)-age_err_parm(5))*
      (1.0-mfexp(-age_err_parm(7)*(r_ages(AgeKey_StartAge,nages)-age_err_parm(1)))) /
      (1.0-mfexp(-age_err_parm(7)*(r_ages(nages)-age_err_parm(1))));
    }
    else
    {
      age_err(Use_AgeKeyZero,2)(AgeKey_StartAge,nages)=age_err_parm(5)+(age_err_parm(6)-age_err_parm(5))*
      (r_ages(AgeKey_StartAge,nages)-age_err_parm(1))/(r_ages(nages)-age_err_parm(1));
    }
  }
}
```

Survival Based SRR Code

Code for the survival based recruitment is shown below:

```
// SS_Label_43.3.7 survival based
case 7: // survival based, so constrained such that recruits cannot exceed fecundity
{
```

```

SRZ_0=log(1.0/(SSB_virgin_adj/Recr_virgin_adj));
SRZ_max=SRZ_0+SR_parm_work(2)*(0.0-SRZ_0);
SRZ_surv=mfexp((1.-pow((SSB_curr_adj/SSB_virgin_adj),SR_parm_work(3)))*(SRZ_max-SRZ_0)+SRZ_0); // survival
NewRecruits=SSB_curr_adj*SRZ_surv;
exp_rec(y,1)=NewRecruits; // expected arithmetic mean recruitment

// SS_Label_43.3.7.1 Do variation in recruitment by adjusting survival
if(recdev_cycle>0)
{
gg=y - (styr+(int((y-styr)/recdev_cycle))*recdev_cycle)+1;
SRZ_surv*=mfexp(recdev_cycle_parm(gg));
}
exp_rec(y,2)=SSB_curr_adj*SRZ_surv;
SRZ_surv*=mfexp(-biasadj(y)*half_sigmaRsq); // bias adjustment
exp_rec(y,3)=SSB_curr_adj*SRZ_surv;
if(y <=recdev_end)
{
if(recdev_doit(y)>0) SRZ_surv*=mfexp(recdev(y)); // recruitment deviation
}
else if(Do_Forecast>0)
{
SRZ_surv *= mfexp(Fcast_recruitments(y));
}
join=1./(1.+mfexp(100*(SRZ_surv-1.)));
SRZ_surv=SRZ_surv*join + (1.-join)*1.0;
NewRecruits=SSB_curr_adj*SRZ_surv;
exp_rec(y,4) = NewRecruits;
break;
}

```

Random Walk Selectivity: Pattern 17

Code for selectivity pattern 17, random walk shown below:

```

// SS_Label_Info_22.7.17 #age selectivity: each age has parameter as random walk
// #41 each age has parameter as random walk scaled by average of values at low age through high age
// transformation as selex=exp(parm); some special codes */
case 41:
scaling_offset = 2;
case 17: //
{
lastsel=0.0; // value is the change in log(selex); this is the reference value for age 0
tempvec_a=-999.;
tempvec_a(0)=0.0; // so do not try to estimate the first value
int lastage;
if(seltype(f,4)==0)
{lastage=nages;}
else
{lastage=abs(seltype(f,4));}

for (a=1;a<=lastage;a++)
{
// with use of -999, lastsel stays constant until changed, so could create a linear change in ln(selex)
// use of (a+1) is because the first element, sp(1), is for age zero
if(sp(a+1+scaling_offset)>-999.) {lastsel=sp(a+1+scaling_offset);}
tempvec_a(a)=tempvec_a(a-1)+lastsel; // cumulative log(selex)
}
if (scaling_offset == 0)
{
temp=max(tempvec_a); // find max so at least one age will have selex=1.
}
}

```

```

}
else
{
int low_bin = int(value(sp(1)));
int high_bin = int(value(sp(2)));
if (low_bin < 0)
{
low_bin = 0;
N_warn++; warning<<" selex pattern 41; value for low bin is less than 0, so set to 0 "<<endl;
}
if (high_bin > nages)
{
high_bin = nages;
N_warn++; warning<<" selex pattern 41; value for high bin is greater than "<<nages<<" , so set to "<<nages<<" "<<endl;
}
if (high_bin < low_bin) high_bin = low_bin;
if (low_bin > high_bin) low_bin = high_bin;
sp(1) = low_bin;
sp(2) = high_bin;
temp=mean(tempvec_a(low_bin,high_bin));
}
sel_a(y,fs,1)=mfexp(tempvec_a-temp);
a=0;
while(sp(a+1+scaling_offset)==-1000) // reset range of young ages to selex=0.0
{
sel_a(y,fs,1,a)=0.0;
a++;
}
scaling_offset = 0; // reset scaling offset
if(lastage<nages)
{
for (a=lastage+1;a<=nages;a++)
{
if(seltype(f,4)>0)
{sel_a(y,fs,1,a)=sel_a(y,fs,1,a-1);}
else
{sel_a(y,fs,1,a)=0.0;}
}
}
break;
}
}

```

Cubic Spline Selectivity

Code for cubic spline selectivity, option 42, shown below:

```

// SS_Label_Info_22.7.27 #age selectivity: cubic spline
// #42 cubic spline scaled by average of values at low age through high age
case 42:
scaling_offset = 2;
case 27:
{
k=seltype(f,4); // n points to include in cubic spline
for (i=1;i<=k;i++)
{
splineX(i)=value(sp(i+3+scaling_offset)); // "value" required to avoid error, but values should be always fixed anyway
splineY(i)=sp(i+3+k+scaling_offset);
}
z=nages;
while(r_ages(z)>splineX(k)) {z--;}
}

```

```

j2=z+1; // first age beyond last node
vcubic_spline_function splinefn=vcubic_spline_function(splineX(1,k),splineY(1,k),sp(2+scaling_offset),sp(3+scaling_offset));
tempvec_a= splinefn(r_ages); // interpolate selectivity at each age
if (scaling_offset == 0)
{
temp=max(tempvec_a(0,j2));
}
else
{
int low_bin = int(value(sp(1)));
int high_bin = int(value(sp(2)));
if (low_bin < 0)
{
low_bin = 0;
N_warn++; warning<<" selex pattern 42; value for low bin is less than 0, so set to 0 "<<endl;
}
if (high_bin > nages)
{
high_bin = nages;
N_warn++; warning<<" selex pattern 42; value for high bin is greater than "<<nages<<", so set to "<<nages<<" "<<endl;
}
if (high_bin < low_bin) high_bin = low_bin;
if (low_bin > high_bin) low_bin = high_bin;
sp(1) = low_bin;
sp(2) = high_bin;
temp=mean(tempvec_a(low_bin,high_bin));
scaling_offset = 0; // reset scaling offset
}
tempvec_a-=temp; // rescale to get max of 0.0
tempvec_a(j2+1,nages) = tempvec_a(j2); // set constant above last node
sel_a(y,fs,1)=mfexp(tempvec_a);
break;
}

```

Deviation Link

Code for alternative deviation links shown below:

```

case 1 (multiplicative):
{
for (j=timevary_setup(10);j<=timevary_setup(11);j++)
{
parm_timevary(tvary,j)*=mfexp(parm_dev(k,j)*parm_dev_stddev(k));
}
break;
}

case 2 (additive):
{
for (j=timevary_setup(10);j<=timevary_setup(11);j++)
{
parm_timevary(tvary,j)+=parm_dev(k,j)*parm_dev_stddev(k);
}
break;
}

case 3 (random walk):
{
parm_dev_rwalk(k,timevary_setup(10))=parm_dev(k,timevary_setup(10))*parm_dev_stddev(k);
}

```

```

    parm_timevary(tvary,timevary_setup(10))+=parm_dev_rwalk(k,timevary_setup(10));
    for (j=timevary_setup(10)+1;j<=timevary_setup(11);j++)
    {
        parm_dev_rwalk(k,j)=parm_dev_rwalk(k,j-1)+parm_dev(k,j)*parm_dev_stddev(k);
        parm_timevary(tvary,j)+=parm_dev_rwalk(k,j);
    }
    break;
}

case 4 (mean reverting random walk)
{
    parm_dev_rwalk(k,timevary_setup(10))=parm_dev(k,timevary_setup(10))*parm_dev_stddev(k);
    parm_timevary(tvary,timevary_setup(10))+=parm_dev_rwalk(k,timevary_setup(10));
    for (j=timevary_setup(10)+1;j<=timevary_setup(11);j++)
    {
        //      =(1-rho)*mean + rho*prevval + dev    //  where mean = 0.0
        parm_dev_rwalk(k,j)=parm_dev_rho(k)*parm_dev_rwalk(k,j-1)+parm_dev(k,j)*parm_dev_stddev(k);
        parm_timevary(tvary,j)+=parm_dev_rwalk(k,j);
    }
    break;
}

```

Bibliography

- Francis, R.C. and Hilborn, R. 2011. Data weighting in statistical fisheries stock assessment models. *Canadian Journal of Fisheries and Aquatic Sciences* **68**(6): 1124–1138. doi:10.1139/f2011-025.
- Gertseva, V. and Thorson, James T 2013. Status of the darkblotched rockfish resource off the continental U.S. Pacific Coast in 2013. Technical report, Pacific Fishery Management Council, 7700 Ambassador Place NE, Suite 200, Portland, OR 97220.
- Grandin, C.J., Johnson, K.F., Edwards, A.M., and Berger, A.M. 2020. Status of Pacific Hake (whiting) stock in the U.S. and Canadian waters in 2020. Technical report, Joint Technical Committee of the Pacific Hake/Whiting Agreement Between the Governments of the United States and Canada, National Marine Fisheries Service and Fisheries and Ocean Canada.
- Johnson, K.F., Councill, E., Thorson, J.T., Brooks, E., Methot, R.D., and Punt, A.E. 2016. Can autocorrelated recruitment be estimated using integrated assessment models and how does it affect population forecasts? *Fisheries Research* **183**: 222–232. doi:10.1016/j.fishres.2016.06.004.
- Lee, H.H., Thomas, L.R., Piner, K.R., and Maunder, M.N. 2017. Effects of age-based movement on the estimation of growth assuming random-at-age or random-at-length data. *Journal of Fish Biology* **90**(1): 222–235. doi:10.1111/jfb.13177.
- Lee, H., Piner, K.R., Taylor, I.G., and Kitakado, T. 2019. On the use of conditional age at length data as a likelihood component in integrated population dynamics models. *Fisheries Research* **216**: 204–211. doi:10.1016/j.fishres.2019.04.007.
- Maunder, M.N. and Deriso, R.B. 2003. Estimation of recruitment in catch-at-age models. *Canadian Journal of Fisheries and Aquatic Sciences* **60**(10): 1204–1216. doi:10.1139/f03-104.
- Maunder, M.N., Deriso, R.B., Schaefer, K.M., Fuller, D.W., Aires-da Silva, A.M., Minte-Vera, C.V., and Campana, S.E. 2018. The growth cessation model: a growth model for species showing a near cessation in growth with application to bigeye tuna (*Thunnus obesus*). *Marine Biology* **165**(4). doi:10.1007/s00227-018-3336-9.
- McAllister, M.K. and Ianelli, J.N. 1997. Bayesian stock assessment using catch-age data and the sampling - importance resampling algorithm. *Canadian Journal of Fisheries and Aquatic Sciences* **54**: 284–300.
- Methot, R.D., Taylor, I.G., and Chen, Y. 2011. Adjusting for bias due to variability of estimated recruitments in fishery assessment models. *Canadian Journal of Fisheries and Aquatic Sciences* **68**(10): 1744–1760. doi:10.1139/f2011-092.
- Methot, R.D. and Wetzel, C.R. 2013. Stock synthesis: A biological and statistical framework for fish stock assessment and fishery management. *Fisheries Research* **142**: 86–99. doi:10.1016/j.fishres.2012.10.012.

- Piner, K.R., Lee, H.H., and Maunder, M.N. 2016. Evaluation of using random-at-length observations and an equilibrium approximation of the population age structure in fitting the von Bertalanffy growth function. *Fisheries Research* **180**: 128–137. doi:10.1016/j.fishres.2015.05.024.
- Punt, A.E. and Cope, J.M. 2019. Extending integrated stock assessment models to use non-depensatory three-parameter stock-recruitment relationships. *Fisheries Research* **217**: 46–57. doi:10.1016/j.fishres.2017.07.007.
- Taylor, I.G., Gertseva, V., Methot, R.D., and Maunder, M.N. 2013. A stock–recruitment relationship based on pre-recruit survival, illustrated with application to spiny dogfish shark. *Fisheries Research* **142**: 15–21. doi:10.1016/j.fishres.2012.04.018.
- Thorson, J.T., Johnson, K.F., Methot, R.D., and Taylor, I.G. 2017. Model-based estimates of effective sample size in stock assessment models using the Dirichlet-multinomial distribution. *Fisheries Research* **192**: 84–93. doi:10.1016/j.fishres.2016.06.005.
- Xu, H., Thorson, J.T., Methot, R.D., and Taylor, I.G. 2019. A new semi-parametric method for autocorrelated age- and time-varying selectivity in age-structured assessment models. *Canadian Journal of Fisheries and Aquatic Sciences* **76**(2): 268–285. doi:10.1139/cjfas-2017-0446.



U.S. Secretary of Commerce
Wilbur L. Ross, Jr.

Acting Under Secretary of Commerce
for Oceans and Atmosphere
Dr. Neil Jacobs

Assistant Administrator for Fisheries
Chris Oliver

April 2020

[fisheries.noaa.gov](https://www.fisheries.noaa.gov)

OFFICIAL BUSINESS

National Marine
Fisheries Service
Northwest Fisheries Science Center
2725 Montlake Boulevard East
Seattle, Washington 98112